

Step-7 (Track A) — Horizon x Cov-Duration Bake-off Master Dashboard (Sigma 30s / 45s / 60s, MASKED)

2022-2025 · cov-dur x horizon(nextwin..1h) x metric(LE/AI) x 15 learners · fixed Step-6 predictor [rv, s_slow, s_fast, v]; realized-VOL target aggregated to horizon · AI headline / LE diagnostic · next-window COINCIDENT, no next-day forecast (deferred Track B) · cells finished: 54 · durations present: 30s / 45s / 60s

What this run is. A **cov-duration x prediction-horizon** bake-off. No new covariance is built: the 30/45/60 s **masked Sigma state** (Steps 3-6) is unchanged; Step 7 moves ONE axis - the target's forward SPAN H (nextwin, 2m, 5m, 10m, 15m, 30m, 1h). For every (cov-dur x horizon x metric) cell we fit the COMPLETE 15-learner panel (GridSearch + purged walk-forward) with the **fixed Step-6 predictor** [rv, s_slow, s_fast, v] and a realized-VOL target aggregated to H. **LE** features = Log-Euclidean Step-6 dial; **AI** features = affine-invariant dial (PC1 of the AI whitened tangent) + AI geodesic velocity **AI** **LE** . **Every number below is labelled with its cov-duration AND its horizon.**

The claim being tested. Holding the covariance state fixed: (1) at which HORIZON does the dial stop being coincident and start adding OOS value over persistence (the **dR2(H)** curve); (2) does a FLEXIBLE learner beat linear at that horizon (capacity); (3) is the answer METRIC-ROBUST (AI vs LE)? **next-window is COINCIDENT** (persistence wins); longer horizons test whether the dial becomes LEADING. There is NO next-day forecast here - that is the deferred Track B.

30s / 45s / 60s
cov durations

nextwin . 30s . 45s . 60s . 2m . 5m . 10m . 15m . 30m . 1h
horizons

LE / AI
metrics

15
learners / cell

54 / 60
cells finished

6
cells N/A (skipped)

30s / 45s / 60s
durations present

Coverage grid (cov-duration x horizon). which cells the detached run has finished (LE+AI green / one metric amber), which are physically SKIPPED (N/A = horizon finer than the cov window), and which are still pending (—).

Sigma \ horizon	nextwin	30s	45s	60s	2m	5m	10m	15m	30m	1h
30s	LE+AI	LE+AI	LE+AI	LE+AI	LE+AI	LE+AI	LE+AI	LE+AI	LE+AI	LE+AI
45s	LE+AI	N/A	LE+AI	LE+AI	LE+AI	LE+AI	LE+AI	LE+AI	LE+AI	LE+AI
60s	LE+AI	N/A	N/A	LE+AI	LE+AI	LE+AI	LE+AI	LE+AI	LE+AI	LE+AI

Peak lead-vs-coincident edge (per cov-dur x metric). the horizon where the best learner's dR2 over persistence is largest, and whether its 95% CI clears 0.


Sigma	metric	horizons done	peak dR2 horizon	peak dR2	sig (CI>0)
30s	LE	10 / 10	1h	0.136	True
30s	AI	10 / 10	1h	0.129	True
45s	LE	9 / 10	1h	0.107	True
45s	AI	9 / 10	1h	0.099	True
60s	LE	8 / 10	1h	0.087	True
60s	AI	8 / 10	1h	0.079	True

THE headline of Track A - lead-vs-coincident. For each cov-duration, plot the dial's OOS edge over persistence **dR2(horizon)** as the target span H grows (x = nextwin, 2m, 5m, 10m, 15m, 30m, 1h). Solid = the BEST learner's dR2; dotted = the linear learner's dR2. LE (blue) vs AI (green) **AI** **LE** . **The question:** at next-window the dial is COINCIDENT (persistence wins, dR2 small); does dR2 GROW with horizon (dial becomes LEADING)? The marked star = the horizon that PEAKS dR2 (the characteristic horizon).

dR2(horizon) per cov-duration — native 30s / 45s / 60s Sigma; LE vs AI

Sigma=30s dR2(horizon) lead-vs-coincident

dR2(horizon) -- lead-vs-coincident [Sigma=30s]



In plain terms How much extra vol-forecast accuracy the covariance dial adds over pure persistence, at each forward span.


Math / formula $dR2(H) = R2_OOS(\text{best learner}) - R2_OOS(\text{persistence}=rv\text{-only})$; block-bootstrap 95% CI over days.

How to read it flat/low near nextwin = coincident; rising with H = the dial becomes leading; star marks the peak H.

Insight LE: peak dR2 **0.136** @ 1h (nextwin 0.075) [sig] | AI: peak dR2 **0.129** @ 1h (nextwin 0.075) [sig]

Sigma=45s dR2(horizon) lead-vs-coincident

dR2(horizon) -- lead-vs-coincident [Sigma=45s]



In plain terms How much extra vol-forecast accuracy the covariance dial adds over pure persistence, at each forward span.


Math / formula $dR2(H) = R2_OOS(\text{best learner}) - R2_OOS(\text{persistence}=rv\text{-only})$; block-bootstrap 95% CI over days.

How to read it flat/low near nextwin = coincident; rising with H = the dial becomes leading; star marks the peak H.

Insight LE: peak dR2 **0.107** @ 1h (nextwin 0.053) [sig] | AI: peak dR2 **0.099** @ 1h (nextwin 0.050) [sig]

Sigma=60s dR2(horizon) lead-vs-coincident

dR2(horizon) -- lead-vs-coincident [Sigma=60s]



In plain terms How much extra vol-forecast accuracy the covariance dial adds over pure persistence, at each forward span.

Math / formula $dR2(H) = R2_OOS(\text{best learner}) - R2_OOS(\text{persistence}=rv\text{-only})$; block-bootstrap 95% CI over days.


How to read it flat/low near nextwin = coincident; rising with H = the dial becomes leading; star marks the peak H.

Insight LE: peak dR2 **0.087** @ 1h (nextwin 0.038) [sig] | AI: peak dR2 **0.079** @ 1h (nextwin 0.037) [sig]

Performance-comparison plots. The same numbers as the tables, drawn for the eye. Every chart is labelled with its cov-duration / horizon / metric; horizons are always in canonical order (nextwin, 30s, 45s, 60s, 2m, 5m, 10m, 15m, 30m, 1h); skipped cells (horizon finer than the cov window) are greyed/omitted. Colors = learners (svr = the usual winner, highlighted).

1 · OOS test R2 vs horizon — 15 learners + persistence — one chart per cov-duration x metric

Sigma=30s OOS R2 vs horizon [Sigma=30s | metric=LE] LE




In plain terms How each of the 15 learners tracks vol as the forecast span grows, vs the persistence floor.

Math / formula $y = R2_OOS(model)$ per horizon; dashed white = persistence (rv-only) baseline.

How to read it lines above the dashed floor beat persistence; watch the interior peak (~10m).

Insight cov 30s, metric LE; missing points = unfinished/skipped horizons.

Sigma=30s OOS R2 vs horizon [Sigma=30s | metric=AI] AI




In plain terms How each of the 15 learners tracks vol as the forecast span grows, vs the persistence floor.

Math / formula $y = R2_OOS(model)$ per horizon; dashed white = persistence (rv-only) baseline.

How to read it lines above the dashed floor beat persistence; watch the interior peak (~10m).

Insight cov 30s, metric AI; missing points = unfinished/skipped horizons.

Sigma=45s OOS R2 vs horizon [Sigma=45s | metric=LE] LE




In plain terms How each of the 15 learners tracks vol as the forecast span grows, vs the persistence floor.

Math / formula $y = R2_OOS(model)$ per horizon; dashed white = persistence (rv-only) baseline.

How to read it lines above the dashed floor beat persistence; watch the interior peak (~10m).

Sigma=45s OOS R2 vs horizon [Sigma=45s | metric=AI] AI




In plain terms How each of the 15 learners tracks vol as the forecast span grows, vs the persistence floor.

Math / formula $y = R2_OOS(model)$ per horizon; dashed white = persistence (rv-only) baseline.

How to read it lines above the dashed floor beat persistence; watch the interior peak (~10m).

Insight cov 45s, metric AI; missing points = unfinished/skipped horizons.

Sigma=60s OOS R2 vs horizon [Sigma=60s | metric=LE] LE




In plain terms How each of the 15 learners tracks vol as the forecast span grows, vs the persistence floor.

Math / formula $y = R2_OOS(model)$ per horizon; dashed white = persistence (rv-only) baseline.

How to read it lines above the dashed floor beat persistence; watch the interior peak (~10m).

Insight cov 60s, metric LE; missing points = unfinished/skipped horizons.

Sigma=60s OOS R2 vs horizon [Sigma=60s | metric=AI] AI




In plain terms How each of the 15 learners tracks vol as the forecast span grows, vs the persistence floor.

Math / formula $y = R2_OOS(model)$ per horizon; dashed white = persistence (rv-only) baseline.

How to read it lines above the dashed floor beat persistence; watch the interior peak (~10m).

Insight cov 60s, metric AI; missing points = unfinished/skipped horizons.

Sigma=30s Cross-duration: best-model & persistence R2 vs horizon [metric=LE] LE




In plain terms Does a coarser covariance window (60s) give a higher absolute R2 than a finer one (30s)?

Math / formula left: best-learner R2 per horizon for 30/45/60s; right: the persistence floor per horizon.

How to read it curves ordered 60s>45s>30s = coarser Sigma is more predictable (both model and floor rise).

Insight metric LE; both panels share the horizon axis.

Sigma=30s Cross-duration: best-model & persistence R2 vs horizon [metric=AI] AI




In plain terms Does a coarser covariance window (60s) give a higher absolute R2 than a finer one (30s)?

Math / formula left: best-learner R2 per horizon for 30/45/60s; right: the persistence floor per horizon.

How to read it curves ordered 60s>45s>30s = coarser Sigma is more predictable (both model and floor rise).

Insight metric AI; both panels share the horizon axis.

Sigma=30s Learner ranking [Sigma=30s | metric=LE] LE




In plain terms Which learner wins at the interesting horizons, and by how much.

Math / formula bars = R2_OOS per learner, sorted descending; amber = svr (the usual winner).

How to read it a tall gap from #1 to linear = capacity helps; a flat top = many models tie.

Insight cov 30s, metric LE, horizons 10m, 1h.

Sigma=30s Learner ranking [Sigma=30s | metric=AI] AI




In plain terms Which learner wins at the interesting horizons, and by how much.

Math / formula bars = R2_OOS per learner, sorted descending; amber = svr (the usual winner).

How to read it a tall gap from #1 to linear = capacity helps; a flat top = many models tie.

Insight cov 30s, metric AI, horizons 10m, 1h.

Sigma=45s Learner ranking [Sigma=45s | metric=LE] LE




In plain terms Which learner wins at the interesting horizons, and by how much.

Math / formula bars = R2_OOS per learner, sorted descending; amber = svr (the usual winner).

How to read it a tall gap from #1 to linear = capacity helps; a flat top = many models tie.

Insight cov 45s, metric LE, horizons 10m, 1h.

Sigma=45s Learner ranking [Sigma=45s | metric=AI] AI




In plain terms Which learner wins at the interesting horizons, and by how much.


Math / formula bars = R2_OOS per learner, sorted descending; amber = svr (the usual winner).

How to read it a tall gap from #1 to linear = capacity helps; a flat top = many models tie.

Insight cov 45s, metric AI, horizons 10m, 1h.

Sigma=60s Learner ranking [Sigma=60s | metric=LE] LE






In plain terms Which learner wins at the interesting horizons, and by how much.

Math / formula bars = R2_OOS per learner, sorted descending; amber = svr (the usual winner).

How to read it a tall gap from #1 to linear = capacity helps; a flat top = many models tie.

Insight cov 60s, metric LE, horizons 10m, 1h.

Sigma=60s Learner ranking [Sigma=60s | metric=AI] **AI**



In plain terms Which learner wins at the interesting horizons, and by how much.


Math / formula bars = R2_OOS per learner, sorted descending; amber = svr (the usual winner).

How to read it a tall gap from #1 to linear = capacity helps; a flat top = many models tie.

Insight cov 60s, metric AI, horizons 10m, 1h.

6 · OOS R2 heatmap — learner x horizon — per cov-duration x metric; skipped horizons greyed

Sigma=30s OOS R2 heatmap [Sigma=30s | metric=LE] **LE**



In plain terms The whole learner x horizon accuracy grid at a glance.


Math / formula cell colour = R2_OOS(learner, horizon); grey = skipped (horizon < cov window).

How to read it bright bands = strong learner/horizon combos; a bright column = an easy horizon.

Insight cov 30s, metric LE; 15 learners x 10 finished horizons.

Sigma=30s OOS R2 heatmap [Sigma=30s | metric=AI] **AI**






In plain terms The whole learner x horizon accuracy grid at a glance.

Math / formula cell colour = $R2_OOS(learner, horizon)$; grey = skipped (horizon < cov window).

How to read it bright bands = strong learner/horizon combos; a bright column = an easy horizon.

Insight cov 30s, metric AI; 15 learners x 10 finished horizons.

Sigma=45s OOS R2 heatmap [Sigma=45s | metric=LE] LE




In plain terms The whole learner x horizon accuracy grid at a glance.

Math / formula cell colour = $R2_OOS(learner, horizon)$; grey = skipped (horizon < cov window).

How to read it bright bands = strong learner/horizon combos; a bright column = an easy horizon.

Insight cov 45s, metric LE; 15 learners x 9 finished horizons.

Sigma=45s OOS R2 heatmap [Sigma=45s | metric=AI] AI



In plain terms The whole learner x horizon accuracy grid at a glance.


Math / formula cell colour = $R2_OOS(learner, horizon)$; grey = skipped (horizon < cov window).

How to read it bright bands = strong learner/horizon combos; a bright column = an easy horizon.

Insight cov 45s, metric AI; 15 learners x 9 finished horizons.

Sigma=60s OOS R2 heatmap [Sigma=60s | metric=LE] LE






In plain terms The whole learner x horizon accuracy grid at a glance.

Math / formula cell colour = $R2_OOS(learner, horizon)$; grey = skipped (horizon < cov window).

How to read it bright bands = strong learner/horizon combos; a bright column = an easy horizon.

Insight cov 60s, metric LE; 15 learners x 8 finished horizons.

Sigma=60s OOS R2 heatmap [Sigma=60s | metric=AI] AI



In plain terms The whole learner x horizon accuracy grid at a glance.


Math / formula cell colour = $R2_OOS(learner, horizon)$; grey = skipped (horizon < cov window).

How to read it bright bands = strong learner/horizon combos; a bright column = an easy horizon.

Insight cov 60s, metric AI; 15 learners x 8 finished horizons.

7 · Overfit view — train R2 vs test R2 per learner — per cov-duration x metric; points off the diagonal overfit

Sigma=30s Overfit: train vs test R2 [Sigma=30s | metric=LE] LE




In plain terms How much each learner's in-sample fit exceeds its out-of-sample accuracy (the gap = overfit).

Math / formula x = train R2, y = test R2; dashed = no gap; distance BELOW the line = overfit.

How to read it points far under the diagonal (watch gbm, red) overfit; on/near it generalise.

Insight cov 30s, metric LE, horizons 2m, 10m.

Sigma=30s Overfit: train vs test R2 [Sigma=30s | metric=AI] AI




In plain terms How much each learner's in-sample fit exceeds its out-of-sample accuracy (the gap = overfit).

Math / formula x = train R2, y = test R2; dashed = no gap; distance BELOW the line = overfit.

How to read it points far under the diagonal (watch gbm, red) overfit; on/near it generalise.

Insight cov 30s, metric AI, horizons 2m, 10m.

Sigma=45s Overfit: train vs test R2 [Sigma=45s | metric=LE] **LE**




In plain terms How much each learner's in-sample fit exceeds its out-of-sample accuracy (the gap = overfit).

Math / formula x = train R2, y = test R2; dashed = no gap; distance BELOW the line = overfit.

How to read it points far under the diagonal (watch gbm, red) overfit; on/near it generalise.

Insight cov 45s, metric LE, horizons 2m, 10m.

Sigma=45s Overfit: train vs test R2 [Sigma=45s | metric=AI] **AI**




In plain terms How much each learner's in-sample fit exceeds its out-of-sample accuracy (the gap = overfit).

Math / formula x = train R2, y = test R2; dashed = no gap; distance BELOW the line = overfit.

How to read it points far under the diagonal (watch gbm, red) overfit; on/near it generalise.

Insight cov 45s, metric AI, horizons 2m, 10m.

Sigma=60s Overfit: train vs test R2 [Sigma=60s | metric=LE] **LE**




In plain terms How much each learner's in-sample fit exceeds its out-of-sample accuracy (the gap = overfit).

Math / formula $x = \text{train R2}, y = \text{test R2}$; dashed = no gap; distance BELOW the line = overfit.

How to read it points far under the diagonal (watch gbm, red) overfit; on/near it generalise.

Insight cov 60s, metric LE, horizons 2m, 10m.

Sigma=60s Overfit: train vs test R2 [Sigma=60s | metric=AI] AI



In plain terms How much each learner's in-sample fit exceeds its out-of-sample accuracy (the gap = overfit).


Math / formula $x = \text{train R2}, y = \text{test R2}$; dashed = no gap; distance BELOW the line = overfit.

How to read it points far under the diagonal (watch gbm, red) overfit; on/near it generalise.

Insight cov 60s, metric AI, horizons 2m, 10m.

8 · QLIKE vs horizon — 15 learners (lower is better) — per cov-duration x metric

Sigma=30s QLIKE vs horizon [Sigma=30s | metric=LE] LE




In plain terms Does the accuracy story survive the stricter QLIKE variance loss (not just R2)?

Math / formula $y = \text{QLIKE_OOS}(\text{model})$ per horizon = $\text{mean}(r - \ln r - 1)$, r = variance ratio; lower = better.

How to read it lines that stay low across horizons are robust; spikes flag loss-sensitive learners.

Insight cov 30s, metric LE.

Sigma=30s QLIKE vs horizon [Sigma=30s | metric=AI] AI




In plain terms Does the accuracy story survive the stricter QLIKE variance loss (not just R2)?

Math / formula $y = \text{QLIKE_OOS}(\text{model}) \text{ per horizon} = \text{mean}(r - \ln r - 1)$, r = variance ratio; lower = better.

How to read it lines that stay low across horizons are robust; spikes flag loss-sensitive learners.

Insight cov 30s, metric AI.

Sigma=45s QLIKE vs horizon [Sigma=45s | metric=LE] LE




In plain terms Does the accuracy story survive the stricter QLIKE variance loss (not just R2)?

Math / formula $y = \text{QLIKE_OOS}(\text{model}) \text{ per horizon} = \text{mean}(r - \ln r - 1)$, r = variance ratio; lower = better.

How to read it lines that stay low across horizons are robust; spikes flag loss-sensitive learners.

Insight cov 45s, metric LE.

Sigma=45s QLIKE vs horizon [Sigma=45s | metric=AI] AI




In plain terms Does the accuracy story survive the stricter QLIKE variance loss (not just R2)?


Math / formula $y = \text{QLIKE_OOS}(\text{model}) \text{ per horizon} = \text{mean}(r - \ln r - 1)$, r = variance ratio; lower = better.

How to read it lines that stay low across horizons are robust; spikes flag loss-sensitive learners.

Insight cov 45s, metric AI.

Sigma=60s QLIKE vs horizon [Sigma=60s | metric=LE] LE





In plain terms

Does the accuracy story survive the stricter QLIKE variance loss (not just R2)?

Math / formula

$y = \text{QLIKE_OOS}(\text{model}) \text{ per horizon} = \text{mean}(r - \ln r - 1)$, r = variance ratio; lower = better.


How to read it

lines that stay low across horizons are robust; spikes flag loss-sensitive learners.

Insight

cov 60s, metric LE.

Sigma=60s QLIKE vs horizon [Sigma=60s | metric=AI] **AI**



In plain terms

Does the accuracy story survive the stricter QLIKE variance loss (not just R2)?

Math / formula

$y = \text{QLIKE_OOS}(\text{model}) \text{ per horizon} = \text{mean}(r - \ln r - 1)$, r = variance ratio; lower = better.

How to read it

lines that stay low across horizons are robust; spikes flag loss-sensitive learners.

Insight

cov 60s, metric AI.

The full matrix. For EVERY (cov-dur x horizon x metric) cell, all 15 learners x {train R2, test R2, gap, dR2 [CI95], sig, RMSE, MAE, QLIKE, bias}. Grouped by cov-duration; within each, a labelled block per horizon with LE and AI side by side. Every table title carries **Sigma=<dur>s | horizon=<H> | metric=<LE/AI>**. The star row = best-by-test-R2. RMSE/MAE/bias in vol units; QLIKE on variance.

Sigma=30s cov-duration 30s — 10 horizon(s) finished

horizon=nextwin Sigma=30s · horizon=nextwin (H=30s) — metrics: LE + AI

LE Sigma=30s | horizon=nextwin | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.404	0.471	-0.067	0.044 [0.040, 0.047]	yes	2.31e-04	1.85e-04	5.33e-08	0.646	5.32e-05	2.75e-04	2.12e-04	7.54e-08	0.901	1.41e-16
ridge	0.404	0.471	-0.067	0.044 [0.040, 0.047]	yes	2.31e-04	1.85e-04	5.33e-08	0.645	5.33e-05	2.75e-04	2.12e-04	7.54e-08	0.901	-5.96e-17
lasso	0.401	0.459	-0.058	0.031 [0.029, 0.034]	yes	2.34e-04	1.87e-04	5.45e-08	0.586	5.74e-05	2.75e-04	2.12e-04	7.59e-08	0.816	-5.63e-21
elasticnet	0.390	0.442	-0.052	0.014 [0.012, 0.016]	yes	2.37e-04	1.89e-04	5.63e-08	0.553	6.13e-05	2.78e-04	2.14e-04	7.72e-08	0.765	-1.90e-20
huber	0.402	0.484	-0.082	0.056 [0.054, 0.059]	yes	2.28e-04	1.82e-04	5.20e-08	0.796	3.58e-05	2.75e-04	2.11e-04	7.57e-08	1.110	-1.69e-05
poly2	0.425	0.487	-0.062	0.060 [0.055, 0.065]	yes	2.27e-04	1.82e-04	5.17e-08	0.547	5.83e-05	2.70e-04	2.08e-04	7.27e-08	0.766	4.62e-15
poly3	0.423	0.478	-0.054	0.050 [0.047, 0.054]	yes	2.29e-04	1.84e-04	5.26e-08	0.568	6.19e-05	2.70e-04	2.09e-04	7.30e-08	0.816	1.45e-13
knn	0.455	0.463	-0.008	0.036 [0.029, 0.042]	yes	2.33e-04	1.86e-04	5.41e-08	0.513	6.78e-05	2.63e-04	2.02e-04	6.90e-08	0.715	-1.32e-06
svr ★	0.453	0.502	-0.050	0.075 [0.070, 0.080]	yes	2.24e-04	1.79e-04	5.01e-08	0.777	5.10e-05	2.63e-04	2.02e-04	6.93e-08	1.032	-1.58e-05
kernelridge	0.452	0.482	-0.029	0.054 [0.045, 0.062]	yes	2.29e-04	1.83e-04	5.22e-08	0.706	6.52e-05	2.63e-04	2.02e-04	6.93e-08	0.986	-1.25e-06
gbm	0.360	0.366	-0.005	-0.062 [-0.070, -0.052]	no	2.53e-04	2.03e-04	6.39e-08	0.405	9.48e-05	2.85e-04	2.20e-04	8.10e-08	0.532	1.26e-08
rf	0.467	0.482	-0.016	0.055 [0.048, 0.062]	yes	2.28e-04	1.84e-04	5.22e-08	0.483	6.78e-05	2.60e-04	2.01e-04	6.75e-08	0.656	-2.98e-08
extratrees	0.492	0.477	0.015	0.050 [0.043, 0.056]	yes	2.29e-04	1.85e-04	5.27e-08	0.473	7.10e-05	2.53e-04	1.95e-04	6.43e-08	0.596	-8.80e-22
mlp	0.466	0.482	-0.017	0.055 [0.047, 0.062]	yes	2.28e-04	1.84e-04	5.22e-08	0.499	7.02e-05	2.60e-04	2.00e-04	6.76e-08	0.671	1.92e-06
gp	0.444	0.466	-0.023	0.039 [0.032, 0.045]	yes	2.32e-04	1.86e-04	5.38e-08	0.502	7.33e-05	2.65e-04	2.04e-04	7.04e-08	0.718	-5.28e-08

best-by-test-R2 = svr (0.502) · linear 0.471 · persistence 0.427 · H=30s · 2774.4s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI Sigma=30s | horizon=nextwin | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.402	0.468	-0.066	0.040 [0.037, 0.044]	yes	2.32e-04	1.85e-04	5.37e-08	0.658	5.28e-05	2.75e-04	2.13e-04	7.57e-08	0.904	-4.56e-16
ridge	0.402	0.468	-0.066	0.040 [0.037, 0.044]	yes	2.32e-04	1.85e-04	5.37e-08	0.657	5.29e-05	2.75e-04	2.13e-04	7.57e-08	0.904	-6.13e-17
lasso	0.398	0.456	-0.058	0.029 [0.026, 0.032]	yes	2.34e-04	1.87e-04	5.48e-08	0.591	5.76e-05	2.76e-04	2.13e-04	7.61e-08	0.818	7.04e-22
elasticnet	0.389	0.440	-0.051	0.013 [0.011, 0.015]	yes	2.37e-04	1.90e-04	5.64e-08	0.552	6.23e-05	2.78e-04	2.14e-04	7.73e-08	0.763	-1.62e-20
huber	0.399	0.480	-0.081	0.052 [0.050, 0.055]	yes	2.29e-04	1.83e-04	5.24e-08	0.807	3.56e-05	2.76e-04	2.12e-04	7.61e-08	1.117	-1.69e-05
poly2	0.422	0.480	-0.058	0.053 [0.048, 0.058]	yes	2.29e-04	1.84e-04	5.24e-08	0.537	5.98e-05	2.70e-04	2.09e-04	7.31e-08	0.757	3.14e-15
poly3	0.421	0.469	-0.048	0.042 [0.037, 0.046]	yes	2.31e-04	1.86e-04	5.35e-08	0.560	6.50e-05	2.71e-04	2.09e-04	7.32e-08	0.792	-1.27e-13
knn	0.453	0.466	-0.013	0.039 [0.033, 0.044]	yes	2.32e-04	1.86e-04	5.38e-08	0.507	6.72e-05	2.63e-04	2.03e-04	6.92e-08	0.717	-1.93e-06
svr ★	0.450	0.502	-0.052	0.075 [0.070, 0.079]	yes	2.24e-04	1.79e-04	5.02e-08	0.747	5.04e-05	2.64e-04	2.02e-04	6.96e-08	1.001	-1.54e-05
kernelridge	0.451	0.482	-0.031	0.054 [0.047, 0.061]	yes	2.29e-04	1.83e-04	5.22e-08	0.694	6.50e-05	2.64e-04	2.03e-04	6.95e-08	0.942	-1.23e-06
gbm	0.359	0.364	-0.005	-0.063 [-0.072, -0.053]	no	2.53e-04	2.03e-04	6.41e-08	0.408	9.38e-05	2.85e-04	2.20e-04	8.12e-08	0.535	1.31e-08
rf	0.463	0.480	-0.016	0.052 [0.046, 0.059]	yes	2.29e-04	1.84e-04	5.24e-08	0.473	6.82e-05	2.61e-04	2.01e-04	6.79e-08	0.657	-3.52e-08
extratrees	0.490	0.474	0.015	0.047 [0.040, 0.053]	yes	2.30e-04	1.85e-04	5.30e-08	0.459	7.29e-05	2.54e-04	1.96e-04	6.46e-08	0.596	-4.05e-21
mlp	0.463	0.481	-0.018	0.054 [0.047, 0.061]	yes	2.29e-04	1.84e-04	5.23e-08	0.504	6.89e-05	2.61e-04	2.01e-04	6.80e-08	0.700	-9.00e-07
gp	0.441	0.469	-0.027	0.042 [0.036, 0.047]	yes	2.31e-04	1.85e-04	5.35e-08	0.509	7.12e-05	2.66e-04	2.05e-04	7.07e-08	0.704	8.62e-08

best-by-test-R2 = svr (0.502) · linear 0.468 · persistence 0.427 · H=30s · 2728.3s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=30s Sigma=30s · horizon=30s (H=30s) — metrics: LE + AI

LE Sigma=30s | horizon=30s | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.400	0.468	-0.068	0.046 [0.042, 0.050]	yes	2.26e-04	1.82e-04	5.09e-08	0.630	5.31e-05	2.67e-04	2.08e-04	7.15e-08	0.883	2.73e-16
ridge	0.400	0.468	-0.068	0.046 [0.042, 0.050]	yes	2.26e-04	1.82e-04	5.09e-08	0.630	5.31e-05	2.67e-04	2.08e-04	7.15e-08	0.883	-4.86e-17
lasso	0.397	0.456	-0.059	0.033 [0.030, 0.036]	yes	2.28e-04	1.84e-04	5.21e-08	0.570	5.72e-05	2.68e-04	2.09e-04	7.20e-08	0.799	-3.79e-20
elasticnet	0.385	0.437	-0.052	0.015 [0.013, 0.017]	yes	2.32e-04	1.87e-04	5.39e-08	0.537	6.12e-05	2.71e-04	2.10e-04	7.33e-08	0.748	-4.31e-20
huber	0.398	0.481	-0.083	0.058 [0.056, 0.061]	yes	2.23e-04	1.79e-04	4.97e-08	0.770	3.65e-05	2.68e-04	2.08e-04	7.18e-08	1.081	-1.60e-05
poly2	0.422	0.485	-0.063	0.062 [0.058, 0.067]	yes	2.22e-04	1.80e-04	4.93e-08	0.557	5.81e-05	2.62e-04	2.05e-04	6.89e-08	0.759	7.81e-15
poly3	0.421	0.476	-0.055	0.054 [0.050, 0.058]	yes	2.24e-04	1.81e-04	5.01e-08	0.521	6.15e-05	2.63e-04	2.05e-04	6.90e-08	0.825	1.43e-13
knn	0.452	0.460	-0.008	0.037 [0.031, 0.043]	yes	2.27e-04	1.84e-04	5.17e-08	0.505	6.66e-05	2.56e-04	1.99e-04	6.54e-08	0.704	-1.53e-06
svr ★	0.451	0.505	-0.055	0.083 [0.078, 0.087]	yes	2.18e-04	1.75e-04	4.74e-08	0.968	4.53e-05	2.56e-04	1.98e-04	6.55e-08	0.962	-1.48e-05
kernelridge	0.451	0.486	-0.034	0.063 [0.057, 0.069]	yes	2.22e-04	1.79e-04	4.93e-08	0.688	6.34e-05	2.56e-04	1.99e-04	6.54e-08	0.718	-1.18e-06
gbm	0.357	0.362	-0.005	-0.061 [-0.069, -0.051]	no	2.47e-04	2.00e-04	6.11e-08	0.395	9.35e-05	2.77e-04	2.16e-04	7.67e-08	0.523	8.70e-08
rf	0.464	0.480	-0.017	0.058 [0.050, 0.065]	yes	2.23e-04	1.81e-04	4.98e-08	0.478	6.65e-05	2.53e-04	1.97e-04	6.40e-08	0.648	-8.85e-10
extratrees	0.490	0.474	0.016	0.051 [0.044, 0.058]	yes	2.24e-04	1.82e-04	5.04e-08	0.461	7.09e-05	2.47e-04	1.92e-04	6.08e-08	0.588	1.86e-19
mlp	0.463	0.483	-0.019	0.060 [0.052, 0.067]	yes	2.23e-04	1.80e-04	4.95e-08	0.509	6.67e-05	2.53e-04	1.97e-04	6.40e-08	0.683	-1.75e-06
gp	0.440	0.489	-0.048	0.066 [0.062, 0.070]	yes	2.21e-04	1.79e-04	4.90e-08	0.546	6.26e-05	2.58e-04	2.01e-04	6.68e-08	0.759	-4.34e-06

best-by-test-R2 = svr (0.505) · linear 0.468 · persistence 0.423 · H=30s · 2676.5s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI Sigma=30s | horizon=30s | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.398	0.465	-0.067	0.042 [0.038, 0.046]	yes	2.26e-04	1.82e-04	5.12e-08	0.639	5.28e-05	2.68e-04	2.09e-04	7.18e-08	0.886	-3.08e-16
ridge	0.398	0.465	-0.067	0.042 [0.038, 0.046]	yes	2.26e-04	1.82e-04	5.12e-08	0.639	5.28e-05	2.68e-04	2.09e-04	7.18e-08	0.886	-5.03e-17
lasso	0.394	0.453	-0.059	0.031 [0.027, 0.034]	yes	2.29e-04	1.84e-04	5.24e-08	0.573	5.76e-05	2.69e-04	2.09e-04	7.22e-08	0.801	-2.98e-20
elasticnet	0.384	0.436	-0.052	0.014 [0.011, 0.016]	yes	2.32e-04	1.87e-04	5.40e-08	0.535	6.23e-05	2.71e-04	2.11e-04	7.34e-08	0.747	-3.02e-20
huber	0.395	0.477	-0.082	0.054 [0.052, 0.057]	yes	2.24e-04	1.80e-04	5.01e-08	0.779	3.64e-05	2.69e-04	2.08e-04	7.21e-08	1.088	-1.59e-05
poly2	0.419	0.477	-0.058	0.055 [0.050, 0.060]	yes	2.24e-04	1.81e-04	5.00e-08	0.548	5.98e-05	2.63e-04	2.05e-04	6.93e-08	0.751	3.97e-15
poly3	0.419	0.467	-0.048	0.045 [0.040, 0.050]	yes	2.26e-04	1.83e-04	5.10e-08	0.513	6.49e-05	2.63e-04	2.05e-04	6.93e-08	0.798	4.99e-14

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
knn	0.450	0.462	-0.013	0.040 [0.035, 0.045]	yes	2.27e-04	1.83e-04	5.15e-08	0.500	6.61e-05	2.56e-04	1.99e-04	6.56e-08	0.708	-2.26e-06
svr ★	0.448	0.504	-0.056	0.081 [0.077, 0.086]	yes	2.18e-04	1.76e-04	4.75e-08	0.921	4.49e-05	2.56e-04	1.99e-04	6.58e-08	0.936	-1.49e-05
kernelridge	0.449	0.482	-0.033	0.060 [0.054, 0.066]	yes	2.23e-04	1.80e-04	4.96e-08	0.572	6.35e-05	2.56e-04	1.99e-04	6.57e-08	0.707	-1.16e-06
gbm	0.355	0.359	-0.004	-0.063 [-0.072, -0.053]	no	2.48e-04	2.00e-04	6.13e-08	0.396	9.33e-05	2.77e-04	2.16e-04	7.69e-08	0.525	8.92e-08
rf	0.461	0.478	-0.017	0.055 [0.048, 0.062]	yes	2.24e-04	1.81e-04	5.00e-08	0.466	6.73e-05	2.54e-04	1.98e-04	6.43e-08	0.648	-4.55e-09
extratrees	0.488	0.471	0.017	0.048 [0.040, 0.056]	yes	2.25e-04	1.83e-04	5.07e-08	0.451	7.24e-05	2.47e-04	1.92e-04	6.11e-08	0.588	1.83e-19
mlp	0.460	0.477	-0.017	0.055 [0.046, 0.063]	yes	2.24e-04	1.81e-04	5.01e-08	0.478	7.01e-05	2.54e-04	1.97e-04	6.44e-08	0.666	2.74e-07
gp	0.439	0.485	-0.046	0.063 [0.058, 0.068]	yes	2.22e-04	1.79e-04	4.93e-08	0.542	6.20e-05	2.59e-04	2.01e-04	6.68e-08	0.758	-4.47e-06

best-by-test-R2 = **svr** (0.504) · linear 0.465 · persistence 0.423 · H=30s · 2661.7s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=45sSigma=30s · horizon=45s (H=45s) — metrics: LE + AI

LESigma=30s | horizon=45s | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.400	0.468	-0.068	0.046 [0.042, 0.050]	yes	2.26e-04	1.82e-04	5.09e-08	0.630	5.30e-05	2.67e-04	2.08e-04	7.15e-08	0.883	6.84e-16
ridge	0.400	0.468	-0.068	0.046 [0.042, 0.050]	yes	2.26e-04	1.82e-04	5.09e-08	0.630	5.30e-05	2.67e-04	2.08e-04	7.15e-08	0.883	-1.30e-17
lasso	0.397	0.455	-0.059	0.033 [0.030, 0.036]	yes	2.28e-04	1.84e-04	5.21e-08	0.570	5.72e-05	2.68e-04	2.09e-04	7.19e-08	0.799	1.47e-21
elasticnet	0.385	0.437	-0.052	0.015 [0.013, 0.017]	yes	2.32e-04	1.87e-04	5.38e-08	0.537	6.12e-05	2.71e-04	2.10e-04	7.33e-08	0.748	-7.37e-22
huber	0.398	0.481	-0.083	0.058 [0.056, 0.061]	yes	2.23e-04	1.79e-04	4.97e-08	0.770	3.65e-05	2.68e-04	2.08e-04	7.17e-08	1.081	-1.60e-05
poly2	0.422	0.485	-0.063	0.062 [0.058, 0.067]	yes	2.22e-04	1.80e-04	4.93e-08	0.557	5.81e-05	2.62e-04	2.05e-04	6.89e-08	0.759	6.41e-15
poly3	0.421	0.476	-0.055	0.054 [0.050, 0.058]	yes	2.24e-04	1.81e-04	5.01e-08	0.521	6.15e-05	2.63e-04	2.05e-04	6.90e-08	0.826	4.11e-13
knn	0.452	0.461	-0.009	0.039 [0.032, 0.045]	yes	2.27e-04	1.83e-04	5.16e-08	0.506	6.59e-05	2.56e-04	1.99e-04	6.53e-08	0.710	-1.95e-06
svr ★	0.450	0.495	-0.045	0.073 [0.066, 0.079]	yes	2.20e-04	1.77e-04	4.83e-08	0.739	5.01e-05	2.56e-04	1.99e-04	6.55e-08	0.939	-1.38e-05
kernelridge	0.450	0.476	-0.026	0.054 [0.044, 0.062]	yes	2.24e-04	1.80e-04	5.01e-08	1.080	6.31e-05	2.56e-04	1.99e-04	6.55e-08	0.925	1.32e-06
gbm	0.357	0.363	-0.005	-0.060 [-0.068, -0.050]	no	2.47e-04	1.99e-04	6.10e-08	0.394	9.33e-05	2.77e-04	2.16e-04	7.66e-08	0.523	-5.71e-08
rf	0.464	0.480	-0.017	0.058 [0.050, 0.065]	yes	2.23e-04	1.81e-04	4.97e-08	0.478	6.64e-05	2.53e-04	1.97e-04	6.39e-08	0.648	-7.57e-09
extratrees	0.491	0.474	0.016	0.052 [0.044, 0.058]	yes	2.24e-04	1.82e-04	5.03e-08	0.462	7.06e-05	2.46e-04	1.92e-04	6.07e-08	0.588	1.85e-19
mlp	0.463	0.488	-0.025	0.065 [0.058, 0.072]	yes	2.21e-04	1.79e-04	4.90e-08	0.506	6.31e-05	2.53e-04	1.97e-04	6.40e-08	0.700	-3.78e-06
gp	0.441	0.472	-0.032	0.050 [0.041, 0.057]	yes	2.25e-04	1.81e-04	5.05e-08	0.490	6.61e-05	2.58e-04	2.01e-04	6.67e-08	0.685	-9.91e-07

best-by-test-R2 = **svr** (0.495) · linear 0.468 · persistence 0.422 · H=45s · 2681.8s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AISigma=30s | horizon=45s | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.398	0.465	-0.067	0.042 [0.038, 0.047]	yes	2.26e-04	1.82e-04	5.12e-08	0.639	5.28e-05	2.68e-04	2.09e-04	7.18e-08	0.887	-7.61e-16
ridge	0.398	0.465	-0.067	0.042 [0.038, 0.047]	yes	2.26e-04	1.82e-04	5.12e-08	0.639	5.28e-05	2.68e-04	2.09e-04	7.18e-08	0.886	-1.36e-17
lasso	0.394	0.453	-0.059	0.031 [0.027, 0.034]	yes	2.29e-04	1.84e-04	5.24e-08	0.574	5.76e-05	2.69e-04	2.09e-04	7.22e-08	0.801	7.00e-21
elasticnet	0.384	0.436	-0.052	0.014 [0.011, 0.016]	yes	2.32e-04	1.87e-04	5.40e-08	0.535	6.23e-05	2.71e-04	2.11e-04	7.34e-08	0.747	-1.03e-20
huber	0.395	0.477	-0.082	0.054 [0.052, 0.057]	yes	2.24e-04	1.80e-04	5.01e-08	0.780	3.64e-05	2.69e-04	2.08e-04	7.21e-08	1.088	-1.59e-05
poly2	0.419	0.477	-0.058	0.055 [0.050, 0.060]	yes	2.24e-04	1.81e-04	5.00e-08	0.548	5.98e-05	2.63e-04	2.05e-04	6.92e-08	0.751	2.56e-15
poly3	0.419	0.467	-0.048	0.045 [0.040, 0.050]	yes	2.26e-04	1.83e-04	5.10e-08	0.514	6.48e-05	2.63e-04	2.05e-04	6.93e-08	0.798	1.11e-13
knn	0.450	0.463	-0.013	0.041 [0.035, 0.046]	yes	2.27e-04	1.83e-04	5.14e-08	0.501	6.55e-05	2.56e-04	1.99e-04	6.56e-08	0.713	-2.64e-06
svr ★	0.448	0.497	-0.049	0.075 [0.070, 0.079]	yes	2.19e-04	1.77e-04	4.81e-08	0.706	4.93e-05	2.56e-04	1.99e-04	6.57e-08	0.920	-1.35e-05
kernelridge	0.449	0.478	-0.029	0.056 [0.049, 0.062]	yes	2.23e-04	1.80e-04	4.99e-08	0.740	6.30e-05	2.56e-04	1.99e-04	6.56e-08	0.746	1.54e-06
gbm	0.355	0.360	-0.004	-0.063 [-0.072, -0.053]	no	2.48e-04	2.00e-04	6.13e-08	0.397	9.30e-05	2.77e-04	2.16e-04	7.69e-08	0.526	-5.19e-08
rf	0.461	0.477	-0.016	0.055 [0.048, 0.062]	yes	2.24e-04	1.81e-04	5.01e-08	0.466	6.75e-05	2.54e-04	1.98e-04	6.43e-08	0.648	-6.11e-09
extratrees	0.487	0.470	0.017	0.048 [0.040, 0.056]	yes	2.25e-04	1.83e-04	5.07e-08	0.449	7.28e-05	2.47e-04	1.92e-04	6.11e-08	0.587	1.88e-19
mlp	0.460	0.482	-0.022	0.060 [0.052, 0.067]	yes	2.23e-04	1.80e-04	4.95e-08	0.476	6.80e-05	2.54e-04	1.98e-04	6.44e-08	0.656	1.33e-06
gp	0.439	0.470	-0.031	0.048 [0.039, 0.055]	yes	2.25e-04	1.81e-04	5.07e-08	0.502	6.49e-05	2.59e-04	2.01e-04	6.68e-08	0.690	-1.36e-06

best-by-test-R2 = **svr** (0.497) · linear 0.465 · persistence 0.422 · H=45s · 2656.5s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=60sSigma=30s · horizon=60s (H=60s) — metrics: LE + AI

LESigma=30s | horizon=60s | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.454	0.533	-0.079	0.048 [0.044, 0.053]	yes	2.66e-04	2.10e-04	7.06e-08	0.410	8.38e-05	3.14e-04	2.42e-04	9.88e-08	0.524	2.09e-16
ridge	0.454	0.533	-0.079	0.048 [0.044, 0.053]	yes	2.66e-04	2.10e-04	7.06e-08	0.410	8.38e-05	3.14e-04	2.42e-04	9.88e-08	0.524	-1.60e-17
lasso	0.452	0.522	-0.070	0.037 [0.033, 0.041]	yes	2.69e-04	2.13e-04	7.23e-08	0.401	8.79e-05	3.15e-04	2.42e-04	9.93e-08	0.501	7.83e-21
elasticnet	0.444	0.507	-0.063	0.022 [0.019, 0.025]	yes	2.73e-04	2.16e-04	7.46e-08	0.397	9.21e-05	3.17e-04	2.43e-04	1.01e-07	0.487	-3.70e-20
huber	0.452	0.551	-0.099	0.066 [0.063, 0.068]	yes	2.61e-04	2.06e-04	6.79e-08	0.419	6.66e-05	3.15e-04	2.41e-04	9.93e-08	0.555	-1.97e-05
poly2	0.479	0.547	-0.068	0.062 [0.055, 0.069]	yes	2.62e-04	2.07e-04	6.85e-08	0.399	9.19e-05	3.07e-04	2.36e-04	9.44e-08	0.479	-3.36e-15
poly3	0.479	0.535	-0.057	0.050 [0.044, 0.057]	yes	2.65e-04	2.10e-04	7.02e-08	0.389	9.73e-05	3.07e-04	2.36e-04	9.44e-08	0.486	2.91e-13
knn	0.516	0.518	-0.002	0.033 [0.026, 0.041]	yes	2.70e-04	2.14e-04	7.28e-08	0.373	1.05e-04	2.96e-04	2.26e-04	8.76e-08	0.421	-2.40e-06
svr ★	0.514	0.558	-0.044	0.073 [0.067, 0.080]	yes	2.58e-04	2.04e-04	6.68e-08	0.365	8.90e-05	2.97e-04	2.26e-04	8.80e-08	0.428	-1.33e-05
kernelridge	0.514	0.534	-0.020	0.049 [0.039, 0.058]	yes	2.65e-04	2.10e-04	7.04e-08	0.535	1.02e-04	2.97e-04	2.26e-04	8.80e-08	0.522	3.17e-07
gbm	0.406	0.400	0.006	-0.085 [-0.093, -0.074]	no	3.01e-04	2.39e-04	9.06e-08	0.399	1.38e-04	3.28e-04	2.50e-04	1.08e-07	0.402	2.04e-07
rf	0.535	0.537	-0.002	0.052 [0.042, 0.062]	yes	2.65e-04	2.11e-04	7.00e-08	0.371	1.03e-04	2.90e-04	2.22e-04	8.43e-08	0.400	2.29e-08
extratrees	0.553	0.531	0.022	0.046 [0.036, 0.055]	yes	2.66e-04	2.12e-04	7.09e-08	0.368	1.08e-04	2.84e-04	2.17e-04	8.09e-08	0.372	5.98e-20
mlp	0.527	0.542	-0.015	0.058 [0.047, 0.067]	yes	2.63e-04	2.10e-04	6.92e-08	0.372	1.03e-04	2.93e-04	2.23e-04	8.56e-08	0.417	-3.04e-06
gp	0.502	0.525	-0.023	0.040 [0.031, 0.049]	yes	2.68e-04	2.11e-04	7.18e-08	0.370	1.07e-04	3.00e-04	2.29e-04	9.01e-08	0.418	-4.53e-07

best-by-test-R2 = **svr** (0.558) · linear 0.533 · persistence 0.485 · H=60s · 2688.0s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AISigma=30s | horizon=60s | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.451	0.529	-0.078	0.044 [0.039, 0.049]	yes	2.67e-04	2.11e-04	7.12e-08	0.415	8.38e-05	3.15e-04	2.42e-04	9.94e-08	0.527	-4.34e-16

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
ridge	0.451	0.529	-0.078	0.044 [0.039, 0.049]	yes	2.67e-04	2.11e-04	7.12e-08	0.415	8.38e-05	3.15e-04	2.42e-04	9.94e-08	0.527	-1.64e-17
lasso	0.449	0.518	-0.070	0.034 [0.029, 0.038]	yes	2.70e-04	2.14e-04	7.28e-08	0.404	8.86e-05	3.16e-04	2.42e-04	9.98e-08	0.503	3.13e-20
elasticnet	0.442	0.504	-0.062	0.019 [0.015, 0.023]	yes	2.74e-04	2.17e-04	7.50e-08	0.399	9.35e-05	3.18e-04	2.44e-04	1.01e-07	0.487	-3.56e-20
huber	0.448	0.546	-0.097	0.061 [0.058, 0.063]	yes	2.62e-04	2.07e-04	6.87e-08	0.425	6.69e-05	3.16e-04	2.41e-04	9.99e-08	0.560	-1.96e-05
poly2	0.475	0.537	-0.062	0.052 [0.045, 0.060]	yes	2.64e-04	2.09e-04	7.00e-08	0.400	9.43e-05	3.08e-04	2.36e-04	9.50e-08	0.478	-1.21e-15
poly3	0.477	0.524	-0.047	0.039 [0.031, 0.047]	yes	2.68e-04	2.13e-04	7.19e-08	0.392	1.02e-04	3.08e-04	2.36e-04	9.48e-08	0.477	-1.14e-13
knn	0.515	0.521	-0.006	0.036 [0.029, 0.043]	yes	2.69e-04	2.14e-04	7.24e-08	0.372	1.04e-04	2.96e-04	2.26e-04	8.79e-08	0.424	-3.49e-06
svr ★	0.512	0.560	-0.048	0.075 [0.070, 0.081]	yes	2.58e-04	2.04e-04	6.65e-08	0.360	8.92e-05	2.97e-04	2.26e-04	8.84e-08	0.430	-1.32e-05
kernelridge	0.512	0.535	-0.023	0.050 [0.042, 0.057]	yes	2.65e-04	2.10e-04	7.03e-08	0.563	1.03e-04	2.97e-04	2.27e-04	8.84e-08	0.504	3.29e-07
gbm	0.404	0.401	0.003	-0.084 [-0.093, -0.074]	no	3.01e-04	2.39e-04	9.06e-08	0.399	1.37e-04	3.29e-04	2.51e-04	1.08e-07	0.404	2.03e-07
rf	0.523	0.533	-0.010	0.048 [0.038, 0.059]	yes	2.66e-04	2.12e-04	7.05e-08	0.366	1.06e-04	2.94e-04	2.25e-04	8.63e-08	0.406	2.73e-08
extratrees	0.551	0.526	0.025	0.041 [0.030, 0.052]	yes	2.68e-04	2.14e-04	7.17e-08	0.367	1.11e-04	2.85e-04	2.18e-04	8.14e-08	0.373	5.86e-20
mlp	0.525	0.532	-0.007	0.047 [0.035, 0.058]	yes	2.66e-04	2.12e-04	7.08e-08	0.369	1.11e-04	2.93e-04	2.24e-04	8.60e-08	0.413	-1.57e-06
gp	0.501	0.522	-0.021	0.037 [0.029, 0.045]	yes	2.69e-04	2.12e-04	7.22e-08	0.367	1.07e-04	3.01e-04	2.29e-04	9.03e-08	0.416	-2.46e-07

best-by-test-R2 = **svr** (0.560) · linear 0.529 · persistence 0.485 · H=60s · 2684.7s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=2mSigma=30s · horizon=2m (H=120s) — metrics: LE + AI

LESigma=30s | horizon=2m | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.505	0.587	-0.082	0.051 [0.046, 0.057]	yes	3.25e-04	2.57e-04	1.06e-07	0.332	1.28e-04	3.83e-04	2.92e-04	1.47e-07	0.381	9.53e-16
ridge	0.505	0.587	-0.082	0.051 [0.046, 0.057]	yes	3.25e-04	2.58e-04	1.06e-07	0.332	1.28e-04	3.83e-04	2.92e-04	1.47e-07	0.381	1.78e-17
lasso	0.504	0.578	-0.074	0.042 [0.037, 0.047]	yes	3.29e-04	2.60e-04	1.08e-07	0.333	1.32e-04	3.84e-04	2.92e-04	1.47e-07	0.374	1.02e-19
elasticnet	0.499	0.566	-0.067	0.031 [0.027, 0.035]	yes	3.33e-04	2.64e-04	1.11e-07	0.335	1.36e-04	3.85e-04	2.93e-04	1.48e-07	0.369	7.93e-20
huber	0.501	0.610	-0.109	0.075 [0.072, 0.078]	yes	3.16e-04	2.48e-04	9.98e-08	0.325	1.04e-04	3.84e-04	2.90e-04	1.48e-07	0.393	-2.94e-05
poly2	0.534	0.600	-0.065	0.064 [0.055, 0.073]	yes	3.20e-04	2.54e-04	1.03e-07	0.334	1.40e-04	3.72e-04	2.82e-04	1.38e-07	0.357	9.22e-15
poly3	0.534	0.584	-0.050	0.049 [0.040, 0.057]	yes	3.26e-04	2.59e-04	1.06e-07	0.333	1.48e-04	3.72e-04	2.83e-04	1.38e-07	0.354	4.80e-13
knn	0.580	0.570	0.010	0.034 [0.026, 0.043]	yes	3.32e-04	2.63e-04	1.10e-07	0.321	1.57e-04	3.53e-04	2.65e-04	1.25e-07	0.301	-3.56e-06
svr ★	0.575	0.629	-0.055	0.094 [0.087, 0.101]	yes	3.08e-04	2.43e-04	9.49e-08	0.309	1.20e-04	3.55e-04	2.65e-04	1.26e-07	0.321	-2.75e-05
kernelridge	0.575	0.593	-0.019	0.058 [0.048, 0.067]	yes	3.23e-04	2.56e-04	1.04e-07	0.395	1.50e-04	3.55e-04	2.67e-04	1.26e-07	0.367	-1.87e-06
gbm	0.455	0.436	0.019	-0.100 [-0.108, -0.089]	no	3.80e-04	3.06e-04	1.44e-07	0.387	2.02e-04	4.02e-04	3.04e-04	1.62e-07	0.336	-1.22e-07
rf	0.619	0.569	0.050	0.033 [0.018, 0.047]	yes	3.32e-04	2.66e-04	1.10e-07	0.328	1.64e-04	3.36e-04	2.53e-04	1.13e-07	0.277	6.40e-08
extratrees	0.615	0.575	0.040	0.039 [0.027, 0.050]	yes	3.30e-04	2.65e-04	1.09e-07	0.325	1.62e-04	3.38e-04	2.54e-04	1.14e-07	0.273	3.01e-21
mlp	0.592	0.590	0.002	0.054 [0.039, 0.066]	yes	3.24e-04	2.59e-04	1.05e-07	0.316	1.56e-04	3.48e-04	2.62e-04	1.21e-07	0.293	1.10e-06
gp	0.560	0.584	-0.024	0.048 [0.035, 0.059]	yes	3.27e-04	2.57e-04	1.07e-07	0.320	1.49e-04	3.61e-04	2.71e-04	1.31e-07	0.316	-4.69e-06

best-by-test-R2 = **svr** (0.629) · linear 0.587 · persistence 0.536 · H=120s · 2735.9s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AISigma=30s | horizon=2m | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.502	0.583	-0.081	0.047 [0.041, 0.053]	yes	3.27e-04	2.59e-04	1.07e-07	0.336	1.28e-04	3.84e-04	2.93e-04	1.48e-07	0.384	2.66e-16
ridge	0.502	0.582	-0.081	0.047 [0.041, 0.053]	yes	3.27e-04	2.59e-04	1.07e-07	0.336	1.28e-04	3.84e-04	2.93e-04	1.48e-07	0.384	1.79e-17
lasso	0.500	0.574	-0.073	0.038 [0.032, 0.043]	yes	3.30e-04	2.62e-04	1.09e-07	0.336	1.33e-04	3.85e-04	2.93e-04	1.48e-07	0.376	8.99e-20
elasticnet	0.496	0.563	-0.066	0.027 [0.022, 0.032]	yes	3.35e-04	2.65e-04	1.12e-07	0.337	1.38e-04	3.86e-04	2.94e-04	1.49e-07	0.371	2.41e-20
huber	0.498	0.605	-0.107	0.069 [0.066, 0.072]	yes	3.18e-04	2.50e-04	1.01e-07	0.330	1.04e-04	3.86e-04	2.91e-04	1.49e-07	0.397	-2.94e-05
poly2	0.531	0.588	-0.057	0.052 [0.042, 0.062]	yes	3.25e-04	2.58e-04	1.06e-07	0.336	1.44e-04	3.73e-04	2.83e-04	1.39e-07	0.358	-4.27e-15
poly3	0.533	0.571	-0.039	0.036 [0.025, 0.046]	yes	3.31e-04	2.64e-04	1.10e-07	0.339	1.55e-04	3.72e-04	2.83e-04	1.39e-07	0.374	-4.99e-13
knn	0.578	0.573	0.005	0.037 [0.028, 0.045]	yes	3.31e-04	2.63e-04	1.09e-07	0.321	1.57e-04	3.54e-04	2.66e-04	1.25e-07	0.303	-5.13e-06
svr ★	0.571	0.624	-0.053	0.089 [0.082, 0.096]	yes	3.10e-04	2.44e-04	9.62e-08	0.302	1.24e-04	3.56e-04	2.66e-04	1.27e-07	0.322	-2.80e-05
kernelridge	0.572	0.587	-0.015	0.051 [0.041, 0.060]	yes	3.25e-04	2.57e-04	1.06e-07	0.398	1.53e-04	3.56e-04	2.67e-04	1.27e-07	0.386	-1.80e-06
gbm	0.451	0.438	0.013	-0.097 [-0.106, -0.088]	no	3.79e-04	3.05e-04	1.44e-07	0.387	2.00e-04	4.03e-04	3.05e-04	1.63e-07	0.339	-1.08e-07
rf	0.616	0.552	0.063	0.017 [0.000, 0.033]	yes	3.39e-04	2.72e-04	1.15e-07	0.334	1.74e-04	3.37e-04	2.54e-04	1.14e-07	0.279	7.08e-08
extratrees	0.613	0.568	0.045	0.032 [0.017, 0.046]	yes	3.33e-04	2.67e-04	1.11e-07	0.328	1.68e-04	3.39e-04	2.55e-04	1.15e-07	0.274	2.23e-20
mlp	0.588	0.575	0.013	0.039 [0.024, 0.055]	yes	3.30e-04	2.65e-04	1.09e-07	0.322	1.70e-04	3.49e-04	2.63e-04	1.22e-07	0.293	2.73e-06
gp	0.559	0.573	-0.014	0.038 [0.026, 0.049]	yes	3.31e-04	2.59e-04	1.09e-07	0.326	1.51e-04	3.61e-04	2.71e-04	1.31e-07	0.317	-4.47e-06

best-by-test-R2 = **svr** (0.624) · linear 0.583 · persistence 0.536 · H=120s · 2732.6s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=5mSigma=30s · horizon=5m (H=300s) — metrics: LE + AI

LESigma=30s | horizon=5m | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.545	0.627	-0.082	0.055 [0.049, 0.061]	yes	4.61e-04	3.68e-04	2.13e-07	0.256	2.16e-04	5.38e-04	4.08e-04	2.90e-07	0.285	4.82e-16
ridge	0.545	0.627	-0.082	0.055 [0.049, 0.061]	yes	4.61e-04	3.68e-04	2.13e-07	0.256	2.16e-04	5.38e-04	4.08e-04	2.90e-07	0.285	7.10e-17
lasso	0.544	0.620	-0.076	0.049 [0.043, 0.054]	yes	4.65e-04	3.71e-04	2.17e-07	0.258	2.20e-04	5.39e-04	4.09e-04	2.90e-07	0.283	-2.99e-20
elasticnet	0.542	0.613	-0.071	0.041 [0.035, 0.046]	yes	4.70e-04	3.75e-04	2.21e-07	0.260	2.24e-04	5.40e-04	4.09e-04	2.91e-07	0.281	-4.49e-20
huber	0.539	0.657	-0.118	0.086 [0.082, 0.089]	yes	4.42e-04	3.49e-04	1.95e-07	0.245	1.75e-04	5.41e-04	4.04e-04	2.93e-07	0.294	-5.14e-05
poly2	0.578	0.638	-0.059	0.066 [0.056, 0.076]	yes	4.55e-04	3.65e-04	2.07e-07	0.261	2.35e-04	5.18e-04	3.90e-04	2.68e-07	0.268	-3.64e-15
poly3	0.578	0.619	-0.041	0.047 [0.037, 0.056]	yes	4.66e-04	3.74e-04	2.17e-07	0.267	2.49e-04	5.18e-04	3.91e-04	2.69e-07	0.278	8.78e-13
knn	0.633	0.604	0.029	0.032 [0.022, 0.042]	yes	4.75e-04	3.83e-04	2.26e-07	0.256	2.63e-04	4.83e-04	3.59e-04	2.33e-07	0.216	-3.87e-06
svr ★	0.626	0.682	-0.056	0.110 [0.102, 0.118]	yes	4.26e-04	3.38e-04	1.81e-07	0.226	1.98e-04	4.87e-04	3.59e-04	2.38e-07	0.229	-4.61e-05
kernelridge	0.626	0.622	0.005	0.050 [0.039, 0.061]	yes	4.64e-04	3.73e-04	2.16e-07	0.405	2.59e-04	4.88e-04	3.63e-04	2.38e-07	0.308	1.81e-06
gbm	0.532	0.507	0.025	-0.064 [-0.072, -0.055]	no	5.30e-04	4.34e-04	2.81e-07	0.313	3.12e-04	5.45e-04	4.10e-04	2.97e-07	0.252	1.95e-07
rf	0.671	0.590	0.081	0.018 [-0.000, 0.034]	no	4.83e-04	3.95e-04	2.34e-07	0.268	2.82e-04	4.57e-04	3.40e-04	2.09e-07	0.197	7.33e-08
extratrees	0.667	0.603	0.064	0.032 [0.017, 0.044]	yes	4.75e-04	3.88e-04	2.26e-07	0.263	2.73e-04	4.60e-04	3.41e-04	2.11e-07	0.194	3.33e-20
mlp	0.644	0.624	0.020	0.053 [0.035, 0.067]	yes	4.63e-04	3.75e-04	2.14e-07	0.249	2.59e-04	4.76e-04	3.53e-04	2.26e-07	0.212	-2.50e-06
gp	0.615	0.629	-0.014	0.057 [0.045, 0.067]	yes	4.60e-04	3.70e-04	2.12e-07	0.249	2.54e-04	4.95e-04	3.67e-04	2.45e-07	0.228	-7.84e-06

best-by-test-R2 = **svr** (0.682) · linear 0.627 · persistence 0.572 · H=300s · 2655.5s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI Sigma=30s | horizon=5m | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.541	0.622	-0.080	0.050 [0.043, 0.057]	yes	4.64e-04	3.71e-04	2.16e-07	0.260	2.17e-04	5.40e-04	4.10e-04	2.92e-07	0.288	-2.76e-16
ridge	0.541	0.622	-0.080	0.050 [0.043, 0.057]	yes	4.64e-04	3.71e-04	2.16e-07	0.260	2.17e-04	5.40e-04	4.10e-04	2.92e-07	0.288	7.31e-17
lasso	0.541	0.615	-0.075	0.044 [0.037, 0.050]	yes	4.68e-04	3.74e-04	2.19e-07	0.261	2.22e-04	5.41e-04	4.10e-04	2.92e-07	0.285	-2.42e-20
elasticnet	0.539	0.608	-0.069	0.037 [0.030, 0.043]	yes	4.73e-04	3.77e-04	2.23e-07	0.263	2.27e-04	5.42e-04	4.11e-04	2.93e-07	0.283	2.71e-20
huber	0.536	0.651	-0.115	0.079 [0.075, 0.083]	yes	4.46e-04	3.52e-04	1.99e-07	0.249	1.75e-04	5.43e-04	4.06e-04	2.95e-07	0.298	-5.16e-05
poly2	0.574	0.624	-0.049	0.052 [0.041, 0.063]	yes	4.63e-04	3.71e-04	2.14e-07	0.266	2.41e-04	5.20e-04	3.92e-04	2.71e-07	0.269	-1.46e-14
poly3	0.577	0.604	-0.027	0.032 [0.019, 0.044]	yes	4.75e-04	3.83e-04	2.26e-07	0.274	2.60e-04	5.19e-04	3.91e-04	2.69e-07	0.308	7.34e-13
knn	0.631	0.606	0.025	0.034 [0.023, 0.044]	yes	4.74e-04	3.83e-04	2.25e-07	0.257	2.64e-04	4.85e-04	3.60e-04	2.35e-07	0.218	-6.20e-06
svr ★	0.624	0.671	-0.048	0.100 [0.091, 0.108]	yes	4.33e-04	3.44e-04	1.87e-07	0.230	2.06e-04	4.89e-04	3.60e-04	2.39e-07	0.231	-4.67e-05
kernelridge	0.624	0.613	0.011	0.041 [0.030, 0.052]	yes	4.70e-04	3.78e-04	2.21e-07	0.314	2.65e-04	4.89e-04	3.64e-04	2.39e-07	0.267	2.02e-06
gbm	0.527	0.502	0.025	-0.069 [-0.077, -0.061]	no	5.33e-04	4.36e-04	2.84e-07	0.315	3.12e-04	5.48e-04	4.13e-04	3.01e-07	0.257	1.96e-07
rf	0.668	0.572	0.096	0.001 [-0.021, 0.020]	no	4.94e-04	4.04e-04	2.44e-07	0.275	2.94e-04	4.59e-04	3.41e-04	2.11e-07	0.199	9.57e-08
extratrees	0.664	0.591	0.073	0.020 [0.002, 0.036]	yes	4.83e-04	3.95e-04	2.33e-07	0.269	2.85e-04	4.62e-04	3.43e-04	2.13e-07	0.196	3.74e-20
mlp	0.641	0.591	0.049	0.020 [-0.001, 0.039]	no	4.83e-04	3.94e-04	2.33e-07	0.266	2.89e-04	4.78e-04	3.55e-04	2.28e-07	0.214	1.10e-06
gp	0.616	0.619	-0.003	0.048 [0.035, 0.058]	yes	4.66e-04	3.74e-04	2.17e-07	0.252	2.58e-04	4.94e-04	3.67e-04	2.44e-07	0.229	-6.72e-06

best-by-test-R2 = **svr** (0.671) · linear 0.622 · persistence 0.572 · H=300s · 2636.2s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=10m Sigma=30s · horizon=10m (H=600s) — metrics: LE + AI

LE Sigma=30s | horizon=10m | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.544	0.630	-0.087	0.058 [0.051, 0.064]	yes	6.34e-04	5.09e-04	4.02e-07	0.227	3.20e-04	7.36e-04	5.56e-04	5.42e-07	0.252	1.31e-16
ridge	0.544	0.630	-0.087	0.058 [0.051, 0.064]	yes	6.34e-04	5.09e-04	4.02e-07	0.227	3.20e-04	7.36e-04	5.56e-04	5.42e-07	0.252	1.19e-16
lasso	0.543	0.625	-0.082	0.053 [0.046, 0.059]	yes	6.39e-04	5.13e-04	4.08e-07	0.228	3.25e-04	7.37e-04	5.56e-04	5.43e-07	0.250	4.29e-19
elasticnet	0.542	0.620	-0.078	0.047 [0.041, 0.053]	yes	6.43e-04	5.17e-04	4.14e-07	0.230	3.29e-04	7.38e-04	5.57e-04	5.44e-07	0.249	2.33e-19
huber	0.537	0.666	-0.130	0.094 [0.090, 0.098]	yes	6.03e-04	4.77e-04	3.63e-07	0.213	2.58e-04	7.42e-04	5.49e-04	5.51e-07	0.260	-7.78e-05
poly2	0.579	0.638	-0.060	0.066 [0.055, 0.077]	yes	6.27e-04	5.07e-04	3.94e-07	0.232	3.48e-04	7.07e-04	5.30e-04	5.00e-07	0.237	1.77e-14
poly3	0.578	0.617	-0.039	0.044 [0.033, 0.054]	yes	6.46e-04	5.22e-04	4.17e-07	0.240	3.68e-04	7.08e-04	5.31e-04	5.01e-07	0.257	3.13e-13
knn	0.638	0.600	0.038	0.027 [0.016, 0.037]	yes	6.60e-04	5.37e-04	4.36e-07	0.232	3.89e-04	6.56e-04	4.82e-04	4.30e-07	0.188	-4.00e-06
svr ★	0.629	0.682	-0.053	0.110 [0.100, 0.119]	yes	5.88e-04	4.68e-04	3.46e-07	0.200	2.88e-04	6.64e-04	4.81e-04	4.40e-07	0.201	-7.27e-05
kernelridge	0.629	0.623	0.006	0.050 [0.037, 0.062]	yes	6.41e-04	5.20e-04	4.11e-07	0.400	3.74e-04	6.64e-04	4.88e-04	4.41e-07	0.263	7.79e-07
gbm	0.568	0.543	0.025	-0.030 [-0.037, -0.020]	no	7.05e-04	5.84e-04	4.97e-07	0.270	4.34e-04	7.17e-04	5.35e-04	5.14e-07	0.213	-1.40e-07
rf	0.690	0.564	0.126	-0.008 [-0.032, 0.012]	no	6.89e-04	5.67e-04	4.74e-07	0.251	4.27e-04	6.06e-04	4.44e-04	3.68e-07	0.161	3.28e-07
extratrees	0.673	0.594	0.079	0.022 [0.004, 0.035]	yes	6.65e-04	5.48e-04	4.42e-07	0.241	4.06e-04	6.23e-04	4.57e-04	3.89e-07	0.168	1.52e-19
mlp	0.647	0.616	0.031	0.044 [0.024, 0.059]	yes	6.46e-04	5.31e-04	4.18e-07	0.227	3.94e-04	6.47e-04	4.75e-04	4.19e-07	0.185	-2.79e-06
gp	0.614	0.582	0.032	0.009 [-0.008, 0.025]	no	6.75e-04	5.45e-04	4.56e-07	0.239	4.00e-04	6.77e-04	4.96e-04	4.59e-07	0.197	1.30e-06

best-by-test-R2 = **svr** (0.682) · linear 0.630 · persistence 0.573 · H=600s · 2654.0s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI Sigma=30s | horizon=10m | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.540	0.625	-0.086	0.053 [0.044, 0.060]	yes	6.39e-04	5.13e-04	4.08e-07	0.230	3.23e-04	7.39e-04	5.59e-04	5.47e-07	0.255	5.78e-16
ridge	0.540	0.625	-0.086	0.053 [0.044, 0.060]	yes	6.39e-04	5.13e-04	4.08e-07	0.230	3.23e-04	7.39e-04	5.59e-04	5.47e-07	0.255	1.23e-16
lasso	0.539	0.621	-0.081	0.048 [0.040, 0.055]	yes	6.43e-04	5.16e-04	4.13e-07	0.231	3.28e-04	7.40e-04	5.59e-04	5.47e-07	0.253	1.57e-19
elasticnet	0.538	0.615	-0.077	0.043 [0.035, 0.050]	yes	6.47e-04	5.20e-04	4.19e-07	0.233	3.33e-04	7.41e-04	5.60e-04	5.48e-07	0.252	3.80e-19
huber	0.533	0.660	-0.127	0.087 [0.083, 0.091]	yes	6.08e-04	4.81e-04	3.70e-07	0.217	2.60e-04	7.45e-04	5.52e-04	5.55e-07	0.265	-7.83e-05
poly2	0.574	0.624	-0.049	0.051 [-0.039, 0.063]	yes	6.40e-04	5.17e-04	4.10e-07	0.238	3.56e-04	7.11e-04	5.33e-04	5.06e-07	0.239	3.18e-14
poly3	0.577	0.601	-0.024	0.028 [0.014, 0.041]	yes	6.59e-04	5.35e-04	4.34e-07	0.259	3.84e-04	7.09e-04	5.30e-04	5.02e-07	0.275	1.06e-12
knn	0.635	0.600	0.035	0.027 [0.013, 0.039]	yes	6.60e-04	5.39e-04	4.36e-07	0.233	3.93e-04	6.58e-04	4.84e-04	4.33e-07	0.190	-7.49e-06
svr ★	0.627	0.674	-0.046	0.101 [0.090, 0.111]	yes	5.96e-04	4.76e-04	3.55e-07	0.204	3.02e-04	6.65e-04	4.82e-04	4.43e-07	0.203	-7.19e-05
kernelridge	0.627	0.616	0.012	0.043 [0.031, 0.054]	yes	6.47e-04	5.25e-04	4.18e-07	0.427	3.84e-04	6.65e-04	4.90e-04	4.43e-07	0.246	1.32e-06
gbm	0.563	0.537	0.026	-0.035 [-0.044, -0.025]	no	7.10e-04	5.88e-04	5.04e-07	0.273	4.37e-04	7.21e-04	5.39e-04	5.19e-07	0.216	-1.54e-07
rf	0.674	0.556	0.118	-0.017 [-0.041, 0.006]	no	6.95e-04	5.73e-04	4.83e-07	0.254	4.40e-04	6.23e-04	4.57e-04	3.88e-07	0.171	1.75e-07
extratrees	0.670	0.579	0.091	0.007 [-0.014, 0.026]	no	6.77e-04	5.60e-04	4.58e-07	0.248	4.25e-04	6.26e-04	4.59e-04	3.92e-07	0.169	5.76e-21
mlp	0.645	0.586	0.059	0.013 [-0.010, 0.035]	no	6.72e-04	5.52e-04	4.51e-07	0.240	4.23e-04	6.50e-04	4.77e-04	4.22e-07	0.187	-4.06e-06
gp	0.613	0.582	0.032	0.009 [-0.004, 0.021]	no	6.75e-04	5.49e-04	4.55e-07	0.241	4.11e-04	6.78e-04	4.97e-04	4.59e-07	0.198	2.59e-06

best-by-test-R2 = **svr** (0.674) · linear 0.625 · persistence 0.573 · H=600s · 2588.6s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=15m Sigma=30s · horizon=15m (H=900s) — metrics: LE + AI

LE Sigma=30s | horizon=15m | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.533	0.625	-0.092	0.060 [0.053, 0.067]	yes	7.75e-04	6.25e-04	6.01e-07	0.219	4.06e-04	8.98e-04	6.75e-04	8.06e-07	0.242	9.28e-16
ridge	0.533	0.625	-0.092	0.060 [0.053, 0.067]	yes	7.75e-04	6.25e-04	6.01e-07	0.219	4.06e-04	8.98e-04	6.75e-04	8.06e-07	0.242	-1.70e-16
lasso	0.533	0.621	-0.088	0.056 [0.049, 0.063]	yes	7.80e-04	6.29e-04	6.08e-07	0.220	4.10e-04	8.98e-04	6.76e-04	8.06e-07	0.241	-3.79e-20
elasticnet	0.532	0.616	-0.084	0.051 [0.044, 0.058]	yes	7.84e-04	6.32e-04	6.15e-07	0.222	4.14e-04	8.99e-04	6.76e-04	8.08e-07	0.240	1.40e-19
huber	0.525	0.665	-0.140	0.100 [0.096, 0.104]	yes	7.32e-04	5.81e-04	5.36e-07	0.203	3.25e-04	9.05e-04	6.66e-04	8.19e-07	0.252	-1.00e-04
poly2	0.568	0.630	-0.062	0.065 [0.053, 0.077]	yes	7.70e-04	6.24e-04	5.93e-07	0.225	4.38e-04	8.63e-04	6.44e-04	7.45e-07	0.228	6.62e-15
poly3	0.568	0.607	-0.039	0.042 [0.031, 0.052]	yes	7.94e-04	6.44e-04	6.30e-07	0.233	4.64e-04	8.63e-04	6.44e-04	7.45e-07	0.248	-7.27e-13
knn	0.628	0.591	0.037	0.026 [0.015, 0.037]	yes	8.09e-04	6.62e-04	6.55e-07	0.226	4.89e-04	8.01e-04	5.85e-04	6.41e-07	0.182	-4.07e-06
svr ★	0.619	0.678	-0.059	0.113 [0.103, 0.123]	yes	7.18e-04	5.76e-04	5.16e-07	0.194	3.67e-04	8.11e-04	5.82e-04	6.57e-07	0.195	-9.32e-05
kernelridge	0.619	0.610	0.010	0.045 [0.031, 0.056]	yes	7.91e-04	6.47e-04	6.26e-07	0.272	4.78e-04	8.10e-04	5.92e-04	6.57e-07	0.222	3.78e-06

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
gbm	0.571	0.546	0.026	-0.019 [-0.027, -0.009]	no	8.53e-04	7.09e-04	7.28e-07	0.257	5.35e-04	8.60e-04	6.38e-04	7.39e-07	0.201	3.43e-07
rf	0.669	0.555	0.114	-0.009 [-0.032, 0.011]	no	8.44e-04	7.01e-04	7.12e-07	0.245	5.40e-04	7.55e-04	5.50e-04	5.70e-07	0.163	2.78e-08
extratrees	0.666	0.576	0.090	0.011 [-0.008, 0.026]	no	8.24e-04	6.83e-04	6.79e-07	0.238	5.17e-04	7.59e-04	5.53e-04	5.77e-07	0.161	-4.00e-20
mlp	0.638	0.602	0.036	0.037 [0.016, 0.053]	yes	7.99e-04	6.59e-04	6.38e-07	0.223	4.98e-04	7.90e-04	5.76e-04	6.24e-07	0.179	-7.68e-06
gp	0.604	0.572	0.032	0.008 [-0.011, 0.023]	no	8.28e-04	6.72e-04	6.85e-07	0.233	4.98e-04	8.26e-04	6.01e-04	6.82e-07	0.191	-4.37e-06

best-by-test-R2 = **svr** (0.678) · linear 0.625 · persistence 0.565 · H=900s · 2605.4s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI

Sigma=30s | horizon=15m | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.529	0.620	-0.091	0.055 [0.046, 0.063]	yes	7.81e-04	6.29e-04	6.09e-07	0.222	4.10e-04	9.02e-04	6.79e-04	8.13e-07	0.245	1.88e-15
ridge	0.529	0.620	-0.091	0.055 [0.046, 0.063]	yes	7.81e-04	6.29e-04	6.09e-07	0.222	4.10e-04	9.02e-04	6.79e-04	8.13e-07	0.245	-1.74e-16
lasso	0.528	0.616	-0.087	0.051 [0.042, 0.058]	yes	7.85e-04	6.33e-04	6.16e-07	0.223	4.14e-04	9.02e-04	6.79e-04	8.14e-07	0.244	1.49e-19
elasticnet	0.528	0.611	-0.083	0.046 [0.037, 0.054]	yes	7.89e-04	6.37e-04	6.23e-07	0.224	4.19e-04	9.03e-04	6.80e-04	8.15e-07	0.243	1.32e-19
huber	0.521	0.659	-0.138	0.094 [0.090, 0.098]	yes	7.39e-04	5.86e-04	5.47e-07	0.208	3.26e-04	9.09e-04	6.70e-04	8.26e-07	0.256	-1.01e-04
poly2	0.563	0.615	-0.051	0.050 [0.037, 0.062]	yes	7.86e-04	6.37e-04	6.17e-07	0.231	4.48e-04	8.68e-04	6.47e-04	7.53e-07	0.231	1.89e-14
poly3	0.567	0.591	-0.024	0.026 [0.011, 0.040]	yes	8.10e-04	6.60e-04	6.56e-07	0.252	4.84e-04	8.64e-04	6.44e-04	7.47e-07	0.264	-2.65e-13
knn	0.626	0.590	0.036	0.025 [0.010, 0.038]	yes	8.10e-04	6.65e-04	6.57e-07	0.227	4.95e-04	8.03e-04	5.86e-04	6.44e-07	0.183	-7.38e-06
svr ★	0.617	0.668	-0.051	0.103 [0.091, 0.114]	yes	7.30e-04	5.86e-04	5.32e-07	0.198	3.83e-04	8.13e-04	5.84e-04	6.61e-07	0.196	-9.27e-05
kernelridge	0.618	0.600	0.018	0.036 [0.018, 0.049]	yes	8.00e-04	6.55e-04	6.40e-07	0.256	4.91e-04	8.12e-04	5.94e-04	6.59e-07	0.248	4.60e-06
gbm	0.565	0.538	0.027	-0.027 [-0.037, -0.016]	no	8.60e-04	7.16e-04	7.40e-07	0.261	5.42e-04	8.66e-04	6.44e-04	7.50e-07	0.205	2.94e-07
rf	0.695	0.512	0.183	-0.053 [-0.083, -0.025]	no	8.84e-04	7.34e-04	7.82e-07	0.258	5.79e-04	7.25e-04	5.25e-04	5.26e-07	0.147	4.49e-07
extratrees	0.663	0.559	0.104	-0.006 [-0.029, 0.015]	no	8.41e-04	6.99e-04	7.06e-07	0.245	5.43e-04	7.62e-04	5.55e-04	5.81e-07	0.162	-3.93e-20
mlp	0.637	0.567	0.070	0.002 [-0.023, 0.025]	no	8.33e-04	6.90e-04	6.94e-07	0.237	5.37e-04	7.91e-04	5.78e-04	6.26e-07	0.179	2.86e-06
gp	0.603	0.567	0.036	0.002 [-0.011, 0.014]	no	8.33e-04	6.79e-04	6.94e-07	0.236	5.10e-04	8.27e-04	6.03e-04	6.85e-07	0.192	-3.34e-06

best-by-test-R2 = **svr** (0.668) · linear 0.620 · persistence 0.565 · H=900s · 2657.2s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=30m

Sigma=30s · horizon=30m (H=1800s) — metrics: LE + AI

LE

Sigma=30s | horizon=30m | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.496	0.601	-0.105	0.067 [0.058, 0.074]	yes	1.11e-03	9.05e-04	1.24e-06	0.216	6.18e-04	1.28e-03	9.56e-04	1.64e-06	0.236	2.78e-15
ridge	0.496	0.601	-0.105	0.067 [0.058, 0.074]	yes	1.11e-03	9.05e-04	1.24e-06	0.216	6.18e-04	1.28e-03	9.56e-04	1.64e-06	0.236	-2.68e-17
lasso	0.496	0.597	-0.102	0.063 [0.055, 0.071]	yes	1.12e-03	9.09e-04	1.25e-06	0.217	6.21e-04	1.28e-03	9.57e-04	1.64e-06	0.236	1.43e-19
elasticnet	0.495	0.594	-0.099	0.060 [0.051, 0.067]	yes	1.12e-03	9.13e-04	1.26e-06	0.218	6.25e-04	1.28e-03	9.57e-04	1.64e-06	0.235	2.59e-19
huber ★	0.486	0.652	-0.166	0.118 [0.113, 0.123]	yes	1.04e-03	8.28e-04	1.08e-06	0.197	4.84e-04	1.29e-03	9.41e-04	1.67e-06	0.248	-1.58e-04
poly2	0.531	0.596	-0.066	0.062 [0.049, 0.075]	yes	1.12e-03	9.13e-04	1.25e-06	0.225	6.61e-04	1.24e-03	9.15e-04	1.53e-06	0.223	-5.88e-15
poly3	0.531	0.570	-0.038	0.036 [0.024, 0.047]	yes	1.16e-03	9.43e-04	1.34e-06	0.233	7.00e-04	1.24e-03	9.15e-04	1.53e-06	0.236	4.88e-13
knn	0.594	0.548	0.046	0.014 [-0.000, 0.025]	no	1.18e-03	9.77e-04	1.40e-06	0.229	7.36e-04	1.15e-03	8.32e-04	1.32e-06	0.183	-8.60e-06
svr	0.581	0.646	-0.065	0.112 [0.095, 0.127]	yes	1.05e-03	8.38e-04	1.10e-06	0.193	5.36e-04	1.17e-03	8.23e-04	1.36e-06	0.199	-1.60e-04
kernelridge	0.583	0.564	0.019	0.030 [0.015, 0.044]	yes	1.16e-03	9.61e-04	1.35e-06	0.248	7.30e-04	1.16e-03	8.45e-04	1.36e-06	0.230	8.27e-06
gbm	0.555	0.530	0.025	-0.004 [-0.014, 0.006]	no	1.21e-03	1.01e-03	1.46e-06	0.247	7.77e-04	1.20e-03	8.83e-04	1.45e-06	0.194	3.62e-07
rf	0.663	0.465	0.198	-0.069 [-0.100, -0.042]	no	1.29e-03	1.08e-03	1.66e-06	0.263	8.63e-04	1.05e-03	7.49e-04	1.10e-06	0.147	5.14e-07
extratrees	0.637	0.515	0.122	-0.019 [-0.041, -0.002]	no	1.23e-03	1.03e-03	1.51e-06	0.246	7.99e-04	1.09e-03	7.83e-04	1.18e-06	0.160	6.02e-21
mlp	0.603	0.557	0.046	0.023 [0.001, 0.041]	yes	1.17e-03	9.76e-04	1.37e-06	0.227	7.52e-04	1.14e-03	8.20e-04	1.29e-06	0.180	-9.27e-06
gp	0.572	0.551	0.021	0.017 [0.001, 0.031]	yes	1.18e-03	9.70e-04	1.39e-06	0.229	7.40e-04	1.18e-03	8.55e-04	1.39e-06	0.190	4.39e-07

best-by-test-R2 = **huber** (0.652) · linear 0.601 · persistence 0.534 · H=1800s · 2605.3s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI

Sigma=30s | horizon=30m | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.491	0.595	-0.104	0.061 [0.050, 0.070]	yes	1.12e-03	9.11e-04	1.26e-06	0.219	6.24e-04	1.29e-03	9.62e-04	1.66e-06	0.239	-1.66e-15
ridge	0.491	0.595	-0.104	0.061 [0.050, 0.070]	yes	1.12e-03	9.11e-04	1.26e-06	0.219	6.24e-04	1.29e-03	9.62e-04	1.66e-06	0.239	-3.02e-17
lasso	0.490	0.592	-0.101	0.058 [0.047, 0.067]	yes	1.13e-03	9.15e-04	1.27e-06	0.220	6.29e-04	1.29e-03	9.63e-04	1.66e-06	0.239	3.82e-19
elasticnet	0.490	0.588	-0.098	0.054 [0.044, 0.063]	yes	1.13e-03	9.19e-04	1.28e-06	0.221	6.33e-04	1.29e-03	9.63e-04	1.66e-06	0.238	2.47e-19
huber ★	0.481	0.645	-0.164	0.111 [0.106, 0.116]	yes	1.05e-03	8.36e-04	1.10e-06	0.201	4.88e-04	1.30e-03	9.47e-04	1.69e-06	0.252	-1.58e-04
poly2	0.525	0.581	-0.056	0.047 [0.031, 0.061]	yes	1.14e-03	9.30e-04	1.30e-06	0.230	6.74e-04	1.24e-03	9.22e-04	1.55e-06	0.226	6.41e-15
poly3	0.531	0.553	-0.023	0.019 [0.002, 0.034]	yes	1.18e-03	9.66e-04	1.39e-06	0.251	7.28e-04	1.24e-03	9.14e-04	1.53e-06	0.253	-8.16e-14
knn	0.592	0.546	0.046	0.012 [-0.006, 0.026]	no	1.19e-03	9.82e-04	1.41e-06	0.230	7.45e-04	1.15e-03	8.34e-04	1.33e-06	0.184	-1.33e-05
svr	0.579	0.642	-0.063	0.108 [0.097, 0.120]	yes	1.05e-03	8.49e-04	1.11e-06	0.196	5.63e-04	1.17e-03	8.27e-04	1.37e-06	0.200	-1.61e-04
kernelridge	0.583	0.557	0.026	0.023 [0.008, 0.037]	yes	1.17e-03	9.73e-04	1.37e-06	0.267	7.49e-04	1.16e-03	8.46e-04	1.36e-06	0.213	8.10e-06
gbm	0.552	0.528	0.024	-0.006 [-0.020, 0.005]	no	1.21e-03	1.02e-03	1.47e-06	0.250	7.88e-04	1.21e-03	8.88e-04	1.46e-06	0.196	2.62e-07
rf	0.660	0.436	0.225	-0.098 [-0.135, -0.066]	no	1.32e-03	1.11e-03	1.75e-06	0.270	8.98e-04	1.05e-03	7.54e-04	1.11e-06	0.148	5.24e-07
extratrees	0.634	0.494	0.140	-0.040 [-0.068, -0.014]	no	1.25e-03	1.05e-03	1.57e-06	0.254	8.38e-04	1.09e-03	7.87e-04	1.19e-06	0.161	4.89e-20
mlp	0.601	0.501	0.100	-0.033 [-0.066, -0.003]	no	1.24e-03	1.04e-03	1.55e-06	0.248	8.44e-04	1.14e-03	8.27e-04	1.30e-06	0.179	8.28e-06
gp	0.573	0.544	0.028	0.010 [-0.003, 0.023]	no	1.19e-03	9.80e-04	1.41e-06	0.232	7.53e-04	1.18e-03	8.56e-04	1.39e-06	0.191	-1.72e-08

best-by-test-R2 = **huber** (0.645) · linear 0.595 · persistence 0.534 · H=1800s · 2548.2s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=1h

Sigma=30s · horizon=1h (H=3600s) — metrics: LE + AI

LE

Sigma=30s | horizon=1h | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.430	0.549	-0.119	0.064 [0.054, 0.073]	yes	1.63e-03	1.34e-03	2.67e-06	0.226	9.50e-04	1.85e-03	1.37e-03	3.43e-06	0.243	3.13e-16
ridge	0.430	0.549	-0.119	0.064 [0.054, 0.073]	yes	1.63e-03	1.34e-03	2.67e-06	0.226	9.50e-04	1.85e-03	1.37e-03	3.43e-06	0.243	3.54e-17
lasso	0.430	0.546	-0.117	0.062 [0.051, 0.070]	yes	1.64e-03	1.34e-03	2.68e-06	0.227	9.54e-04	1.85e-03	1.37e-03	3.43e-06	0.243	-2.85e-19
elasticnet	0.430	0.544	-0.114	0.059 [0.048, 0.067]	yes	1.64e-03	1.34e-03	2.70e-06	0.228	9.58e-04	1.85e-03	1.37e-03	3.43e-06	0.243	-9.71e-20

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
huber ★	0.417	0.621	-0.204	0.136 [0.129, 0.143]	yes	1.50e-03	1.19e-03	2.25e-06	0.200	7.18e-04	1.87e-03	1.34e-03	3.51e-06	0.259	-2.57e-04
poly2	0.463	0.529	-0.065	0.044 [0.029, 0.059]	yes	1.67e-03	1.37e-03	2.79e-06	0.237	1.01e-03	1.80e-03	1.32e-03	3.23e-06	0.231	1.52e-14
poly3	0.465	0.494	-0.029	0.009 [-0.006, 0.022]	no	1.73e-03	1.42e-03	3.00e-06	0.253	1.08e-03	1.79e-03	1.32e-03	3.22e-06	0.240	1.65e-12
knn	0.532	0.474	0.058	-0.011 [-0.028, 0.002]	no	1.76e-03	1.46e-03	3.11e-06	0.244	1.12e-03	1.68e-03	1.21e-03	2.81e-06	0.194	-1.98e-05
svr	0.511	0.616	-0.105	0.131 [0.119, 0.145]	yes	1.51e-03	1.21e-03	2.27e-06	0.196	7.90e-04	1.71e-03	1.19e-03	2.94e-06	0.216	-2.89e-04
kernelridge	0.520	0.499	0.021	0.014 [-0.003, 0.027]	no	1.72e-03	1.44e-03	2.97e-06	0.259	1.11e-03	1.70e-03	1.22e-03	2.89e-06	0.250	-8.76e-06
gbm	0.508	0.478	0.030	-0.007 [-0.021, 0.004]	no	1.76e-03	1.48e-03	3.09e-06	0.253	1.15e-03	1.72e-03	1.25e-03	2.96e-06	0.199	-1.07e-07
rf	0.586	0.400	0.186	-0.085 [-0.119, -0.057]	no	1.88e-03	1.58e-03	3.55e-06	0.272	1.26e-03	1.58e-03	1.13e-03	2.49e-06	0.169	2.18e-07
extratrees	0.583	0.426	0.157	-0.059 [-0.085, -0.039]	no	1.84e-03	1.55e-03	3.40e-06	0.265	1.22e-03	1.58e-03	1.13e-03	2.51e-06	0.167	1.05e-19
mlp	0.545	0.460	0.085	-0.025 [-0.055, 0.001]	no	1.79e-03	1.50e-03	3.20e-06	0.251	1.19e-03	1.65e-03	1.19e-03	2.73e-06	0.188	-6.62e-06
gp	0.508	0.464	0.043	-0.021 [-0.035, -0.008]	no	1.78e-03	1.47e-03	3.17e-06	0.249	1.14e-03	1.72e-03	1.24e-03	2.96e-06	0.202	-1.75e-05

best-by-test-R2 = **huber** (0.621) · linear 0.549 · persistence 0.485 · H=3600s · 2600.2s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI Sigma=30s | horizon=1h | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.425	0.543	-0.118	0.058 [0.045, 0.069]	yes	1.64e-03	1.35e-03	2.71e-06	0.229	9.60e-04	1.86e-03	1.38e-03	3.46e-06	0.246	2.61e-16
ridge	0.425	0.543	-0.118	0.058 [0.045, 0.069]	yes	1.64e-03	1.35e-03	2.71e-06	0.229	9.60e-04	1.86e-03	1.38e-03	3.46e-06	0.246	3.53e-17
lasso	0.425	0.540	-0.115	0.055 [0.043, 0.066]	yes	1.65e-03	1.35e-03	2.72e-06	0.229	9.65e-04	1.86e-03	1.38e-03	3.46e-06	0.246	2.27e-19
elasticnet	0.425	0.537	-0.113	0.052 [0.040, 0.063]	yes	1.65e-03	1.35e-03	2.74e-06	0.230	9.69e-04	1.86e-03	1.38e-03	3.46e-06	0.246	-6.47e-21
huber ★	0.412	0.614	-0.202	0.129 [0.123, 0.136]	yes	1.51e-03	1.20e-03	2.28e-06	0.204	7.25e-04	1.88e-03	1.35e-03	3.54e-06	0.263	-2.57e-04
poly2	0.457	0.512	-0.055	0.027 [0.009, 0.045]	yes	1.70e-03	1.39e-03	2.89e-06	0.242	1.03e-03	1.81e-03	1.33e-03	3.26e-06	0.233	4.20e-14
poly3	0.466	0.478	-0.012	-0.007 [-0.027, 0.011]	no	1.76e-03	1.45e-03	3.09e-06	0.257	1.12e-03	1.79e-03	1.32e-03	3.21e-06	0.256	6.36e-13
knn	0.532	0.470	0.062	-0.015 [-0.037, 0.005]	no	1.77e-03	1.47e-03	3.14e-06	0.245	1.14e-03	1.68e-03	1.21e-03	2.81e-06	0.194	-2.57e-05
svr	0.511	0.598	-0.087	0.113 [0.100, 0.128]	yes	1.54e-03	1.24e-03	2.38e-06	0.203	8.30e-04	1.71e-03	1.19e-03	2.94e-06	0.215	-2.85e-04
kernelridge	0.521	0.486	0.035	0.001 [-0.019, 0.018]	no	1.74e-03	1.46e-03	3.04e-06	0.336	1.14e-03	1.70e-03	1.22e-03	2.88e-06	0.251	-1.08e-05
gbm	0.505	0.448	0.057	-0.037 [-0.057, -0.019]	no	1.81e-03	1.52e-03	3.27e-06	0.263	1.20e-03	1.73e-03	1.26e-03	2.98e-06	0.201	-1.26e-07
rf	0.586	0.340	0.246	-0.145 [-0.190, -0.107]	no	1.98e-03	1.65e-03	3.91e-06	0.284	1.34e-03	1.58e-03	1.13e-03	2.49e-06	0.170	1.54e-07
extratrees	0.583	0.392	0.191	-0.093 [-0.128, -0.060]	no	1.90e-03	1.60e-03	3.60e-06	0.274	1.29e-03	1.58e-03	1.13e-03	2.51e-06	0.168	4.21e-20
mlp	0.546	0.402	0.144	-0.083 [-0.123, -0.047]	no	1.88e-03	1.57e-03	3.54e-06	0.265	1.26e-03	1.65e-03	1.19e-03	2.73e-06	0.186	7.25e-06
gp	0.510	0.460	0.050	-0.025 [-0.042, -0.011]	no	1.79e-03	1.49e-03	3.20e-06	0.251	1.17e-03	1.72e-03	1.24e-03	2.95e-06	0.201	-1.66e-05

best-by-test-R2 = **huber** (0.614) · linear 0.543 · persistence 0.485 · H=3600s · 2489.9s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

Sigma=45s cov-duration 45s — 10 horizon(s) finished

horizon=nextwin Sigma=45s · horizon=nextwin (H=45s) — metrics: LE + AI

LE Sigma=45s | horizon=nextwin | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.485	0.554	-0.069	0.031 [0.029, 0.034]	yes	2.44e-04	1.92e-04	5.97e-08	0.505	5.78e-05	2.91e-04	2.22e-04	8.49e-08	0.640	1.78e-16
ridge	0.485	0.554	-0.069	0.031 [0.029, 0.034]	yes	2.44e-04	1.92e-04	5.97e-08	0.505	5.79e-05	2.91e-04	2.22e-04	8.49e-08	0.640	-2.60e-18
lasso	0.482	0.544	-0.062	0.022 [0.020, 0.024]	yes	2.47e-04	1.94e-04	6.10e-08	0.474	6.11e-05	2.92e-04	2.22e-04	8.55e-08	0.594	6.33e-21
elasticnet	0.474	0.531	-0.057	0.009 [0.007, 0.010]	yes	2.51e-04	1.97e-04	6.28e-08	0.464	6.46e-05	2.95e-04	2.24e-04	8.67e-08	0.576	-3.06e-20
huber	0.483	0.563	-0.080	0.041 [0.039, 0.042]	yes	2.42e-04	1.90e-04	5.85e-08	0.561	4.28e-05	2.92e-04	2.21e-04	8.52e-08	0.716	-1.52e-05
poly2	0.503	0.566	-0.063	0.043 [0.039, 0.047]	yes	2.41e-04	1.90e-04	5.81e-08	0.413	6.62e-05	2.86e-04	2.18e-04	8.20e-08	0.530	7.55e-15
poly3	0.501	0.557	-0.056	0.035 [0.032, 0.037]	yes	2.44e-04	1.92e-04	5.93e-08	0.478	6.91e-05	2.87e-04	2.18e-04	8.23e-08	0.555	2.35e-14
knn	0.532	0.545	-0.013	0.022 [0.016, 0.028]	yes	2.47e-04	1.95e-04	6.10e-08	0.404	7.67e-05	2.78e-04	2.10e-04	7.73e-08	0.485	-2.17e-06
svr ★	0.528	0.576	-0.048	0.053 [0.046, 0.060]	yes	2.38e-04	1.87e-04	5.68e-08	0.415	6.15e-05	2.79e-04	2.10e-04	7.79e-08	0.533	-1.50e-05
kernelridge	0.526	0.558	-0.033	0.036 [0.024, 0.045]	yes	2.43e-04	1.91e-04	5.92e-08	0.867	7.26e-05	2.80e-04	2.11e-04	7.83e-08	0.790	-1.81e-06
gbm	0.427	0.433	-0.006	-0.089 [-0.098, -0.079]	no	2.75e-04	2.18e-04	7.59e-08	0.388	1.10e-04	3.07e-04	2.34e-04	9.45e-08	0.420	-2.77e-08
rf	0.542	0.559	-0.017	0.036 [0.031, 0.041]	yes	2.43e-04	1.93e-04	5.91e-08	0.392	7.88e-05	2.75e-04	2.09e-04	7.56e-08	0.460	2.10e-08
extratrees	0.564	0.558	0.006	0.035 [0.030, 0.040]	yes	2.43e-04	1.93e-04	5.92e-08	0.388	8.01e-05	2.68e-04	2.03e-04	7.19e-08	0.422	3.10e-21
mlp	0.541	0.566	-0.025	0.043 [0.037, 0.049]	yes	2.41e-04	1.91e-04	5.82e-08	0.408	7.56e-05	2.75e-04	2.09e-04	7.58e-08	0.475	-1.56e-06
gp	0.515	0.541	-0.026	0.019 [0.004, 0.030]	yes	2.48e-04	1.94e-04	6.14e-08	0.410	7.83e-05	2.83e-04	2.13e-04	8.00e-08	0.524	-3.74e-06

best-by-test-R2 = **svr** (0.576) · linear 0.554 · persistence 0.523 · H=45s · 2249.6s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI Sigma=45s | horizon=nextwin | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.483	0.550	-0.068	0.028 [0.025, 0.030]	yes	2.45e-04	1.93e-04	6.02e-08	0.518	5.67e-05	2.92e-04	2.23e-04	8.54e-08	0.645	-9.46e-17
ridge	0.483	0.550	-0.067	0.028 [0.025, 0.030]	yes	2.45e-04	1.93e-04	6.02e-08	0.517	5.67e-05	2.92e-04	2.23e-04	8.54e-08	0.645	-2.33e-18
lasso	0.480	0.542	-0.061	0.019 [0.017, 0.021]	yes	2.48e-04	1.95e-04	6.14e-08	0.479	6.11e-05	2.93e-04	2.23e-04	8.58e-08	0.596	-4.17e-20
elasticnet	0.474	0.530	-0.056	0.007 [0.005, 0.009]	yes	2.51e-04	1.97e-04	6.30e-08	0.463	6.59e-05	2.95e-04	2.24e-04	8.68e-08	0.573	0.00e+00
huber	0.481	0.559	-0.078	0.036 [0.035, 0.038]	yes	2.43e-04	1.91e-04	5.91e-08	0.573	4.20e-05	2.93e-04	2.22e-04	8.57e-08	0.723	-1.51e-05
poly2	0.499	0.558	-0.059	0.036 [0.032, 0.040]	yes	2.43e-04	1.92e-04	5.91e-08	0.414	6.74e-05	2.87e-04	2.19e-04	8.26e-08	0.522	7.84e-15
poly3	0.502	0.552	-0.051	0.030 [0.026, 0.034]	yes	2.45e-04	1.93e-04	5.99e-08	0.573	7.26e-05	2.87e-04	2.18e-04	8.22e-08	0.527	2.38e-13
knn	0.528	0.544	-0.016	0.022 [0.016, 0.027]	yes	2.47e-04	1.95e-04	6.10e-08	0.402	7.62e-05	2.79e-04	2.11e-04	7.78e-08	0.488	-2.81e-06
svr ★	0.525	0.573	-0.048	0.050 [0.043, 0.056]	yes	2.39e-04	1.88e-04	5.72e-08	0.412	6.15e-05	2.80e-04	2.11e-04	7.84e-08	0.528	-1.45e-05
kernelridge	0.523	0.554	-0.031	0.031 [0.019, 0.041]	yes	2.44e-04	1.92e-04	5.97e-08	0.985	7.31e-05	2.80e-04	2.11e-04	7.86e-08	0.820	-1.66e-06
gbm	0.426	0.433	-0.007	-0.090 [-0.098, -0.079]	no	2.76e-04	2.18e-04	7.59e-08	0.389	1.09e-04	3.08e-04	2.34e-04	9.47e-08	0.422	-2.98e-08
rf	0.538	0.556	-0.017	0.033 [0.027, 0.038]	yes	2.44e-04	1.93e-04	5.95e-08	0.389	7.91e-05	2.76e-04	2.09e-04	7.62e-08	0.461	2.27e-08
extratrees	0.561	0.553	0.008	0.031 [0.025, 0.036]	yes	2.45e-04	1.94e-04	5.98e-08	0.385	8.16e-05	2.69e-04	2.04e-04	7.24e-08	0.423	1.58e-21
mlp	0.537	0.560	-0.023	0.038 [0.031, 0.043]	yes	2.43e-04	1.92e-04	5.89e-08	0.400	7.90e-05	2.77e-04	2.10e-04	7.65e-08	0.473	1.92e-06
gp	0.515	0.545	-0.029	0.022 [0.010, 0.031]	yes	2.47e-04	1.93e-04	6.10e-08	0.395	7.70e-05	2.83e-04	2.14e-04	8.00e-08	0.519	-3.38e-06

best-by-test-R2 = **svr** (0.573) · linear 0.550 · persistence 0.523 · H=45s · 2172.8s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=30s

Sigma=45s · horizon=30s (H=30s) — metrics: LE + AI

LE

Sigma=45s | horizon=30s | metric=LE

N/A — horizon < cov window (H=30s ≤ Sigma=45s). horizon 30s < cov 45s (cover 0.0%)

AI

Sigma=45s | horizon=30s | metric=AI

N/A — horizon < cov window (H=30s ≤ Sigma=45s). horizon 30s < cov 45s (cover 0.0%)

horizon=45s

Sigma=45s · horizon=45s (H=45s) — metrics: LE + AI

LE

Sigma=45s | horizon=45s | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.483	0.552	-0.069	0.033 [0.030, 0.036]	yes	2.38e-04	1.89e-04	5.67e-08	0.491	5.81e-05	2.82e-04	2.17e-04	7.96e-08	0.623	7.07e-16
ridge	0.483	0.552	-0.069	0.033 [0.030, 0.036]	yes	2.38e-04	1.89e-04	5.67e-08	0.491	5.81e-05	2.82e-04	2.17e-04	7.96e-08	0.623	3.08e-17
lasso	0.480	0.542	-0.062	0.023 [0.020, 0.025]	yes	2.41e-04	1.91e-04	5.80e-08	0.460	6.13e-05	2.83e-04	2.18e-04	8.01e-08	0.577	-3.12e-20
elasticnet	0.471	0.528	-0.058	0.009 [0.007, 0.011]	yes	2.44e-04	1.93e-04	5.97e-08	0.450	6.48e-05	2.85e-04	2.19e-04	8.15e-08	0.559	-4.58e-20
huber	0.481	0.562	-0.080	0.042 [0.040, 0.044]	yes	2.36e-04	1.87e-04	5.55e-08	0.539	4.42e-05	2.83e-04	2.17e-04	7.99e-08	0.690	-1.42e-05
poly2	0.502	0.565	-0.063	0.046 [0.042, 0.050]	yes	2.35e-04	1.86e-04	5.50e-08	0.404	6.65e-05	2.77e-04	2.13e-04	7.66e-08	0.518	2.07e-15
poly3	0.501	0.557	-0.056	0.038 [0.034, 0.040]	yes	2.37e-04	1.88e-04	5.61e-08	0.410	6.93e-05	2.77e-04	2.14e-04	7.69e-08	0.542	1.52e-13
knn	0.532	0.543	-0.011	0.023 [0.016, 0.029]	yes	2.41e-04	1.91e-04	5.79e-08	0.395	7.67e-05	2.68e-04	2.06e-04	7.21e-08	0.472	-1.91e-06
svr ★	0.530	0.578	-0.048	0.059 [0.055, 0.063]	yes	2.31e-04	1.84e-04	5.33e-08	0.410	6.18e-05	2.69e-04	2.06e-04	7.24e-08	0.515	-1.52e-05
kernelridge	0.530	0.562	-0.032	0.043 [0.037, 0.048]	yes	2.35e-04	1.87e-04	5.54e-08	0.443	7.28e-05	2.69e-04	2.06e-04	7.23e-08	0.576	-2.88e-06
gbm	0.426	0.430	-0.004	-0.089 [-0.098, -0.079]	no	2.69e-04	2.13e-04	7.21e-08	0.378	1.08e-04	2.97e-04	2.29e-04	8.84e-08	0.409	-6.20e-08
rf	0.541	0.559	-0.018	0.039 [0.033, 0.045]	yes	2.36e-04	1.89e-04	5.58e-08	0.384	7.79e-05	2.66e-04	2.04e-04	7.06e-08	0.448	2.92e-08
extratrees	0.564	0.557	0.007	0.037 [0.031, 0.043]	yes	2.37e-04	1.89e-04	5.61e-08	0.380	7.93e-05	2.59e-04	1.99e-04	6.72e-08	0.411	2.06e-19
mlp	0.540	0.564	-0.024	0.045 [0.037, 0.050]	yes	2.35e-04	1.88e-04	5.52e-08	0.403	7.41e-05	2.66e-04	2.04e-04	7.09e-08	0.456	1.22e-06
gp	0.522	0.562	-0.041	0.043 [0.037, 0.049]	yes	2.35e-04	1.87e-04	5.54e-08	0.380	7.54e-05	2.71e-04	2.08e-04	7.36e-08	0.472	-2.34e-06

best-by-test-R2 = **svr** (0.578) · linear 0.552 · persistence 0.519 · H=45s · 2060.2s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI

Sigma=45s | horizon=45s | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.480	0.549	-0.068	0.029 [0.026, 0.032]	yes	2.39e-04	1.90e-04	5.71e-08	0.501	5.72e-05	2.83e-04	2.18e-04	8.00e-08	0.629	-4.16e-16
ridge	0.480	0.549	-0.068	0.029 [0.026, 0.032]	yes	2.39e-04	1.90e-04	5.71e-08	0.500	5.72e-05	2.83e-04	2.18e-04	8.00e-08	0.628	3.16e-17
lasso	0.478	0.539	-0.062	0.020 [0.017, 0.023]	yes	2.41e-04	1.91e-04	5.83e-08	0.463	6.16e-05	2.84e-04	2.18e-04	8.04e-08	0.579	4.42e-21
elasticnet	0.471	0.527	-0.056	0.008 [0.006, 0.010]	yes	2.45e-04	1.94e-04	5.98e-08	0.447	6.64e-05	2.86e-04	2.19e-04	8.15e-08	0.557	-2.48e-20
huber	0.479	0.557	-0.079	0.038 [0.036, 0.040]	yes	2.37e-04	1.87e-04	5.60e-08	0.548	4.35e-05	2.83e-04	2.18e-04	8.03e-08	0.697	-1.41e-05
poly2	0.499	0.557	-0.058	0.038 [0.033, 0.042]	yes	2.37e-04	1.88e-04	5.61e-08	0.403	6.82e-05	2.78e-04	2.14e-04	7.72e-08	0.510	6.04e-15
poly3	0.502	0.552	-0.051	0.033 [0.028, 0.037]	yes	2.38e-04	1.90e-04	5.67e-08	0.394	7.33e-05	2.77e-04	2.14e-04	7.68e-08	0.514	3.24e-13
knn	0.528	0.543	-0.015	0.024 [0.018, 0.029]	yes	2.40e-04	1.91e-04	5.78e-08	0.393	7.67e-05	2.69e-04	2.07e-04	7.26e-08	0.475	-2.41e-06
svr ★	0.527	0.574	-0.047	0.055 [0.050, 0.060]	yes	2.32e-04	1.84e-04	5.39e-08	0.411	6.11e-05	2.70e-04	2.06e-04	7.28e-08	0.520	-1.48e-05
kernelridge	0.527	0.556	-0.029	0.037 [0.029, 0.043]	yes	2.37e-04	1.88e-04	5.62e-08	0.494	7.27e-05	2.70e-04	2.07e-04	7.28e-08	0.626	-2.81e-06
gbm	0.424	0.428	-0.004	-0.091 [-0.100, -0.082]	no	2.69e-04	2.14e-04	7.24e-08	0.379	1.09e-04	2.98e-04	2.29e-04	8.87e-08	0.411	-6.48e-08
rf	0.538	0.555	-0.018	0.036 [0.030, 0.041]	yes	2.37e-04	1.89e-04	5.63e-08	0.380	7.84e-05	2.67e-04	2.05e-04	7.12e-08	0.449	2.41e-08
extratrees	0.561	0.552	0.009	0.032 [0.026, 0.038]	yes	2.38e-04	1.90e-04	5.67e-08	0.377	8.14e-05	2.60e-04	1.99e-04	6.76e-08	0.414	1.97e-19
mlp	0.536	0.564	-0.029	0.045 [0.038, 0.051]	yes	2.35e-04	1.88e-04	5.51e-08	0.394	7.38e-05	2.67e-04	2.05e-04	7.15e-08	0.459	-1.01e-06
gp	0.520	0.558	-0.038	0.038 [0.030, 0.045]	yes	2.37e-04	1.88e-04	5.60e-08	0.389	7.43e-05	2.72e-04	2.09e-04	7.39e-08	0.469	-2.72e-06

best-by-test-R2 = **svr** (0.574) · linear 0.549 · persistence 0.519 · H=45s · 2497.0s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=60s

Sigma=45s · horizon=60s (H=60s) — metrics: LE + AI

LE

Sigma=45s | horizon=60s | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.483	0.552	-0.069	0.033 [0.030, 0.036]	yes	2.38e-04	1.89e-04	5.66e-08	0.491	5.81e-05	2.82e-04	2.17e-04	7.96e-08	0.624	-3.65e-16
ridge	0.483	0.552	-0.069	0.033 [0.030, 0.036]	yes	2.38e-04	1.89e-04	5.66e-08	0.491	5.81e-05	2.82e-04	2.17e-04	7.96e-08	0.623	-2.58e-17
lasso	0.480	0.542	-0.062	0.023 [0.020, 0.025]	yes	2.41e-04	1.91e-04	5.79e-08	0.460	6.13e-05	2.83e-04	2.18e-04	8.01e-08	0.577	1.33e-20
elasticnet	0.471	0.528	-0.057	0.009 [0.008, 0.011]	yes	2.44e-04	1.93e-04	5.97e-08	0.450	6.48e-05	2.85e-04	2.19e-04	8.15e-08	0.559	2.43e-20
huber	0.481	0.562	-0.080	0.042 [0.040, 0.044]	yes	2.35e-04	1.87e-04	5.55e-08	0.539	4.42e-05	2.83e-04	2.17e-04	7.98e-08	0.690	-1.42e-05
poly2	0.502	0.565	-0.063	0.046 [0.042, 0.050]	yes	2.35e-04	1.86e-04	5.50e-08	0.404	6.65e-05	2.77e-04	2.13e-04	7.66e-08	0.518	2.78e-16
poly3	0.501	0.557	-0.056	0.038 [0.034, 0.040]	yes	2.37e-04	1.88e-04	5.61e-08	0.410	6.93e-05	2.77e-04	2.14e-04	7.68e-08	0.542	1.74e-13
knn	0.532	0.544	-0.012	0.025 [0.018, 0.031]	yes	2.40e-04	1.91e-04	5.76e-08	0.395	7.62e-05	2.68e-04	2.06e-04	7.21e-08	0.477	-2.81e-06
svr ★	0.530	0.579	-0.048	0.059 [0.055, 0.064]	yes	2.31e-04	1.84e-04	5.33e-08	0.402	6.26e-05	2.69e-04	2.06e-04	7.23e-08	0.522	-1.16e-05
kernelridge	0.528	0.562	-0.033	0.042 [0.034, 0.049]	yes	2.35e-04	1.87e-04	5.54e-08	0.415	7.18e-05	2.69e-04	2.06e-04	7.26e-08	0.679	5.31e-08
gbm	0.426	0.431	-0.005	-0.088 [-0.096, -0.078]	no	2.68e-04	2.13e-04	7.19e-08	0.378	1.08e-04	2.97e-04	2.28e-04	8.83e-08	0.409	2.82e-08
rf	0.541	0.559	-0.018	0.040 [0.034, 0.046]	yes	2.36e-04	1.89e-04	5.58e-08	0.384	7.76e-05	2.66e-04	2.04e-04	7.06e-08	0.448	-6.52e-09
extratrees	0.564	0.556	0.008	0.037 [0.031, 0.042]	yes	2.37e-04	1.89e-04	5.61e-08	0.380	7.95e-05	2.59e-04	1.99e-04	6.71e-08	0.412	2.06e-19
mlp	0.540	0.561	-0.022	0.042 [0.035, 0.048]	yes	2.36e-04	1.88e-04	5.55e-08	0.405	7.63e-05	2.66e-04	2.04e-04	7.09e-08	0.459	1.45e-06
gp	0.523	0.569	-0.046	0.050 [0.045, 0.054]	yes	2.33e-04	1.86e-04	5.45e-08	0.384	7.32e-05	2.71e-04	2.09e-04	7.35e-08	0.486	-1.37e-06

best-by-test-R2 = **svr** (0.579) · linear 0.552 · persistence 0.519 · H=60s · 2070.2s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI

Sigma=45s | horizon=60s | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.480	0.548	-0.068	0.029 [0.026, 0.032]	yes	2.39e-04	1.90e-04	5.71e-08	0.501	5.72e-05	2.83e-04	2.18e-04	8.00e-08	0.629	-1.76e-17
ridge	0.480	0.548	-0.068	0.029 [0.026, 0.032]	yes	2.39e-04	1.90e-04	5.71e-08	0.500	5.72e-05	2.83e-04	2.18e-04	8.00e-08	0.629	-2.69e-17

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
lasso	0.478	0.539	-0.062	0.020 [0.018, 0.023]	yes	2.41e-04	1.91e-04	5.83e-08	0.463	6.16e-05	2.84e-04	2.18e-04	8.04e-08	0.579	2.43e-20
elasticnet	0.471	0.527	-0.056	0.008 [0.006, 0.010]	yes	2.45e-04	1.94e-04	5.98e-08	0.447	6.64e-05	2.85e-04	2.19e-04	8.15e-08	0.557	-2.21e-21
huber	0.479	0.557	-0.079	0.038 [0.036, 0.040]	yes	2.37e-04	1.87e-04	5.60e-08	0.548	4.35e-05	2.83e-04	2.17e-04	8.03e-08	0.697	-1.41e-05
poly2	0.499	0.557	-0.058	0.038 [0.033, 0.042]	yes	2.37e-04	1.88e-04	5.60e-08	0.403	6.82e-05	2.78e-04	2.14e-04	7.72e-08	0.511	-1.69e-15
poly3	0.501	0.552	-0.050	0.033 [0.029, 0.037]	yes	2.38e-04	1.89e-04	5.67e-08	0.394	7.33e-05	2.77e-04	2.13e-04	7.67e-08	0.514	-8.00e-14
knn	0.528	0.544	-0.016	0.025 [0.019, 0.030]	yes	2.40e-04	1.91e-04	5.76e-08	0.394	7.65e-05	2.69e-04	2.07e-04	7.26e-08	0.481	-3.36e-06
svr ★	0.528	0.575	-0.047	0.056 [0.051, 0.060]	yes	2.32e-04	1.84e-04	5.38e-08	0.400	6.35e-05	2.70e-04	2.06e-04	7.27e-08	0.524	-1.08e-05
kernelridge	0.526	0.558	-0.032	0.039 [0.032, 0.045]	yes	2.36e-04	1.88e-04	5.59e-08	0.465	7.31e-05	2.70e-04	2.07e-04	7.29e-08	0.740	1.62e-07
gbm	0.424	0.429	-0.004	-0.091 [-0.099, -0.081]	no	2.69e-04	2.13e-04	7.23e-08	0.379	1.08e-04	2.98e-04	2.29e-04	8.86e-08	0.411	3.72e-08
rf	0.538	0.555	-0.018	0.036 [0.030, 0.041]	yes	2.37e-04	1.89e-04	5.63e-08	0.380	7.83e-05	2.67e-04	2.05e-04	7.12e-08	0.449	-6.48e-09
extratrees	0.561	0.552	0.009	0.032 [0.026, 0.038]	yes	2.38e-04	1.90e-04	5.67e-08	0.377	8.15e-05	2.60e-04	1.99e-04	6.76e-08	0.414	1.98e-19
mlp	0.536	0.557	-0.020	0.038 [0.030, 0.044]	yes	2.37e-04	1.89e-04	5.61e-08	0.383	8.04e-05	2.67e-04	2.05e-04	7.14e-08	0.463	6.98e-07
gp	0.520	0.561	-0.042	0.042 [0.037, 0.047]	yes	2.35e-04	1.87e-04	5.55e-08	0.384	7.44e-05	2.72e-04	2.09e-04	7.39e-08	0.485	-1.10e-06

best-by-test-R2 = **svr** (0.575) · linear 0.548 · persistence 0.519 · H=60s · 2088.6s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=2mSigma=45s · horizon=2m (H=120s) — metrics: LE + AI

LESigma=45s | horizon=2m | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.526	0.605	-0.079	0.034 [0.030, 0.038]	yes	2.85e-04	2.25e-04	8.12e-08	0.357	9.19e-05	3.39e-04	2.58e-04	1.15e-07	0.415	9.42e-17
ridge	0.526	0.605	-0.079	0.034 [0.030, 0.038]	yes	2.85e-04	2.25e-04	8.13e-08	0.357	9.20e-05	3.39e-04	2.58e-04	1.15e-07	0.415	4.83e-18
lasso	0.524	0.597	-0.073	0.026 [0.022, 0.029]	yes	2.88e-04	2.27e-04	8.30e-08	0.352	9.52e-05	3.40e-04	2.59e-04	1.15e-07	0.401	-2.40e-20
elasticnet	0.518	0.586	-0.068	0.015 [0.012, 0.017]	yes	2.92e-04	2.30e-04	8.53e-08	0.352	9.86e-05	3.42e-04	2.60e-04	1.17e-07	0.395	-2.62e-20
huber	0.523	0.620	-0.097	0.049 [0.047, 0.051]	yes	2.80e-04	2.19e-04	7.82e-08	0.357	7.44e-05	3.40e-04	2.57e-04	1.15e-07	0.427	-2.11e-05
poly2	0.547	0.615	-0.068	0.044 [0.037, 0.050]	yes	2.82e-04	2.22e-04	7.94e-08	0.337	1.05e-04	3.31e-04	2.52e-04	1.10e-07	0.369	-6.39e-15
poly3	0.547	0.604	-0.058	0.033 [0.028, 0.038]	yes	2.85e-04	2.25e-04	8.15e-08	0.342	1.10e-04	3.31e-04	2.52e-04	1.10e-07	0.370	-9.13e-15
knn	0.581	0.587	-0.006	0.016 [0.008, 0.022]	yes	2.92e-04	2.30e-04	8.51e-08	0.330	1.20e-04	3.19e-04	2.41e-04	1.02e-07	0.327	-2.34e-06
svr ★	0.578	0.635	-0.058	0.064 [0.059, 0.070]	yes	2.74e-04	2.15e-04	7.51e-08	0.316	9.38e-05	3.20e-04	2.41e-04	1.02e-07	0.342	-2.13e-05
kernelridge	0.576	0.612	-0.036	0.041 [0.032, 0.049]	yes	2.83e-04	2.22e-04	7.99e-08	0.633	1.10e-04	3.20e-04	2.41e-04	1.03e-07	0.454	-3.14e-06
gbm	0.465	0.461	0.005	-0.110 [-0.119, -0.099]	no	3.33e-04	2.65e-04	1.11e-07	0.382	1.60e-04	3.60e-04	2.72e-04	1.29e-07	0.345	8.17e-09
rf	0.590	0.603	-0.012	0.032 [0.022, 0.040]	yes	2.86e-04	2.27e-04	8.18e-08	0.327	1.20e-04	3.15e-04	2.38e-04	9.92e-08	0.318	-2.18e-09
extratrees	0.613	0.600	0.013	0.029 [0.020, 0.037]	yes	2.87e-04	2.28e-04	8.24e-08	0.327	1.22e-04	3.06e-04	2.31e-04	9.38e-08	0.295	6.86e-20
mlp	0.589	0.610	-0.021	0.039 [0.028, 0.049]	yes	2.83e-04	2.25e-04	8.03e-08	0.323	1.17e-04	3.15e-04	2.38e-04	9.94e-08	0.322	-1.50e-06
gp	0.567	0.591	-0.024	0.020 [0.014, 0.025]	yes	2.90e-04	2.29e-04	8.42e-08	0.331	1.25e-04	3.24e-04	2.44e-04	1.05e-07	0.340	-4.90e-06

best-by-test-R2 = **svr** (0.635) · linear 0.605 · persistence 0.571 · H=120s · 2107.2s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AISigma=45s | horizon=2m | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.523	0.601	-0.078	0.030 [0.026, 0.034]	yes	2.87e-04	2.26e-04	8.21e-08	0.361	9.18e-05	3.40e-04	2.60e-04	1.16e-07	0.418	-2.63e-16
ridge	0.523	0.601	-0.078	0.030 [0.026, 0.034]	yes	2.87e-04	2.26e-04	8.21e-08	0.361	9.18e-05	3.40e-04	2.60e-04	1.16e-07	0.418	5.24e-18
lasso	0.521	0.593	-0.072	0.022 [0.018, 0.026]	yes	2.89e-04	2.28e-04	8.38e-08	0.355	9.63e-05	3.41e-04	2.60e-04	1.16e-07	0.404	-2.35e-20
elasticnet	0.516	0.583	-0.067	0.012 [0.009, 0.015]	yes	2.93e-04	2.31e-04	8.59e-08	0.353	1.01e-04	3.42e-04	2.61e-04	1.17e-07	0.395	-4.70e-20
huber	0.520	0.615	-0.095	0.044 [0.042, 0.046]	yes	2.82e-04	2.21e-04	7.92e-08	0.362	7.42e-05	3.41e-04	2.58e-04	1.16e-07	0.432	-2.10e-05
poly2	0.543	0.604	-0.061	0.033 [0.026, 0.040]	yes	2.85e-04	2.25e-04	8.15e-08	0.341	1.08e-04	3.33e-04	2.53e-04	1.11e-07	0.368	8.99e-16
poly3	0.548	0.597	-0.050	0.026 [0.018, 0.033]	yes	2.88e-04	2.28e-04	8.29e-08	0.364	1.16e-04	3.31e-04	2.52e-04	1.10e-07	0.360	1.06e-13
knn	0.578	0.586	-0.008	0.015 [0.007, 0.021]	yes	2.92e-04	2.31e-04	8.53e-08	0.331	1.21e-04	3.20e-04	2.42e-04	1.02e-07	0.329	-3.24e-06
svr ★	0.575	0.631	-0.055	0.060 [0.054, 0.065]	yes	2.76e-04	2.17e-04	7.60e-08	0.316	9.41e-05	3.21e-04	2.41e-04	1.03e-07	0.344	-2.10e-05
kernelridge	0.575	0.607	-0.032	0.036 [0.026, 0.044]	yes	2.85e-04	2.24e-04	8.10e-08	0.543	1.11e-04	3.21e-04	2.42e-04	1.03e-07	0.424	-3.03e-06
gbm	0.463	0.462	0.000	-0.109 [-0.117, -0.099]	no	3.33e-04	2.64e-04	1.11e-07	0.382	1.59e-04	3.61e-04	2.73e-04	1.30e-07	0.348	1.05e-08
rf	0.587	0.596	-0.009	0.025 [0.014, 0.034]	yes	2.88e-04	2.29e-04	8.31e-08	0.329	1.23e-04	3.16e-04	2.39e-04	1.00e-07	0.320	-8.56e-10
extratrees	0.610	0.593	0.017	0.022 [0.012, 0.031]	yes	2.89e-04	2.30e-04	8.38e-08	0.330	1.26e-04	3.07e-04	2.32e-04	9.44e-08	0.296	6.84e-20
mlp	0.586	0.598	-0.012	0.027 [0.015, 0.037]	yes	2.88e-04	2.29e-04	8.28e-08	0.329	1.26e-04	3.17e-04	2.39e-04	1.00e-07	0.320	1.27e-06
gp	0.565	0.587	-0.022	0.016 [0.010, 0.022]	yes	2.92e-04	2.30e-04	8.50e-08	0.332	1.25e-04	3.25e-04	2.45e-04	1.05e-07	0.340	-4.96e-06

best-by-test-R2 = **svr** (0.631) · linear 0.601 · persistence 0.571 · H=120s · 2105.6s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=5mSigma=45s · horizon=5m (H=300s) — metrics: LE + AI

LESigma=45s | horizon=5m | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.589	0.667	-0.078	0.037 [0.032, 0.042]	yes	4.15e-04	3.29e-04	1.73e-07	0.245	1.78e-04	4.88e-04	3.70e-04	2.38e-07	0.270	1.57e-15
ridge	0.589	0.667	-0.078	0.037 [0.032, 0.042]	yes	4.15e-04	3.29e-04	1.73e-07	0.245	1.78e-04	4.88e-04	3.70e-04	2.38e-07	0.270	2.67e-17
lasso	0.588	0.662	-0.074	0.032 [0.027, 0.036]	yes	4.19e-04	3.32e-04	1.75e-07	0.246	1.81e-04	4.89e-04	3.70e-04	2.39e-07	0.267	-2.56e-20
elasticnet	0.586	0.655	-0.069	0.025 [0.021, 0.029]	yes	4.23e-04	3.35e-04	1.79e-07	0.248	1.84e-04	4.90e-04	3.71e-04	2.41e-07	0.265	3.84e-20
huber	0.585	0.689	-0.104	0.059 [0.056, 0.061]	yes	4.02e-04	3.15e-04	1.61e-07	0.237	1.45e-04	4.91e-04	3.67e-04	2.41e-07	0.277	-4.13e-05
poly2	0.615	0.674	-0.059	0.044 [0.036, 0.053]	yes	4.11e-04	3.28e-04	1.69e-07	0.247	2.01e-04	4.73e-04	3.55e-04	2.23e-07	0.243	-1.24e-14
poly3	0.615	0.660	-0.045	0.030 [0.023, 0.037]	yes	4.20e-04	3.35e-04	1.76e-07	0.280	2.10e-04	4.73e-04	3.55e-04	2.24e-07	0.241	2.13e-13
knn	0.658	0.641	0.016	0.011 [0.002, 0.019]	yes	4.31e-04	3.45e-04	1.86e-07	0.244	2.26e-04	4.46e-04	3.30e-04	1.99e-07	0.206	-3.75e-06
svr ★	0.653	0.702	-0.048	0.071 [0.064, 0.079]	yes	3.93e-04	3.10e-04	1.55e-07	0.219	1.70e-04	4.49e-04	3.29e-04	2.01e-07	0.216	-3.82e-05
kernelridge	0.652	0.664	-0.012	0.033 [0.024, 0.043]	yes	4.18e-04	3.33e-04	1.75e-07	0.569	2.14e-04	4.50e-04	3.33e-04	2.02e-07	0.375	2.55e-07
gbm	0.562	0.549	0.013	-0.081 [-0.089, -0.072]	no	4.84e-04	3.93e-04	2.34e-07	0.298	2.71e-04	5.04e-04	3.79e-04	2.54e-07	0.241	1.30e-07
rf	0.696	0.641	0.054	0.011 [-0.004, 0.025]	no	4.31e-04	3.50e-04	1.86e-07	0.250	2.38e-04	4.20e-04	3.12e-04	1.77e-07	0.186	1.09e-07
extratrees	0.688	0.650	0.037	0.020 [0.008, 0.030]	yes	4.26e-04	3.45e-04	1.81e-07	0.247	2.32e-04	4.26e-04	3.15e-04	1.81e-07	0.186	-3.14e-20
mlp	0.667	0.665	0.002	0.035 [0.020, 0.046]	yes	4.17e-04	3.36e-04	1.74e-07	0.235	2.18e-04	4.40e-04	3.26e-04	1.93e-07	0.202	-1.36e-06
gp	0.641	0.651	-0.010	0.021 [0.005, 0.035]	yes	4.25e-04	3.37e-04	1.81e-07	0.240	2.15e-04	4.57e-04	3.37e-04	2.09e-07	0.214	-2.33e-06

best-by-test-R2 = **svr** (0.702) · linear 0.667 · persistence 0.630 · H=300s · 2117.3s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI Sigma=45s | horizon=5m | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.586	0.663	-0.077	0.032 [0.026, 0.038]	yes	4.18e-04	3.32e-04	1.75e-07	0.248	1.79e-04	4.90e-04	3.72e-04	2.40e-07	0.273	-8.30e-16
ridge	0.586	0.663	-0.077	0.032 [0.026, 0.038]	yes	4.18e-04	3.32e-04	1.75e-07	0.248	1.79e-04	4.90e-04	3.72e-04	2.40e-07	0.273	2.67e-17
lasso	0.585	0.657	-0.072	0.027 [0.021, 0.032]	yes	4.22e-04	3.34e-04	1.78e-07	0.249	1.83e-04	4.91e-04	3.72e-04	2.41e-07	0.269	4.27e-21
elasticnet	0.583	0.651	-0.068	0.020 [0.015, 0.026]	yes	4.26e-04	3.37e-04	1.81e-07	0.250	1.88e-04	4.92e-04	3.73e-04	2.42e-07	0.267	5.98e-20
huber	0.582	0.683	-0.101	0.053 [0.050, 0.056]	yes	4.06e-04	3.18e-04	1.64e-07	0.242	1.45e-04	4.93e-04	3.69e-04	2.43e-07	0.281	-4.14e-05
poly2	0.611	0.661	-0.050	0.031 [0.022, 0.039]	yes	4.19e-04	3.34e-04	1.76e-07	0.252	2.07e-04	4.75e-04	3.57e-04	2.26e-07	0.245	-2.41e-15
poly3	0.618	0.651	-0.034	0.021 [0.011, 0.030]	yes	4.25e-04	3.41e-04	1.81e-07	0.283	2.22e-04	4.71e-04	3.53e-04	2.22e-07	0.245	-2.91e-13
knn	0.655	0.639	0.016	0.009 [-0.002, 0.017]	no	4.33e-04	3.48e-04	1.87e-07	0.246	2.30e-04	4.48e-04	3.32e-04	2.00e-07	0.208	-5.21e-06
svr ★	0.651	0.699	-0.048	0.068 [0.060, 0.075]	yes	3.95e-04	3.13e-04	1.56e-07	0.221	1.75e-04	4.50e-04	3.31e-04	2.03e-07	0.217	-3.77e-05
kernelridge	0.651	0.658	-0.007	0.028 [0.017, 0.036]	yes	4.21e-04	3.37e-04	1.77e-07	0.375	2.21e-04	4.50e-04	3.33e-04	2.03e-07	0.373	6.84e-07
gbm	0.559	0.542	0.017	-0.088 [-0.096, -0.078]	no	4.88e-04	3.96e-04	2.38e-07	0.302	2.74e-04	5.06e-04	3.81e-04	2.56e-07	0.244	1.39e-07
rf	0.692	0.629	0.063	-0.001 [-0.018, 0.014]	no	4.39e-04	3.56e-04	1.93e-07	0.255	2.46e-04	4.23e-04	3.14e-04	1.79e-07	0.187	1.45e-07
extratrees	0.685	0.640	0.045	0.010 [-0.005, 0.022]	no	4.32e-04	3.51e-04	1.87e-07	0.252	2.41e-04	4.27e-04	3.17e-04	1.83e-07	0.187	-1.13e-20
mlp	0.664	0.658	0.006	0.027 [0.013, 0.040]	yes	4.21e-04	3.40e-04	1.78e-07	0.240	2.28e-04	4.42e-04	3.28e-04	1.95e-07	0.204	-4.80e-06
gp	0.639	0.648	-0.009	0.018 [0.005, 0.030]	yes	4.27e-04	3.39e-04	1.83e-07	0.251	2.20e-04	4.58e-04	3.38e-04	2.10e-07	0.219	-2.55e-06

best-by-test-R2 = **svr** (0.699) · linear 0.663 · persistence 0.630 · H=300s · 2111.2s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=10m Sigma=45s · horizon=10m (H=600s) — metrics: LE + AI

LE Sigma=45s | horizon=10m | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.594	0.677	-0.083	0.040 [0.034, 0.045]	yes	5.86e-04	4.67e-04	3.43e-07	0.207	2.80e-04	6.86e-04	5.16e-04	4.70e-07	0.229	-5.05e-16
ridge	0.594	0.677	-0.083	0.040 [0.034, 0.045]	yes	5.86e-04	4.67e-04	3.43e-07	0.207	2.80e-04	6.86e-04	5.16e-04	4.70e-07	0.229	6.87e-17
lasso	0.594	0.673	-0.079	0.036 [0.031, 0.041]	yes	5.89e-04	4.70e-04	3.47e-07	0.208	2.83e-04	6.86e-04	5.16e-04	4.70e-07	0.227	-1.54e-19
elasticnet	0.592	0.668	-0.076	0.031 [0.026, 0.036]	yes	5.94e-04	4.73e-04	3.52e-07	0.209	2.86e-04	6.87e-04	5.17e-04	4.72e-07	0.226	-2.92e-20
huber	0.589	0.704	-0.115	0.067 [0.064, 0.070]	yes	5.61e-04	4.42e-04	3.15e-07	0.196	2.27e-04	6.90e-04	5.10e-04	4.76e-07	0.235	-6.65e-05
poly2	0.622	0.680	-0.058	0.043 [0.034, 0.053]	yes	5.83e-04	4.68e-04	3.39e-07	0.212	3.13e-04	6.61e-04	4.93e-04	4.37e-07	0.208	1.27e-14
poly3	0.622	0.664	-0.042	0.027 [0.019, 0.034]	yes	5.97e-04	4.80e-04	3.57e-07	0.252	3.28e-04	6.61e-04	4.93e-04	4.37e-07	0.208	6.13e-13
knn	0.670	0.651	0.019	0.014 [0.004, 0.023]	yes	6.09e-04	4.94e-04	3.70e-07	0.210	3.48e-04	6.18e-04	4.52e-04	3.82e-07	0.171	-5.89e-06
svr ★	0.662	0.727	-0.065	0.090 [0.083, 0.098]	yes	5.38e-04	4.27e-04	2.90e-07	0.179	2.47e-04	6.25e-04	4.52e-04	3.91e-07	0.182	-6.87e-05
kernelridge	0.662	0.675	-0.013	0.038 [0.027, 0.048]	yes	5.88e-04	4.74e-04	3.46e-07	0.276	3.26e-04	6.26e-04	4.57e-04	3.91e-07	0.297	-4.85e-06
gbm	0.604	0.590	0.014	-0.047 [-0.054, -0.037]	no	6.59e-04	5.42e-04	4.35e-07	0.249	3.94e-04	6.77e-04	5.04e-04	4.58e-07	0.195	-1.61e-07
rf	0.714	0.626	0.088	-0.011 [-0.030, 0.006]	no	6.30e-04	5.17e-04	3.97e-07	0.225	3.81e-04	5.76e-04	4.21e-04	3.32e-07	0.150	2.11e-07
extratrees	0.700	0.647	0.053	0.010 [-0.005, 0.022]	no	6.12e-04	5.03e-04	3.75e-07	0.218	3.65e-04	5.89e-04	4.31e-04	3.47e-07	0.152	4.59e-21
mlp	0.678	0.664	0.013	0.028 [0.011, 0.040]	yes	5.97e-04	4.87e-04	3.56e-07	0.205	3.50e-04	6.11e-04	4.48e-04	3.73e-07	0.167	1.80e-06
gp	0.651	0.674	-0.024	0.037 [0.023, 0.049]	yes	5.88e-04	4.81e-04	3.46e-07	0.204	3.47e-04	6.36e-04	4.64e-04	4.04e-07	0.178	-4.67e-06

best-by-test-R2 = **svr** (0.727) · linear 0.677 · persistence 0.637 · H=600s · 2056.4s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI Sigma=45s | horizon=10m | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.591	0.672	-0.081	0.035 [0.028, 0.041]	yes	5.90e-04	4.71e-04	3.48e-07	0.210	2.83e-04	6.89e-04	5.19e-04	4.74e-07	0.232	-1.49e-16
ridge	0.591	0.672	-0.081	0.035 [0.028, 0.041]	yes	5.90e-04	4.71e-04	3.48e-07	0.210	2.83e-04	6.89e-04	5.19e-04	4.74e-07	0.232	6.89e-17
lasso	0.590	0.668	-0.078	0.031 [0.024, 0.037]	yes	5.94e-04	4.74e-04	3.53e-07	0.211	2.87e-04	6.89e-04	5.19e-04	4.74e-07	0.230	-2.05e-20
elasticnet	0.589	0.663	-0.074	0.026 [0.020, 0.032]	yes	5.98e-04	4.77e-04	3.58e-07	0.212	2.92e-04	6.90e-04	5.20e-04	4.76e-07	0.228	2.27e-20
huber	0.586	0.697	-0.112	0.060 [0.057, 0.064]	yes	5.67e-04	4.46e-04	3.21e-07	0.201	2.28e-04	6.93e-04	5.13e-04	4.80e-07	0.240	-6.68e-05
poly2	0.618	0.666	-0.048	0.029 [0.018, 0.038]	yes	5.96e-04	4.79e-04	3.55e-07	0.217	3.23e-04	6.65e-04	4.96e-04	4.43e-07	0.210	-7.53e-15
poly3	0.625	0.654	-0.029	0.017 [0.006, 0.028]	yes	6.06e-04	4.90e-04	3.67e-07	0.277	3.46e-04	6.59e-04	4.90e-04	4.34e-07	0.207	2.23e-13
knn	0.667	0.646	0.021	0.009 [-0.003, 0.019]	no	6.13e-04	4.99e-04	3.75e-07	0.214	3.55e-04	6.21e-04	4.55e-04	3.85e-07	0.172	-7.97e-06
svr ★	0.660	0.717	-0.057	0.080 [0.071, 0.088]	yes	5.48e-04	4.37e-04	3.01e-07	0.184	2.63e-04	6.28e-04	4.53e-04	3.94e-07	0.183	-6.79e-05
kernelridge	0.660	0.665	-0.005	0.028 [0.016, 0.039]	yes	5.96e-04	4.82e-04	3.56e-07	0.381	3.37e-04	6.27e-04	4.58e-04	3.93e-07	0.299	-4.68e-06
gbm	0.600	0.582	0.018	-0.055 [-0.063, -0.045]	no	6.66e-04	5.48e-04	4.44e-07	0.253	3.99e-04	6.80e-04	5.08e-04	4.63e-07	0.198	-1.82e-07
rf	0.710	0.609	0.101	-0.028 [-0.048, -0.009]	no	6.44e-04	5.29e-04	4.15e-07	0.231	3.96e-04	5.79e-04	4.24e-04	3.35e-07	0.151	2.10e-07
extratrees	0.697	0.634	0.063	-0.003 [-0.020, 0.012]	no	6.23e-04	5.13e-04	3.88e-07	0.224	3.80e-04	5.92e-04	4.33e-04	3.50e-07	0.154	7.29e-21
mlp	0.673	0.655	0.018	0.018 [0.001, 0.035]	yes	6.05e-04	4.96e-04	3.66e-07	0.210	3.65e-04	6.15e-04	4.51e-04	3.78e-07	0.169	3.22e-06
gp	0.649	0.660	-0.011	0.023 [0.009, 0.035]	yes	6.01e-04	4.92e-04	3.61e-07	0.211	3.61e-04	6.38e-04	4.65e-04	4.06e-07	0.180	-4.50e-06

best-by-test-R2 = **svr** (0.717) · linear 0.672 · persistence 0.637 · H=600s · 2079.4s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=15m Sigma=45s · horizon=15m (H=900s) — metrics: LE + AI

LE Sigma=45s | horizon=15m | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.583	0.672	-0.089	0.042 [0.036, 0.048]	yes	7.24e-04	5.80e-04	5.25e-07	0.197	3.63e-04	8.46e-04	6.33e-04	7.16e-07	0.218	-1.66e-15
ridge	0.583	0.672	-0.089	0.042 [0.036, 0.048]	yes	7.24e-04	5.80e-04	5.25e-07	0.197	3.63e-04	8.46e-04	6.33e-04	7.16e-07	0.218	-6.53e-17
lasso	0.583	0.669	-0.086	0.039 [0.033, 0.044]	yes	7.28e-04	5.83e-04	5.30e-07	0.198	3.66e-04	8.47e-04	6.34e-04	7.17e-07	0.217	1.03e-19
elasticnet	0.582	0.665	-0.083	0.035 [0.029, 0.040]	yes	7.32e-04	5.86e-04	5.36e-07	0.200	3.69e-04	8.48e-04	6.35e-04	7.18e-07	0.216	3.60e-20
huber	0.578	0.703	-0.126	0.073 [0.070, 0.077]	yes	6.89e-04	5.43e-04	4.75e-07	0.185	2.92e-04	8.52e-04	6.26e-04	7.26e-07	0.226	-8.80e-05
poly2	0.612	0.672	-0.060	0.042 [0.032, 0.052]	yes	7.25e-04	5.85e-04	5.25e-07	0.225	4.04e-04	8.17e-04	6.06e-04	6.67e-07	0.201	3.78e-14
poly3	0.613	0.654	-0.042	0.024 [0.016, 0.032]	yes	7.44e-04	6.01e-04	5.53e-07	0.249	4.24e-04	8.16e-04	6.05e-04	6.66e-07	0.198	8.69e-13
knn	0.661	0.634	0.027	0.005 [-0.006, 0.014]	no	7.65e-04	6.24e-04	5.85e-07	0.206	4.50e-04	7.63e-04	5.55e-04	5.82e-07	0.164	-6.15e-06
svr ★	0.652	0.722	-0.070	0.092 [0.084, 0.102]	yes	6.67e-04	5.31e-04	4.44e-07	0.171	3.21e-04	7.74e-04	5.52e-04	5.99e-07	0.176	-9.33e-05
kernelridge	0.653	0.662	-0.009	0.032 [0.020, 0.043]	yes	7.36e-04	5.98e-04	5.42e-07	0.319	4.28e-04	7.73e-04	5.61e-04	5.97e-07	0.294	-1.55e-06

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
gbm	0.607	0.598	0.010	-0.032 [-0.040, -0.023]	no	8.02e-04	6.64e-04	6.44e-07	0.235	4.90e-04	8.22e-04	6.08e-04	6.75e-07	0.183	1.45e-07
rf	0.707	0.600	0.107	-0.030 [-0.051, -0.010]	no	8.00e-04	6.61e-04	6.40e-07	0.225	5.03e-04	7.10e-04	5.15e-04	5.03e-07	0.142	2.35e-07
extratrees	0.693	0.631	0.063	0.001 [-0.016, 0.014]	no	7.69e-04	6.35e-04	5.91e-07	0.214	4.75e-04	7.26e-04	5.27e-04	5.28e-07	0.145	-4.47e-20
mlp	0.671	0.646	0.025	0.016 [-0.003, 0.031]	no	7.53e-04	6.18e-04	5.67e-07	0.203	4.57e-04	7.52e-04	5.46e-04	5.66e-07	0.160	-2.52e-06
gp	0.645	0.655	-0.010	0.025 [0.010, 0.038]	yes	7.43e-04	5.99e-04	5.52e-07	0.196	4.21e-04	7.81e-04	5.67e-04	6.10e-07	0.171	-7.37e-06

best-by-test-R2 = **svr** (0.722) · linear 0.672 · persistence 0.630 · H=900s · 2082.4s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI Sigma=45s | horizon=15m | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.580	0.667	-0.087	0.037 [0.030, 0.044]	yes	7.30e-04	5.85e-04	5.33e-07	0.200	3.67e-04	8.50e-04	6.37e-04	7.23e-07	0.221	-2.08e-15
ridge	0.580	0.667	-0.087	0.037 [0.030, 0.044]	yes	7.30e-04	5.85e-04	5.33e-07	0.200	3.67e-04	8.50e-04	6.37e-04	7.23e-07	0.221	-6.89e-17
lasso	0.579	0.663	-0.084	0.034 [0.026, 0.040]	yes	7.34e-04	5.88e-04	5.39e-07	0.201	3.72e-04	8.50e-04	6.38e-04	7.23e-07	0.220	5.67e-20
elasticnet	0.579	0.659	-0.081	0.030 [0.022, 0.036]	yes	7.38e-04	5.91e-04	5.45e-07	0.203	3.76e-04	8.51e-04	6.38e-04	7.25e-07	0.218	1.11e-19
huber	0.574	0.697	-0.123	0.067 [0.063, 0.070]	yes	6.97e-04	5.49e-04	4.86e-07	0.190	2.93e-04	8.56e-04	6.30e-04	7.33e-07	0.230	-8.84e-05
poly2	0.607	0.656	-0.050	0.027 [0.015, 0.037]	yes	7.41e-04	5.98e-04	5.50e-07	0.230	4.16e-04	8.22e-04	6.10e-04	6.76e-07	0.202	-1.83e-14
poly3	0.616	0.644	-0.029	0.014 [0.002, 0.025]	yes	7.55e-04	6.14e-04	5.69e-07	0.275	4.45e-04	8.13e-04	6.01e-04	6.61e-07	0.198	3.23e-13
knn	0.659	0.629	0.030	-0.001 [-0.015, 0.009]	no	7.70e-04	6.30e-04	5.94e-07	0.209	4.61e-04	7.66e-04	5.57e-04	5.87e-07	0.165	-8.41e-06
svr ★	0.650	0.708	-0.058	0.078 [0.069, 0.087]	yes	6.84e-04	5.46e-04	4.67e-07	0.178	3.43e-04	7.76e-04	5.54e-04	6.02e-07	0.176	-9.25e-05
kernelridge	0.652	0.650	0.002	0.020 [0.007, 0.031]	yes	7.49e-04	6.08e-04	5.60e-07	0.446	4.39e-04	7.74e-04	5.62e-04	5.99e-07	0.279	-1.25e-06
gbm	0.603	0.593	0.010	-0.037 [-0.045, -0.027]	no	8.07e-04	6.67e-04	6.51e-07	0.237	4.92e-04	8.26e-04	6.12e-04	6.82e-07	0.186	1.47e-07
rf	0.704	0.580	0.124	-0.050 [-0.074, -0.029]	no	8.20e-04	6.78e-04	6.73e-07	0.232	5.24e-04	7.13e-04	5.18e-04	5.09e-07	0.144	3.23e-07
extratrees	0.691	0.614	0.077	-0.016 [-0.035, 0.000]	no	7.86e-04	6.52e-04	6.18e-07	0.221	4.97e-04	7.29e-04	5.29e-04	5.32e-07	0.146	-2.73e-20
mlp	0.667	0.651	0.016	0.021 [0.003, 0.036]	yes	7.48e-04	6.14e-04	5.59e-07	0.202	4.54e-04	7.57e-04	5.51e-04	5.73e-07	0.161	-2.37e-06
gp	0.644	0.652	-0.007	0.022 [0.012, 0.031]	yes	7.47e-04	6.06e-04	5.58e-07	0.201	4.38e-04	7.82e-04	5.68e-04	6.12e-07	0.171	-6.14e-06

best-by-test-R2 = **svr** (0.708) · linear 0.667 · persistence 0.630 · H=900s · 2115.5s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=30m Sigma=45s · horizon=30m (H=1800s) — metrics: LE + AI

LE Sigma=45s | horizon=30m | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.544	0.647	-0.103	0.048 [0.041, 0.055]	yes	1.05e-03	8.46e-04	1.09e-06	0.195	5.64e-04	1.22e-03	9.02e-04	1.48e-06	0.213	-2.73e-15
ridge	0.544	0.647	-0.103	0.048 [0.041, 0.055]	yes	1.05e-03	8.46e-04	1.09e-06	0.195	5.64e-04	1.22e-03	9.02e-04	1.48e-06	0.213	3.27e-17
lasso	0.544	0.644	-0.100	0.046 [0.038, 0.052]	yes	1.05e-03	8.49e-04	1.10e-06	0.196	5.67e-04	1.22e-03	9.02e-04	1.48e-06	0.212	5.19e-19
elasticnet	0.544	0.641	-0.098	0.043 [0.036, 0.049]	yes	1.05e-03	8.53e-04	1.11e-06	0.197	5.70e-04	1.22e-03	9.03e-04	1.48e-06	0.212	4.06e-19
huber	0.537	0.689	-0.152	0.090 [0.086, 0.094]	yes	9.82e-04	7.78e-04	9.65e-07	0.179	4.43e-04	1.23e-03	8.88e-04	1.50e-06	0.222	-1.42e-04
poly2	0.573	0.637	-0.064	0.038 [0.027, 0.050]	yes	1.06e-03	8.63e-04	1.13e-06	0.223	6.19e-04	1.18e-03	8.65e-04	1.38e-06	0.198	2.43e-14
poly3	0.575	0.616	-0.041	0.017 [0.008, 0.026]	yes	1.09e-03	8.88e-04	1.19e-06	0.247	6.51e-04	1.17e-03	8.62e-04	1.38e-06	0.196	1.66e-12
knn	0.626	0.597	0.029	-0.001 [-0.014, 0.009]	no	1.12e-03	9.21e-04	1.25e-06	0.209	6.84e-04	1.10e-03	7.91e-04	1.21e-06	0.166	-1.42e-05
svr ★	0.614	0.691	-0.077	0.093 [0.081, 0.104]	yes	9.78e-04	7.85e-04	9.56e-07	0.174	4.96e-04	1.12e-03	7.83e-04	1.25e-06	0.179	-1.52e-04
kernelridge	0.617	0.626	-0.009	0.027 [0.014, 0.039]	yes	1.08e-03	8.86e-04	1.16e-06	0.291	6.56e-04	1.11e-03	8.01e-04	1.24e-06	0.245	3.83e-07
gbm	0.588	0.576	0.012	-0.023 [-0.033, -0.014]	no	1.15e-03	9.57e-04	1.31e-06	0.228	7.26e-04	1.16e-03	8.44e-04	1.34e-06	0.177	-4.76e-07
rf	0.690	0.518	0.172	-0.080 [-0.108, -0.054]	no	1.22e-03	1.01e-03	1.49e-06	0.239	7.98e-04	1.00e-03	7.16e-04	1.00e-06	0.134	6.99e-07
extratrees	0.662	0.573	0.089	-0.025 [-0.046, -0.008]	no	1.15e-03	9.61e-04	1.32e-06	0.223	7.42e-04	1.05e-03	7.50e-04	1.09e-06	0.145	1.69e-21
mlp	0.634	0.594	0.040	-0.004 [-0.026, 0.015]	no	1.12e-03	9.38e-04	1.26e-06	0.213	7.32e-04	1.09e-03	7.83e-04	1.18e-06	0.162	3.88e-06
gp	0.606	0.622	-0.016	0.024 [0.011, 0.034]	yes	1.08e-03	8.92e-04	1.17e-06	0.201	6.68e-04	1.13e-03	8.12e-04	1.28e-06	0.171	8.11e-07

best-by-test-R2 = **svr** (0.691) · linear 0.647 · persistence 0.599 · H=1800s · 2044.7s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI Sigma=45s | horizon=30m | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.539	0.641	-0.101	0.042 [0.033, 0.050]	yes	1.06e-03	8.53e-04	1.11e-06	0.198	5.70e-04	1.22e-03	9.08e-04	1.49e-06	0.216	-9.42e-16
ridge	0.539	0.641	-0.101	0.042 [0.033, 0.050]	yes	1.06e-03	8.53e-04	1.11e-06	0.198	5.70e-04	1.22e-03	9.08e-04	1.49e-06	0.216	3.35e-17
lasso	0.539	0.638	-0.099	0.039 [0.030, 0.047]	yes	1.06e-03	8.57e-04	1.12e-06	0.199	5.75e-04	1.22e-03	9.08e-04	1.49e-06	0.215	4.53e-19
elasticnet	0.539	0.635	-0.096	0.036 [0.027, 0.044]	yes	1.06e-03	8.60e-04	1.13e-06	0.200	5.79e-04	1.22e-03	9.09e-04	1.49e-06	0.215	5.01e-19
huber ★	0.532	0.682	-0.150	0.083 [0.079, 0.088]	yes	9.93e-04	7.87e-04	9.87e-07	0.183	4.46e-04	1.23e-03	8.95e-04	1.52e-06	0.227	-1.42e-04
poly2	0.567	0.620	-0.054	0.022 [0.009, 0.033]	yes	1.08e-03	8.81e-04	1.18e-06	0.228	6.35e-04	1.19e-03	8.72e-04	1.40e-06	0.200	5.19e-14
poly3	0.578	0.606	-0.028	0.007 [-0.007, 0.020]	no	1.11e-03	9.08e-04	1.22e-06	0.270	6.81e-04	1.17e-03	8.56e-04	1.37e-06	0.213	-4.03e-13
knn	0.623	0.590	0.033	-0.008 [-0.025, 0.005]	no	1.13e-03	9.33e-04	1.27e-06	0.212	7.02e-04	1.10e-03	7.94e-04	1.22e-06	0.167	-1.70e-05
svr	0.613	0.679	-0.066	0.080 [0.068, 0.090]	yes	9.98e-04	8.04e-04	9.96e-07	0.179	5.18e-04	1.12e-03	7.85e-04	1.25e-06	0.179	-1.50e-04
kernelridge	0.616	0.610	0.006	0.012 [-0.005, 0.025]	no	1.10e-03	9.02e-04	1.21e-06	0.441	6.71e-04	1.12e-03	8.02e-04	1.24e-06	0.291	8.82e-07
gbm	0.585	0.567	0.017	-0.031 [-0.043, -0.021]	no	1.16e-03	9.67e-04	1.34e-06	0.231	7.38e-04	1.16e-03	8.49e-04	1.35e-06	0.179	-4.93e-07
rf	0.668	0.514	0.154	-0.084 [-0.114, -0.059]	no	1.23e-03	1.02e-03	1.51e-06	0.241	8.13e-04	1.04e-03	7.45e-04	1.07e-06	0.145	2.72e-07
extratrees	0.661	0.551	0.111	-0.048 [-0.071, -0.028]	no	1.18e-03	9.89e-04	1.39e-06	0.232	7.78e-04	1.05e-03	7.52e-04	1.10e-06	0.145	6.99e-20
mlp	0.631	0.571	0.060	-0.028 [-0.055, -0.004]	no	1.15e-03	9.62e-04	1.33e-06	0.220	7.59e-04	1.09e-03	7.89e-04	1.20e-06	0.163	2.98e-06
gp	0.605	0.613	-0.007	0.014 [-0.001, 0.027]	no	1.10e-03	9.05e-04	1.20e-06	0.205	6.84e-04	1.13e-03	8.15e-04	1.28e-06	0.172	2.65e-06

best-by-test-R2 = **huber** (0.682) · linear 0.641 · persistence 0.599 · H=1800s · 2049.5s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=1h Sigma=45s · horizon=1h (H=3600s) — metrics: LE + AI

LE Sigma=45s | horizon=1h | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.475	0.594	-0.118	0.045 [0.036, 0.053]	yes	1.55e-03	1.26e-03	2.40e-06	0.206	8.80e-04	1.77e-03	1.30e-03	3.13e-06	0.221	-1.08e-15
ridge	0.475	0.594	-0.118	0.045 [0.036, 0.053]	yes	1.55e-03	1.26e-03	2.40e-06	0.206	8.80e-04	1.77e-03	1.30e-03	3.13e-06	0.221	1.21e-16
lasso	0.475	0.592	-0.116	0.043 [0.034, 0.050]	yes	1.55e-03	1.27e-03	2.41e-06	0.207	8.83e-04	1.77e-03	1.30e-03	3.13e-06	0.221	2.18e-19
elasticnet	0.475	0.589	-0.114	0.040 [0.032, 0.048]	yes	1.56e-03	1.27e-03	2.43e-06	0.208	8.86e-04	1.77e-03	1.30e-03	3.14e-06	0.221	4.99e-19

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
huber ★	0.464	0.656	-0.192	0.107 [0.102, 0.113]	yes	1.43e-03	1.13e-03	2.03e-06	0.184	6.64e-04	1.79e-03	1.28e-03	3.20e-06	0.234	-2.37e-04
poly2	0.504	0.568	-0.064	0.019 [0.006, 0.033]	yes	1.60e-03	1.30e-03	2.55e-06	0.234	9.61e-04	1.72e-03	1.26e-03	2.96e-06	0.210	2.36e-14
poly3	0.507	0.539	-0.031	-0.010 [-0.023, 0.002]	no	1.65e-03	1.35e-03	2.72e-06	0.249	1.02e-03	1.72e-03	1.25e-03	2.94e-06	0.207	1.39e-12
knn	0.562	0.517	0.045	-0.032 [-0.049, -0.020]	no	1.69e-03	1.40e-03	2.85e-06	0.227	1.07e-03	1.62e-03	1.16e-03	2.62e-06	0.177	-1.58e-05
svr	0.544	0.649	-0.105	0.100 [0.086, 0.115]	yes	1.44e-03	1.15e-03	2.07e-06	0.181	7.24e-04	1.65e-03	1.13e-03	2.72e-06	0.197	-2.76e-04
kernelridge	0.553	0.540	0.013	-0.009 [-0.024, 0.006]	no	1.65e-03	1.36e-03	2.72e-06	0.286	1.03e-03	1.63e-03	1.17e-03	2.67e-06	0.215	-8.22e-06
gbm	0.538	0.519	0.018	-0.030 [-0.043, -0.019]	no	1.68e-03	1.41e-03	2.84e-06	0.235	1.09e-03	1.66e-03	1.20e-03	2.76e-06	0.183	6.97e-07
rf	0.630	0.434	0.196	-0.115 [-0.145, -0.085]	no	1.83e-03	1.52e-03	3.34e-06	0.253	1.21e-03	1.49e-03	1.06e-03	2.21e-06	0.145	6.97e-07
extratrees	0.606	0.488	0.118	-0.061 [-0.084, -0.042]	no	1.74e-03	1.46e-03	3.02e-06	0.241	1.14e-03	1.53e-03	1.09e-03	2.35e-06	0.153	-1.45e-20
mlp	0.571	0.528	0.043	-0.021 [-0.046, 0.001]	no	1.67e-03	1.40e-03	2.79e-06	0.224	1.09e-03	1.60e-03	1.14e-03	2.56e-06	0.173	2.88e-06
gp	0.541	0.526	0.015	-0.023 [-0.038, -0.009]	no	1.67e-03	1.38e-03	2.80e-06	0.225	1.05e-03	1.65e-03	1.18e-03	2.74e-06	0.183	-1.65e-05

best-by-test-R2 = **huber** (0.656) · linear 0.594 · persistence 0.549 · H=3600s · 1998.6s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI Sigma=45s | horizon=1h | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.471	0.586	-0.115	0.038 [0.027, 0.047]	yes	1.56e-03	1.27e-03	2.44e-06	0.210	8.92e-04	1.78e-03	1.31e-03	3.16e-06	0.224	3.62e-17
ridge	0.471	0.586	-0.115	0.038 [0.027, 0.047]	yes	1.56e-03	1.27e-03	2.44e-06	0.210	8.92e-04	1.78e-03	1.31e-03	3.16e-06	0.224	1.23e-16
lasso	0.471	0.584	-0.113	0.035 [0.024, 0.044]	yes	1.57e-03	1.28e-03	2.46e-06	0.210	8.96e-04	1.78e-03	1.31e-03	3.16e-06	0.223	2.18e-19
elasticnet	0.471	0.582	-0.111	0.033 [0.022, 0.042]	yes	1.57e-03	1.28e-03	2.47e-06	0.211	9.01e-04	1.78e-03	1.31e-03	3.16e-06	0.223	4.07e-19
huber ★	0.460	0.648	-0.188	0.099 [0.094, 0.105]	yes	1.44e-03	1.14e-03	2.08e-06	0.188	6.72e-04	1.80e-03	1.29e-03	3.23e-06	0.238	-2.37e-04
poly2	0.497	0.549	-0.052	0.001 [-0.017, 0.016]	no	1.63e-03	1.33e-03	2.66e-06	0.239	9.83e-04	1.73e-03	1.27e-03	3.00e-06	0.212	-1.00e-14
poly3	0.512	0.529	-0.017	-0.020 [-0.038, -0.003]	no	1.67e-03	1.38e-03	2.78e-06	0.269	1.06e-03	1.71e-03	1.24e-03	2.91e-06	0.221	6.46e-13
knn	0.562	0.504	0.058	-0.045 [-0.066, -0.028]	no	1.71e-03	1.42e-03	2.93e-06	0.231	1.10e-03	1.62e-03	1.16e-03	2.62e-06	0.177	-1.91e-05
svr	0.543	0.641	-0.098	0.092 [0.079, 0.107]	yes	1.46e-03	1.17e-03	2.12e-06	0.185	7.55e-04	1.65e-03	1.14e-03	2.73e-06	0.196	-2.75e-04
kernelridge	0.553	0.529	0.024	-0.020 [-0.038, -0.004]	no	1.67e-03	1.39e-03	2.78e-06	0.458	1.06e-03	1.63e-03	1.17e-03	2.67e-06	0.244	-7.96e-06
gbm	0.536	0.517	0.019	-0.032 [-0.047, -0.020]	no	1.69e-03	1.42e-03	2.85e-06	0.237	1.10e-03	1.67e-03	1.21e-03	2.77e-06	0.184	7.57e-07
rf	0.615	0.407	0.208	-0.141 [-0.180, -0.109]	no	1.87e-03	1.56e-03	3.50e-06	0.261	1.26e-03	1.52e-03	1.08e-03	2.30e-06	0.153	1.92e-07
extratrees	0.608	0.459	0.148	-0.089 [-0.120, -0.064]	no	1.79e-03	1.50e-03	3.19e-06	0.250	1.20e-03	1.53e-03	1.09e-03	2.34e-06	0.153	1.55e-19
mlp	0.570	0.464	0.106	-0.085 [-0.117, -0.057]	no	1.78e-03	1.49e-03	3.16e-06	0.244	1.20e-03	1.60e-03	1.15e-03	2.57e-06	0.172	4.96e-06
gp	0.544	0.521	0.024	-0.028 [-0.046, -0.014]	no	1.68e-03	1.40e-03	2.83e-06	0.227	1.07e-03	1.65e-03	1.18e-03	2.72e-06	0.182	-9.61e-06

best-by-test-R2 = **huber** (0.648) · linear 0.586 · persistence 0.549 · H=3600s · 2026.4s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

Sigma=60s cov-duration 60s — 10 horizon(s) finished

horizon=nextwin Sigma=60s · horizon=nextwin (H=60s) — metrics: LE + AI

LE Sigma=60s | horizon=nextwin | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.542	0.610	-0.068	0.023 [0.021, 0.025]	yes	2.54e-04	1.98e-04	6.44e-08	0.442	5.92e-05	3.03e-04	2.29e-04	9.19e-08	0.518	9.45e-17
ridge	0.542	0.610	-0.068	0.023 [0.021, 0.025]	yes	2.54e-04	1.98e-04	6.44e-08	0.442	5.92e-05	3.03e-04	2.29e-04	9.19e-08	0.518	-5.46e-18
lasso	0.540	0.602	-0.063	0.015 [0.014, 0.017]	yes	2.56e-04	2.00e-04	6.57e-08	0.423	6.22e-05	3.04e-04	2.29e-04	9.24e-08	0.493	4.29e-20
elasticnet	0.533	0.592	-0.059	0.004 [0.003, 0.006]	yes	2.60e-04	2.03e-04	6.75e-08	0.420	6.56e-05	3.06e-04	2.31e-04	9.38e-08	0.486	8.43e-21
huber	0.540	0.619	-0.078	0.032 [0.030, 0.033]	yes	2.51e-04	1.96e-04	6.30e-08	0.472	4.45e-05	3.04e-04	2.28e-04	9.22e-08	0.555	-1.57e-05
poly2	0.558	0.619	-0.061	0.031 [0.028, 0.035]	yes	2.51e-04	1.96e-04	6.30e-08	0.359	7.15e-05	2.98e-04	2.24e-04	8.88e-08	0.425	1.68e-15
poly3	0.558	0.614	-0.056	0.027 [0.024, 0.029]	yes	2.53e-04	1.98e-04	6.38e-08	0.437	7.39e-05	2.98e-04	2.25e-04	8.87e-08	0.434	2.39e-13
knn	0.582	0.595	-0.013	0.008 [0.002, 0.015]	yes	2.59e-04	2.03e-04	6.69e-08	0.361	8.36e-05	2.90e-04	2.17e-04	8.39e-08	0.392	-3.65e-06
svr ★	0.579	0.626	-0.046	0.038 [0.031, 0.045]	yes	2.49e-04	1.94e-04	6.19e-08	0.375	6.55e-05	2.91e-04	2.17e-04	8.44e-08	0.416	-1.55e-05
kernelridge	0.577	0.606	-0.028	0.018 [0.004, 0.030]	yes	2.55e-04	1.98e-04	6.52e-08	0.952	7.72e-05	2.91e-04	2.18e-04	8.49e-08	0.802	-1.77e-06
gbm	0.474	0.478	-0.004	-0.109 [-0.117, -0.098]	no	2.94e-04	2.31e-04	8.62e-08	0.382	1.21e-04	3.25e-04	2.45e-04	1.06e-07	0.372	-2.12e-07
rf	0.592	0.606	-0.014	0.019 [0.013, 0.024]	yes	2.55e-04	2.01e-04	6.51e-08	0.353	8.68e-05	2.86e-04	2.15e-04	8.18e-08	0.376	6.51e-08
extratrees	0.611	0.606	0.005	0.019 [0.014, 0.023]	yes	2.55e-04	2.01e-04	6.51e-08	0.350	8.75e-05	2.79e-04	2.09e-04	7.81e-08	0.348	8.43e-21
mlp	0.589	0.613	-0.024	0.026 [0.020, 0.031]	yes	2.53e-04	1.99e-04	6.39e-08	0.358	8.38e-05	2.87e-04	2.16e-04	8.25e-08	0.379	2.12e-06
gp	0.570	0.614	-0.044	0.026 [0.019, 0.033]	yes	2.53e-04	1.97e-04	6.39e-08	0.362	7.79e-05	2.94e-04	2.20e-04	8.63e-08	0.417	-5.77e-06

best-by-test-R2 = **svr** (0.626) · linear 0.610 · persistence 0.587 · H=60s · 1801.4s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI Sigma=60s | horizon=nextwin | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.540	0.607	-0.067	0.020 [0.018, 0.022]	yes	2.55e-04	1.99e-04	6.49e-08	0.451	5.76e-05	3.04e-04	2.30e-04	9.24e-08	0.523	-3.25e-16
ridge	0.540	0.607	-0.067	0.020 [0.018, 0.022]	yes	2.55e-04	1.99e-04	6.49e-08	0.451	5.76e-05	3.04e-04	2.30e-04	9.24e-08	0.523	-5.10e-18
lasso	0.538	0.600	-0.062	0.013 [0.011, 0.015]	yes	2.57e-04	2.01e-04	6.61e-08	0.427	6.23e-05	3.05e-04	2.30e-04	9.28e-08	0.495	2.81e-21
elasticnet	0.533	0.590	-0.058	0.003 [0.002, 0.005]	yes	2.60e-04	2.03e-04	6.77e-08	0.417	6.72e-05	3.06e-04	2.31e-04	9.38e-08	0.482	2.81e-21
huber	0.538	0.615	-0.077	0.028 [0.026, 0.029]	yes	2.52e-04	1.96e-04	6.36e-08	0.479	4.31e-05	3.05e-04	2.29e-04	9.27e-08	0.562	-1.58e-05
poly2	0.554	0.612	-0.058	0.024 [0.021, 0.028]	yes	2.53e-04	1.98e-04	6.42e-08	0.363	7.20e-05	2.99e-04	2.25e-04	8.95e-08	0.421	-4.68e-16
poly3	0.557	0.608	-0.051	0.020 [0.016, 0.024]	yes	2.55e-04	1.99e-04	6.48e-08	0.439	7.63e-05	2.98e-04	2.25e-04	8.89e-08	0.427	-1.20e-14
knn	0.579	0.594	-0.015	0.007 [0.002, 0.012]	yes	2.59e-04	2.03e-04	6.70e-08	0.362	8.37e-05	2.91e-04	2.18e-04	8.45e-08	0.394	-4.26e-06
svr ★	0.577	0.624	-0.047	0.037 [0.032, 0.042]	yes	2.49e-04	1.94e-04	6.21e-08	0.371	6.48e-05	2.91e-04	2.18e-04	8.49e-08	0.415	-1.48e-05
kernelridge	0.575	0.606	-0.031	0.019 [0.010, 0.027]	yes	2.55e-04	1.99e-04	6.51e-08	0.700	7.76e-05	2.92e-04	2.18e-04	8.53e-08	0.871	-1.44e-06
gbm	0.473	0.482	-0.009	-0.105 [-0.113, -0.095]	no	2.92e-04	2.30e-04	8.55e-08	0.380	1.19e-04	3.25e-04	2.45e-04	1.06e-07	0.373	-2.15e-07
rf	0.589	0.603	-0.014	0.016 [0.009, 0.021]	yes	2.56e-04	2.02e-04	6.57e-08	0.353	8.70e-05	2.87e-04	2.16e-04	8.25e-08	0.376	6.02e-08
extratrees	0.608	0.602	0.007	0.015 [0.009, 0.019]	yes	2.57e-04	2.02e-04	6.58e-08	0.351	8.89e-05	2.80e-04	2.10e-04	7.86e-08	0.348	-3.78e-21
mlp	0.586	0.620	-0.034	0.033 [0.027, 0.037]	yes	2.51e-04	1.97e-04	6.28e-08	0.360	7.65e-05	2.88e-04	2.17e-04	8.32e-08	0.392	-6.26e-07
gp	0.567	0.610	-0.042	0.023 [0.014, 0.030]	yes	2.54e-04	1.98e-04	6.45e-08	0.366	7.61e-05	2.95e-04	2.20e-04	8.68e-08	0.415	-5.71e-06

best-by-test-R2 = **svr** (0.624) · linear 0.607 · persistence 0.587 · H=60s · 1840.1s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=30s

Sigma=60s · horizon=30s (H=30s) — metrics: LE + AI

LE

Sigma=60s | horizon=30s | metric=LE

N/A — horizon < cov window (H=30s ≤ Sigma=60s). horizon 30s < cov 60s (cover 0.0%)

AI

Sigma=60s | horizon=30s | metric=AI

N/A — horizon < cov window (H=30s ≤ Sigma=60s). horizon 30s < cov 60s (cover 0.0%)

horizon=45s

Sigma=60s · horizon=45s (H=45s) — metrics: LE + AI

LE

Sigma=60s | horizon=45s | metric=LE

N/A — horizon < cov window (H=45s ≤ Sigma=60s). horizon 45s < cov 60s (cover 0.0%)

AI

Sigma=60s | horizon=45s | metric=AI

N/A — horizon < cov window (H=45s ≤ Sigma=60s). horizon 45s < cov 60s (cover 0.0%)

horizon=60s

Sigma=60s · horizon=60s (H=60s) — metrics: LE + AI

LE

Sigma=60s | horizon=60s | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.544	0.612	-0.067	0.024 [0.022, 0.026]	yes	2.46e-04	1.94e-04	6.04e-08	0.429	5.96e-05	2.92e-04	2.24e-04	8.50e-08	0.501	-1.41e-16
ridge	0.544	0.612	-0.067	0.024 [0.022, 0.026]	yes	2.46e-04	1.94e-04	6.04e-08	0.428	5.97e-05	2.92e-04	2.24e-04	8.50e-08	0.501	1.67e-17
lasso	0.542	0.603	-0.062	0.016 [0.014, 0.017]	yes	2.48e-04	1.96e-04	6.17e-08	0.408	6.26e-05	2.93e-04	2.24e-04	8.56e-08	0.474	-5.89e-21
elasticnet	0.534	0.592	-0.058	0.004 [0.003, 0.006]	yes	2.52e-04	1.98e-04	6.34e-08	0.406	6.59e-05	2.95e-04	2.25e-04	8.70e-08	0.468	3.31e-21
huber	0.543	0.620	-0.077	0.033 [0.031, 0.034]	yes	2.43e-04	1.91e-04	5.90e-08	0.451	4.62e-05	2.92e-04	2.23e-04	8.53e-08	0.531	-1.46e-05
poly2	0.561	0.620	-0.059	0.033 [0.029, 0.036]	yes	2.43e-04	1.92e-04	5.90e-08	0.347	7.21e-05	2.86e-04	2.19e-04	8.20e-08	0.410	8.59e-15
poly3	0.562	0.615	-0.054	0.028 [0.025, 0.031]	yes	2.45e-04	1.93e-04	5.98e-08	0.431	7.46e-05	2.86e-04	2.19e-04	8.19e-08	0.418	2.41e-13
knn	0.586	0.594	-0.008	0.007 [-0.000, 0.013]	no	2.51e-04	1.99e-04	6.31e-08	0.352	8.43e-05	2.78e-04	2.12e-04	7.72e-08	0.376	-2.41e-06
svr ★	0.583	0.628	-0.045	0.040 [0.033, 0.047]	yes	2.41e-04	1.89e-04	5.78e-08	0.356	6.62e-05	2.79e-04	2.12e-04	7.78e-08	0.411	-1.41e-05
kernelridge	0.581	0.610	-0.028	0.022 [0.011, 0.032]	yes	2.46e-04	1.94e-04	6.07e-08	0.500	7.68e-05	2.80e-04	2.12e-04	7.81e-08	0.696	1.60e-07
gbm	0.477	0.482	-0.005	-0.106 [-0.114, -0.096]	no	2.84e-04	2.25e-04	8.06e-08	0.371	1.19e-04	3.12e-04	2.38e-04	9.76e-08	0.359	-2.74e-07
rf	0.596	0.608	-0.012	0.020 [0.014, 0.026]	yes	2.47e-04	1.96e-04	6.09e-08	0.346	8.60e-05	2.75e-04	2.10e-04	7.54e-08	0.363	2.62e-08
extratrees	0.615	0.608	0.007	0.020 [0.015, 0.025]	yes	2.47e-04	1.96e-04	6.10e-08	0.342	8.69e-05	2.68e-04	2.04e-04	7.18e-08	0.337	2.21e-19
mlp	0.593	0.614	-0.021	0.026 [0.018, 0.032]	yes	2.45e-04	1.95e-04	6.01e-08	0.357	8.42e-05	2.76e-04	2.11e-04	7.60e-08	0.371	2.63e-06
gp	0.572	0.605	-0.034	0.017 [0.007, 0.027]	yes	2.48e-04	1.95e-04	6.14e-08	0.346	8.20e-05	2.83e-04	2.15e-04	8.00e-08	0.394	-2.86e-06

best-by-test-R2 = svr (0.628) · linear 0.612 · persistence 0.588 · H=60s · 1771.6s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI

Sigma=60s | horizon=60s | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.542	0.608	-0.066	0.021 [0.018, 0.023]	yes	2.47e-04	1.95e-04	6.09e-08	0.435	5.84e-05	2.92e-04	2.24e-04	8.55e-08	0.506	6.29e-18
ridge	0.542	0.608	-0.066	0.021 [0.018, 0.023]	yes	2.47e-04	1.95e-04	6.09e-08	0.435	5.85e-05	2.92e-04	2.24e-04	8.55e-08	0.506	1.75e-17
lasso	0.540	0.601	-0.061	0.013 [0.011, 0.015]	yes	2.49e-04	1.96e-04	6.20e-08	0.411	6.31e-05	2.93e-04	2.24e-04	8.59e-08	0.477	6.63e-21
elasticnet	0.534	0.591	-0.056	0.003 [0.001, 0.005]	yes	2.52e-04	1.99e-04	6.36e-08	0.402	6.81e-05	2.95e-04	2.25e-04	8.69e-08	0.465	9.94e-21
huber	0.541	0.616	-0.076	0.029 [0.027, 0.030]	yes	2.44e-04	1.92e-04	5.96e-08	0.457	4.50e-05	2.93e-04	2.24e-04	8.58e-08	0.537	-1.46e-05
poly2	0.557	0.613	-0.055	0.025 [0.021, 0.029]	yes	2.45e-04	1.94e-04	6.02e-08	0.350	7.34e-05	2.87e-04	2.20e-04	8.26e-08	0.406	-9.05e-16
poly3	0.561	0.609	-0.048	0.021 [0.016, 0.025]	yes	2.47e-04	1.95e-04	6.08e-08	0.432	7.77e-05	2.86e-04	2.19e-04	8.20e-08	0.411	1.45e-14
knn	0.583	0.594	-0.011	0.006 [0.001, 0.012]	yes	2.51e-04	1.99e-04	6.31e-08	0.353	8.49e-05	2.79e-04	2.13e-04	7.79e-08	0.378	-3.03e-06
svr ★	0.580	0.625	-0.045	0.038 [0.031, 0.044]	yes	2.41e-04	1.90e-04	5.83e-08	0.356	6.62e-05	2.80e-04	2.12e-04	7.84e-08	0.411	-1.40e-05
kernelridge	0.578	0.607	-0.029	0.020 [0.010, 0.028]	yes	2.47e-04	1.95e-04	6.10e-08	0.548	7.84e-05	2.81e-04	2.13e-04	7.87e-08	0.784	2.33e-07
gbm	0.475	0.482	-0.006	-0.106 [-0.114, -0.096]	no	2.84e-04	2.25e-04	8.06e-08	0.371	1.18e-04	3.13e-04	2.39e-04	9.79e-08	0.361	-2.74e-07
rf	0.593	0.604	-0.011	0.016 [0.009, 0.022]	yes	2.48e-04	1.97e-04	6.16e-08	0.345	8.68e-05	2.76e-04	2.11e-04	7.60e-08	0.364	3.44e-08
extratrees	0.612	0.602	0.010	0.014 [0.008, 0.019]	yes	2.49e-04	1.98e-04	6.19e-08	0.343	8.93e-05	2.69e-04	2.05e-04	7.24e-08	0.338	2.11e-19
mlp	0.589	0.615	-0.025	0.027 [0.020, 0.032]	yes	2.45e-04	1.95e-04	5.99e-08	0.347	8.30e-05	2.77e-04	2.11e-04	7.66e-08	0.363	3.17e-06
gp	0.569	0.602	-0.033	0.015 [0.006, 0.023]	yes	2.49e-04	1.96e-04	6.18e-08	0.353	8.30e-05	2.84e-04	2.15e-04	8.04e-08	0.398	-2.66e-06

best-by-test-R2 = svr (0.625) · linear 0.608 · persistence 0.588 · H=60s · 1814.1s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=2m

Sigma=60s · horizon=2m (H=120s) — metrics: LE + AI

LE

Sigma=60s | horizon=2m | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.580	0.656	-0.076	0.025 [0.022, 0.028]	yes	2.99e-04	2.35e-04	8.95e-08	0.316	9.53e-05	3.57e-04	2.71e-04	1.27e-07	0.352	-5.02e-16
ridge	0.580	0.656	-0.076	0.025 [0.022, 0.028]	yes	2.99e-04	2.35e-04	8.95e-08	0.316	9.54e-05	3.57e-04	2.71e-04	1.27e-07	0.352	1.34e-17
lasso	0.578	0.649	-0.071	0.018 [0.016, 0.021]	yes	3.02e-04	2.37e-04	9.12e-08	0.313	9.82e-05	3.57e-04	2.72e-04	1.28e-07	0.343	-2.45e-20
elasticnet	0.573	0.640	-0.067	0.010 [0.008, 0.012]	yes	3.06e-04	2.40e-04	9.35e-08	0.314	1.02e-04	3.59e-04	2.73e-04	1.29e-07	0.341	1.21e-20
huber	0.577	0.669	-0.091	0.038 [0.036, 0.040]	yes	2.93e-04	2.30e-04	8.61e-08	0.315	7.76e-05	3.58e-04	2.70e-04	1.28e-07	0.359	-2.19e-05
poly2	0.597	0.661	-0.064	0.030 [0.024, 0.035]	yes	2.97e-04	2.34e-04	8.82e-08	0.289	1.14e-04	3.49e-04	2.64e-04	1.22e-07	0.306	-1.17e-15
poly3	0.599	0.655	-0.056	0.024 [0.019, 0.028]	yes	3.00e-04	2.36e-04	8.98e-08	0.325	1.18e-04	3.48e-04	2.64e-04	1.21e-07	0.305	-9.43e-14
knn	0.626	0.634	-0.008	0.003 [-0.003, 0.009]	no	3.08e-04	2.44e-04	9.51e-08	0.295	1.29e-04	3.36e-04	2.53e-04	1.13e-07	0.276	-3.75e-06
svr ★	0.622	0.676	-0.054	0.045 [0.040, 0.051]	yes	2.90e-04	2.28e-04	8.43e-08	0.283	1.03e-04	3.38e-04	2.53e-04	1.14e-07	0.287	-2.30e-05
kernelridge	0.621	0.651	-0.030	0.021 [0.010, 0.030]	yes	3.01e-04	2.36e-04	9.06e-08	0.426	1.19e-04	3.38e-04	2.53e-04	1.15e-07	0.521	-2.93e-06
gbm	0.510	0.504	0.006	-0.127 [-0.136, -0.116]	no	3.59e-04	2.86e-04	1.29e-07	0.358	1.77e-04	3.85e-04	2.91e-04	1.48e-07	0.308	-1.45e-07
rf	0.635	0.645	-0.010	0.014 [0.006, 0.023]	yes	3.04e-04	2.42e-04	9.22e-08	0.295	1.31e-04	3.32e-04	2.50e-04	1.10e-07	0.270	1.19e-07

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
extratrees	0.654	0.644	0.010	0.013 [0.005, 0.019]	yes	3.04e-04	2.42e-04	9.27e-08	0.294	1.33e-04	3.24e-04	2.43e-04	1.05e-07	0.252	2.52e-20
mlp	0.633	0.653	-0.020	0.022 [0.012, 0.031]	yes	3.00e-04	2.39e-04	9.03e-08	0.291	1.29e-04	3.33e-04	2.51e-04	1.11e-07	0.276	-1.72e-06
gp	0.616	0.649	-0.033	0.018 [0.012, 0.024]	yes	3.02e-04	2.38e-04	9.12e-08	0.295	1.26e-04	3.41e-04	2.56e-04	1.16e-07	0.295	-4.29e-06

best-by-test-R2 = **svr** (0.676) · linear 0.656 · persistence 0.631 · H=120s · 1780.9s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI Sigma=60s | horizon=2m | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.577	0.652	-0.075	0.021 [0.018, 0.025]	yes	3.01e-04	2.36e-04	9.04e-08	0.320	9.51e-05	3.58e-04	2.72e-04	1.28e-07	0.356	-4.22e-17
ridge	0.577	0.652	-0.075	0.021 [0.018, 0.025]	yes	3.01e-04	2.36e-04	9.04e-08	0.320	9.52e-05	3.58e-04	2.72e-04	1.28e-07	0.355	1.34e-17
lasso	0.576	0.646	-0.070	0.015 [0.011, 0.018]	yes	3.03e-04	2.39e-04	9.21e-08	0.314	9.98e-05	3.58e-04	2.73e-04	1.28e-07	0.345	2.13e-21
elasticnet	0.572	0.638	-0.065	0.007 [0.004, 0.009]	yes	3.07e-04	2.41e-04	9.42e-08	0.315	1.05e-04	3.60e-04	2.73e-04	1.29e-07	0.341	-4.98e-20
huber	0.575	0.664	-0.089	0.033 [0.031, 0.035]	yes	2.95e-04	2.31e-04	8.73e-08	0.319	7.70e-05	3.59e-04	2.71e-04	1.29e-07	0.364	-2.19e-05
poly2	0.594	0.651	-0.057	0.020 [0.014, 0.026]	yes	3.01e-04	2.38e-04	9.07e-08	0.294	1.17e-04	3.51e-04	2.65e-04	1.23e-07	0.305	-5.66e-15
poly3	0.599	0.646	-0.047	0.015 [0.009, 0.021]	yes	3.03e-04	2.40e-04	9.20e-08	0.330	1.24e-04	3.48e-04	2.64e-04	1.21e-07	0.302	2.03e-13
knn	0.624	0.634	-0.010	0.003 [-0.005, 0.009]	no	3.09e-04	2.44e-04	9.52e-08	0.297	1.31e-04	3.37e-04	2.54e-04	1.14e-07	0.277	-4.52e-06
svr ★	0.619	0.671	-0.052	0.040 [0.034, 0.045]	yes	2.93e-04	2.30e-04	8.56e-08	0.285	1.03e-04	3.40e-04	2.54e-04	1.15e-07	0.288	-2.29e-05
kernelridge	0.618	0.647	-0.029	0.017 [0.008, 0.025]	yes	3.03e-04	2.38e-04	9.16e-08	0.461	1.21e-04	3.40e-04	2.54e-04	1.15e-07	0.533	-3.06e-06
gbm	0.508	0.504	0.004	-0.127 [-0.136, -0.116]	no	3.59e-04	2.86e-04	1.29e-07	0.358	1.76e-04	3.86e-04	2.91e-04	1.49e-07	0.309	-1.41e-07
rf	0.632	0.640	-0.009	0.010 [-0.000, 0.018]	no	3.06e-04	2.43e-04	9.34e-08	0.298	1.34e-04	3.34e-04	2.51e-04	1.11e-07	0.271	1.27e-07
extratrees	0.651	0.636	0.015	0.006 [-0.004, 0.013]	no	3.07e-04	2.45e-04	9.45e-08	0.298	1.38e-04	3.25e-04	2.44e-04	1.06e-07	0.252	4.17e-20
mlp	0.630	0.654	-0.025	0.024 [0.014, 0.032]	yes	3.00e-04	2.38e-04	8.98e-08	0.289	1.26e-04	3.35e-04	2.52e-04	1.12e-07	0.274	-2.74e-06
gp	0.613	0.641	-0.028	0.011 [0.003, 0.017]	yes	3.05e-04	2.41e-04	9.32e-08	0.297	1.28e-04	3.42e-04	2.57e-04	1.17e-07	0.297	-4.50e-06

best-by-test-R2 = **svr** (0.671) · linear 0.652 · persistence 0.631 · H=120s · 1836.1s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=5m Sigma=60s · horizon=5m (H=300s) — metrics: LE + AI

LE Sigma=60s | horizon=5m | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.626	0.703	-0.077	0.027 [0.023, 0.031]	yes	4.12e-04	3.26e-04	1.70e-07	0.225	1.67e-04	4.89e-04	3.69e-04	2.39e-07	0.246	-1.96e-16
ridge	0.626	0.703	-0.077	0.027 [0.023, 0.031]	yes	4.12e-04	3.26e-04	1.70e-07	0.225	1.67e-04	4.89e-04	3.69e-04	2.39e-07	0.246	1.98e-17
lasso	0.625	0.698	-0.073	0.022 [0.019, 0.026]	yes	4.15e-04	3.28e-04	1.73e-07	0.225	1.70e-04	4.89e-04	3.69e-04	2.39e-07	0.243	-4.41e-20
elasticnet	0.623	0.692	-0.069	0.016 [0.013, 0.019]	yes	4.20e-04	3.31e-04	1.76e-07	0.226	1.73e-04	4.91e-04	3.71e-04	2.41e-07	0.242	-8.53e-21
huber	0.623	0.720	-0.097	0.044 [0.042, 0.046]	yes	4.00e-04	3.13e-04	1.60e-07	0.218	1.36e-04	4.91e-04	3.66e-04	2.41e-07	0.251	-3.85e-05
poly2	0.647	0.705	-0.058	0.029 [0.022, 0.036]	yes	4.11e-04	3.27e-04	1.69e-07	0.218	1.97e-04	4.75e-04	3.56e-04	2.26e-07	0.216	2.17e-15
poly3	0.649	0.697	-0.048	0.021 [0.016, 0.027]	yes	4.16e-04	3.31e-04	1.73e-07	0.240	2.04e-04	4.73e-04	3.55e-04	2.24e-07	0.213	3.79e-13
knn	0.682	0.674	0.008	-0.002 [-0.009, 0.006]	no	4.32e-04	3.45e-04	1.86e-07	0.224	2.21e-04	4.51e-04	3.33e-04	2.03e-07	0.187	-4.33e-06
svr ★	0.677	0.733	-0.056	0.058 [0.052, 0.065]	yes	3.90e-04	3.07e-04	1.52e-07	0.200	1.57e-04	4.54e-04	3.33e-04	2.06e-07	0.197	-4.02e-05
kernelridge	0.678	0.699	-0.022	0.024 [0.014, 0.033]	yes	4.15e-04	3.29e-04	1.72e-07	0.281	2.02e-04	4.54e-04	3.35e-04	2.06e-07	0.283	-2.30e-06
gbm	0.601	0.597	0.004	-0.079 [-0.087, -0.070]	no	4.80e-04	3.88e-04	2.30e-07	0.272	2.62e-04	5.05e-04	3.79e-04	2.55e-07	0.219	1.72e-08
rf	0.708	0.675	0.033	-0.001 [-0.014, 0.011]	no	4.31e-04	3.48e-04	1.86e-07	0.231	2.31e-04	4.32e-04	3.20e-04	1.87e-07	0.175	5.53e-08
extratrees	0.708	0.682	0.027	0.006 [-0.005, 0.015]	no	4.27e-04	3.45e-04	1.82e-07	0.227	2.28e-04	4.32e-04	3.19e-04	1.86e-07	0.170	-3.96e-20
mlp	0.689	0.696	-0.006	0.020 [0.007, 0.031]	yes	4.17e-04	3.36e-04	1.74e-07	0.216	2.17e-04	4.45e-04	3.30e-04	1.98e-07	0.184	2.29e-06
gp	0.665	0.695	-0.029	0.019 [0.010, 0.028]	yes	4.18e-04	3.32e-04	1.75e-07	0.217	2.04e-04	4.62e-04	3.40e-04	2.14e-07	0.195	-2.92e-06

best-by-test-R2 = **svr** (0.733) · linear 0.703 · persistence 0.676 · H=300s · 1803.2s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI Sigma=60s | horizon=5m | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.623	0.698	-0.075	0.023 [0.018, 0.027]	yes	4.15e-04	3.28e-04	1.72e-07	0.227	1.68e-04	4.90e-04	3.71e-04	2.41e-07	0.249	-5.21e-16
ridge	0.623	0.698	-0.075	0.023 [0.018, 0.027]	yes	4.15e-04	3.28e-04	1.72e-07	0.227	1.68e-04	4.90e-04	3.71e-04	2.41e-07	0.249	1.96e-17
lasso	0.623	0.694	-0.071	0.018 [0.013, 0.022]	yes	4.18e-04	3.30e-04	1.75e-07	0.227	1.73e-04	4.91e-04	3.71e-04	2.41e-07	0.245	-5.12e-20
elasticnet	0.621	0.688	-0.067	0.012 [0.008, 0.016]	yes	4.22e-04	3.33e-04	1.78e-07	0.228	1.78e-04	4.92e-04	3.72e-04	2.42e-07	0.243	-2.84e-20
huber	0.620	0.715	-0.094	0.039 [0.036, 0.041]	yes	4.04e-04	3.16e-04	1.63e-07	0.222	1.36e-04	4.92e-04	3.68e-04	2.42e-07	0.255	-3.86e-05
poly2	0.643	0.693	-0.050	0.017 [0.010, 0.024]	yes	4.19e-04	3.33e-04	1.76e-07	0.224	2.03e-04	4.77e-04	3.57e-04	2.28e-07	0.218	5.35e-15
poly3	0.650	0.687	-0.037	0.011 [0.003, 0.019]	yes	4.23e-04	3.38e-04	1.79e-07	0.245	2.14e-04	4.73e-04	3.54e-04	2.24e-07	0.222	6.07e-13
knn	0.679	0.671	0.008	-0.005 [-0.015, 0.003]	no	4.34e-04	3.48e-04	1.88e-07	0.228	2.26e-04	4.53e-04	3.35e-04	2.05e-07	0.189	-5.68e-06
svr ★	0.674	0.725	-0.051	0.049 [0.042, 0.057]	yes	3.96e-04	3.12e-04	1.57e-07	0.203	1.67e-04	4.56e-04	3.34e-04	2.08e-07	0.198	-3.97e-05
kernelridge	0.675	0.695	-0.019	0.019 [0.008, 0.028]	yes	4.18e-04	3.34e-04	1.75e-07	0.307	2.09e-04	4.55e-04	3.36e-04	2.07e-07	0.299	-1.95e-06
gbm	0.598	0.591	0.007	-0.085 [-0.092, -0.075]	no	4.84e-04	3.91e-04	2.34e-07	0.274	2.65e-04	5.07e-04	3.81e-04	2.57e-07	0.221	3.20e-08
rf	0.715	0.658	0.056	-0.017 [-0.032, -0.004]	no	4.42e-04	3.57e-04	1.95e-07	0.238	2.43e-04	4.27e-04	3.16e-04	1.82e-07	0.170	6.66e-08
extratrees	0.705	0.670	0.035	-0.006 [-0.019, 0.005]	no	4.34e-04	3.52e-04	1.89e-07	0.233	2.38e-04	4.34e-04	3.21e-04	1.88e-07	0.171	1.51e-20
mlp	0.685	0.691	-0.006	0.015 [0.002, 0.027]	yes	4.20e-04	3.38e-04	1.77e-07	0.220	2.20e-04	4.49e-04	3.32e-04	2.01e-07	0.187	-5.20e-06
gp	0.664	0.689	-0.024	0.013 [0.003, 0.021]	yes	4.22e-04	3.37e-04	1.78e-07	0.220	2.13e-04	4.63e-04	3.41e-04	2.14e-07	0.196	-1.72e-06

best-by-test-R2 = **svr** (0.725) · linear 0.698 · persistence 0.676 · H=300s · 1850.7s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=10m Sigma=60s · horizon=10m (H=600s) — metrics: LE + AI

LE Sigma=60s | horizon=10m | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.632	0.711	-0.079	0.029 [0.024, 0.033]	yes	5.61e-04	4.45e-04	3.15e-07	0.191	2.54e-04	6.60e-04	4.96e-04	4.36e-07	0.210	5.56e-16
ridge	0.632	0.711	-0.079	0.029 [0.024, 0.033]	yes	5.61e-04	4.46e-04	3.15e-07	0.191	2.54e-04	6.60e-04	4.96e-04	4.36e-07	0.210	-4.74e-17
lasso	0.631	0.708	-0.076	0.025 [0.021, 0.029]	yes	5.64e-04	4.48e-04	3.18e-07	0.192	2.57e-04	6.61e-04	4.97e-04	4.37e-07	0.208	1.13e-19
elasticnet	0.630	0.703	-0.073	0.021 [0.017, 0.025]	yes	5.68e-04	4.51e-04	3.23e-07	0.193	2.60e-04	6.62e-04	4.98e-04	4.38e-07	0.208	1.73e-20
huber	0.628	0.733	-0.105	0.051 [0.048, 0.053]	yes	5.39e-04	4.23e-04	2.91e-07	0.182	2.05e-04	6.64e-04	4.92e-04	4.41e-07	0.216	-6.07e-05
poly2	0.654	0.710	-0.056	0.028 [0.020, 0.035]	yes	5.62e-04	4.50e-04	3.16e-07	0.190	2.95e-04	6.40e-04	4.76e-04	4.09e-07	0.187	4.13e-15

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
poly3	0.657	0.701	-0.043	0.018 [0.012, 0.025]	yes	5.71e-04	4.58e-04	3.26e-07	0.218	3.07e-04	6.37e-04	4.74e-04	4.06e-07	0.182	1.23e-12
knn	0.694	0.678	0.016	-0.004 [-0.013, 0.004]	no	5.92e-04	4.79e-04	3.50e-07	0.197	3.30e-04	6.02e-04	4.41e-04	3.63e-07	0.158	-6.37e-06
svr ★	0.687	0.748	-0.061	0.066 [0.059, 0.074]	yes	5.23e-04	4.13e-04	2.74e-07	0.167	2.28e-04	6.09e-04	4.40e-04	3.71e-07	0.167	-6.26e-05
kernelridge	0.687	0.703	-0.016	0.021 [0.008, 0.033]	yes	5.68e-04	4.53e-04	3.23e-07	0.356	2.98e-04	6.09e-04	4.44e-04	3.71e-07	0.283	-2.28e-06
gbm	0.632	0.625	0.008	-0.057 [-0.065, -0.049]	no	6.39e-04	5.22e-04	4.08e-07	0.232	3.69e-04	6.60e-04	4.92e-04	4.36e-07	0.181	2.17e-07
rf	0.730	0.661	0.069	-0.022 [-0.039, -0.007]	no	6.08e-04	4.97e-04	3.69e-07	0.210	3.59e-04	5.66e-04	4.14e-04	3.20e-07	0.140	2.06e-07
extratrees	0.720	0.678	0.042	-0.004 [-0.018, 0.006]	no	5.92e-04	4.84e-04	3.51e-07	0.203	3.45e-04	5.76e-04	4.21e-04	3.32e-07	0.142	-2.93e-20
mlp	0.701	0.694	0.007	0.012 [-0.004, 0.025]	no	5.77e-04	4.70e-04	3.33e-07	0.191	3.28e-04	5.95e-04	4.37e-04	3.55e-07	0.154	7.58e-06
gp	0.674	0.692	-0.018	0.010 [-0.001, 0.020]	no	5.79e-04	4.66e-04	3.35e-07	0.191	3.17e-04	6.21e-04	4.52e-04	3.86e-07	0.165	-3.72e-06

best-by-test-R2 = **svr** (0.748) · linear 0.711 · persistence 0.682 · H=600s · 1803.3s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI Sigma=60s | horizon=10m | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.629	0.706	-0.077	0.024 [0.019, 0.029]	yes	5.65e-04	4.49e-04	3.20e-07	0.194	2.57e-04	6.63e-04	4.99e-04	4.40e-07	0.213	-1.39e-15
ridge	0.629	0.706	-0.077	0.024 [0.018, 0.029]	yes	5.65e-04	4.49e-04	3.20e-07	0.194	2.57e-04	6.63e-04	4.99e-04	4.40e-07	0.213	-4.85e-17
lasso	0.628	0.703	-0.074	0.021 [0.015, 0.025]	yes	5.69e-04	4.51e-04	3.23e-07	0.194	2.61e-04	6.63e-04	4.99e-04	4.40e-07	0.211	1.95e-19
elasticnet	0.627	0.698	-0.071	0.016 [0.011, 0.021]	yes	5.73e-04	4.55e-04	3.28e-07	0.195	2.66e-04	6.64e-04	5.00e-04	4.41e-07	0.209	1.19e-19
huber	0.625	0.727	-0.102	0.045 [0.042, 0.048]	yes	5.45e-04	4.27e-04	2.97e-07	0.186	2.06e-04	6.67e-04	4.94e-04	4.44e-07	0.219	-6.09e-05
poly2	0.651	0.697	-0.046	0.015 [0.006, 0.023]	yes	5.74e-04	4.60e-04	3.30e-07	0.195	3.05e-04	6.43e-04	4.79e-04	4.14e-07	0.188	2.05e-14
poly3	0.658	0.690	-0.032	0.008 [-0.002, 0.017]	no	5.80e-04	4.69e-04	3.37e-07	0.223	3.22e-04	6.36e-04	4.72e-04	4.05e-07	0.190	5.14e-13
knn	0.691	0.672	0.019	-0.011 [-0.022, -0.002]	no	5.98e-04	4.85e-04	3.57e-07	0.201	3.39e-04	6.05e-04	4.43e-04	3.67e-07	0.159	-8.04e-06
svr ★	0.685	0.736	-0.051	0.054 [0.045, 0.061]	yes	5.36e-04	4.26e-04	2.87e-07	0.173	2.49e-04	6.11e-04	4.41e-04	3.74e-07	0.168	-6.22e-05
kernelridge	0.685	0.696	-0.011	0.014 [0.002, 0.024]	yes	5.75e-04	4.61e-04	3.31e-07	0.340	3.08e-04	6.11e-04	4.46e-04	3.73e-07	0.274	-1.86e-06
gbm	0.629	0.618	0.011	-0.064 [-0.071, -0.055]	no	6.44e-04	5.27e-04	4.15e-07	0.235	3.76e-04	6.63e-04	4.94e-04	4.39e-07	0.183	2.25e-07
rf	0.726	0.648	0.078	-0.034 [-0.052, -0.017]	no	6.18e-04	5.06e-04	3.82e-07	0.215	3.70e-04	5.70e-04	4.17e-04	3.24e-07	0.141	1.99e-07
extratrees	0.717	0.664	0.054	-0.019 [-0.034, -0.005]	no	6.05e-04	4.96e-04	3.66e-07	0.210	3.61e-04	5.79e-04	4.23e-04	3.35e-07	0.143	-3.50e-20
mlp	0.697	0.688	0.009	0.006 [-0.010, 0.020]	no	5.83e-04	4.76e-04	3.40e-07	0.195	3.39e-04	5.99e-04	4.40e-04	3.59e-07	0.155	7.65e-06
gp	0.677	0.686	-0.009	0.004 [-0.008, 0.014]	no	5.85e-04	4.72e-04	3.42e-07	0.196	3.24e-04	6.19e-04	4.52e-04	3.83e-07	0.164	-3.63e-06

best-by-test-R2 = **svr** (0.736) · linear 0.706 · persistence 0.682 · H=600s · 1867.5s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=15m Sigma=60s · horizon=15m (H=900s) — metrics: LE + AI

LE Sigma=60s | horizon=15m | metric=LE

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.622	0.707	-0.085	0.031 [0.026, 0.035]	yes	6.85e-04	5.46e-04	4.69e-07	0.182	3.27e-04	8.05e-04	6.02e-04	6.48e-07	0.200	-7.75e-17
ridge	0.622	0.707	-0.085	0.031 [0.026, 0.035]	yes	6.85e-04	5.46e-04	4.69e-07	0.182	3.27e-04	8.05e-04	6.02e-04	6.48e-07	0.200	5.23e-17
lasso	0.622	0.704	-0.082	0.028 [0.023, 0.032]	yes	6.88e-04	5.48e-04	4.73e-07	0.183	3.30e-04	8.05e-04	6.02e-04	6.49e-07	0.199	2.95e-19
elasticnet	0.621	0.700	-0.080	0.024 [0.020, 0.028]	yes	6.92e-04	5.51e-04	4.79e-07	0.184	3.33e-04	8.06e-04	6.03e-04	6.50e-07	0.198	2.99e-19
huber	0.617	0.732	-0.115	0.056 [0.054, 0.059]	yes	6.54e-04	5.14e-04	4.28e-07	0.172	2.62e-04	8.10e-04	5.95e-04	6.57e-07	0.206	-7.99e-05
poly2	0.645	0.703	-0.058	0.027 [0.018, 0.035]	yes	6.89e-04	5.55e-04	4.75e-07	0.183	3.77e-04	7.80e-04	5.77e-04	6.09e-07	0.180	2.38e-15
poly3	0.648	0.692	-0.044	0.016 [0.009, 0.022]	yes	7.01e-04	5.66e-04	4.92e-07	0.212	3.92e-04	7.77e-04	5.74e-04	6.03e-07	0.174	8.48e-14
knn	0.687	0.667	0.020	-0.009 [-0.019, -0.000]	no	7.30e-04	5.93e-04	5.32e-07	0.191	4.20e-04	7.33e-04	5.32e-04	5.37e-07	0.151	-6.73e-06
svr ★	0.680	0.737	-0.058	0.061 [0.052, 0.070]	yes	6.48e-04	5.16e-04	4.19e-07	0.162	3.07e-04	7.41e-04	5.29e-04	5.49e-07	0.160	-7.79e-05
kernelridge	0.681	0.698	-0.017	0.022 [0.009, 0.031]	yes	6.95e-04	5.63e-04	4.83e-07	0.222	3.92e-04	7.40e-04	5.38e-04	5.48e-07	0.206	6.39e-06
gbm	0.635	0.624	0.011	-0.052 [-0.060, -0.043]	no	7.75e-04	6.38e-04	6.01e-07	0.222	4.64e-04	7.91e-04	5.85e-04	6.26e-07	0.171	2.54e-07
rf	0.712	0.649	0.063	-0.027 [-0.046, -0.012]	no	7.49e-04	6.14e-04	5.61e-07	0.205	4.51e-04	7.03e-04	5.12e-04	4.95e-07	0.141	1.33e-07
extratrees	0.714	0.663	0.051	-0.013 [-0.029, -0.001]	no	7.34e-04	6.05e-04	5.39e-07	0.200	4.44e-04	7.01e-04	5.08e-04	4.91e-07	0.135	-3.56e-20
mlp	0.692	0.678	0.014	0.002 [-0.015, 0.016]	no	7.17e-04	5.90e-04	5.15e-07	0.189	4.33e-04	7.27e-04	5.30e-04	5.29e-07	0.148	1.33e-05
gp	0.671	0.680	-0.009	0.004 [-0.011, 0.016]	no	7.15e-04	5.72e-04	5.12e-07	0.182	3.92e-04	7.52e-04	5.46e-04	5.65e-07	0.156	8.27e-06

best-by-test-R2 = **svr** (0.737) · linear 0.707 · persistence 0.676 · H=900s · 1777.1s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI Sigma=60s | horizon=15m | metric=AI

learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.619	0.702	-0.083	0.026 [0.019, 0.031]	yes	6.90e-04	5.50e-04	4.76e-07	0.185	3.31e-04	8.09e-04	6.05e-04	6.54e-07	0.203	-2.34e-17
ridge	0.619	0.702	-0.083	0.026 [0.019, 0.031]	yes	6.90e-04	5.50e-04	4.76e-07	0.185	3.31e-04	8.09e-04	6.05e-04	6.54e-07	0.203	5.21e-17
lasso	0.619	0.699	-0.080	0.023 [0.017, 0.028]	yes	6.94e-04	5.53e-04	4.81e-07	0.185	3.36e-04	8.09e-04	6.05e-04	6.54e-07	0.201	2.62e-19
elasticnet	0.618	0.695	-0.077	0.019 [0.013, 0.024]	yes	6.98e-04	5.56e-04	4.87e-07	0.187	3.41e-04	8.10e-04	6.06e-04	6.55e-07	0.200	2.22e-19
huber	0.614	0.726	-0.112	0.050 [0.047, 0.053]	yes	6.61e-04	5.19e-04	4.37e-07	0.176	2.63e-04	8.14e-04	5.99e-04	6.62e-07	0.210	-8.00e-05
poly2	0.641	0.689	-0.048	0.013 [0.004, 0.022]	yes	7.05e-04	5.67e-04	4.96e-07	0.189	3.88e-04	7.85e-04	5.81e-04	6.16e-07	0.182	1.51e-14
poly3	0.650	0.682	-0.032	0.006 [-0.005, 0.016]	no	7.13e-04	5.79e-04	5.08e-07	0.217	4.10e-04	7.75e-04	5.71e-04	6.01e-07	0.194	4.54e-13
knn	0.684	0.659	0.025	-0.017 [-0.030, -0.007]	no	7.38e-04	6.02e-04	5.44e-07	0.196	4.33e-04	7.36e-04	5.35e-04	5.42e-07	0.152	-8.69e-06
svr ★	0.677	0.732	-0.055	0.056 [0.046, 0.065]	yes	6.54e-04	5.22e-04	4.28e-07	0.165	3.16e-04	7.44e-04	5.32e-04	5.54e-07	0.161	-7.55e-05
kernelridge	0.678	0.686	-0.007	0.009 [-0.005, 0.020]	no	7.09e-04	5.75e-04	5.03e-07	0.231	4.03e-04	7.43e-04	5.40e-04	5.52e-07	0.235	6.97e-06
gbm	0.631	0.621	0.010	-0.055 [-0.064, -0.047]	no	7.78e-04	6.40e-04	6.06e-07	0.223	4.65e-04	7.95e-04	5.89e-04	6.33e-07	0.173	2.20e-07
rf	0.720	0.628	0.092	-0.048 [-0.068, -0.029]	no	7.71e-04	6.34e-04	5.94e-07	0.213	4.79e-04	6.93e-04	5.03e-04	4.80e-07	0.134	1.67e-07
extratrees	0.711	0.648	0.063	-0.028 [-0.045, -0.013]	no	7.50e-04	6.19e-04	5.62e-07	0.206	4.64e-04	7.04e-04	5.10e-04	4.95e-07	0.136	3.63e-22
mlp	0.690	0.667	0.023	-0.009 [-0.028, 0.008]	no	7.29e-04	6.00e-04	5.32e-07	0.194	4.45e-04	7.29e-04	5.30e-04	5.31e-07	0.150	3.41e-06
gp	0.670	0.684	-0.014	0.008 [-0.005, 0.018]	no	7.11e-04	5.77e-04	5.05e-07	0.186	4.10e-04	7.53e-04	5.47e-04	5.67e-07	0.157	9.42e-06

best-by-test-R2 = **svr** (0.732) · linear 0.702 · persistence 0.676 · H=900s · 1859.8s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=30m Sigma=60s · horizon=30m (H=1800s) — metrics: LE + AI

LE Sigma=60s | horizon=30m | metric=LE

	learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
	linear	0.582	0.681	-0.100	0.036 [0.030, 0.041]	yes	9.93e-04	7.99e-04	9.86e-07	0.180	5.17e-04	1.16e-03	8.60e-04	1.35e-06	0.196	7.72e-16
	ridge	0.582	0.681	-0.100	0.036 [0.030, 0.041]	yes	9.93e-04	8.00e-04	9.86e-07	0.180	5.17e-04	1.16e-03	8.60e-04	1.35e-06	0.196	-6.92e-17
	lasso	0.581	0.679	-0.098	0.033 [0.027, 0.039]	yes	9.96e-04	8.02e-04	9.93e-07	0.181	5.19e-04	1.16e-03	8.61e-04	1.35e-06	0.195	4.96e-19
	elasticnet	0.581	0.676	-0.095	0.031 [0.025, 0.036]	yes	1.00e-03	8.05e-04	1.00e-06	0.182	5.22e-04	1.16e-03	8.61e-04	1.35e-06	0.195	4.51e-19
	huber	0.575	0.717	-0.142	0.071 [0.068, 0.075]	yes	9.35e-04	7.38e-04	8.75e-07	0.166	4.04e-04	1.17e-03	8.48e-04	1.37e-06	0.204	-1.31e-04
	poly2	0.606	0.668	-0.062	0.022 [0.012, 0.032]	yes	1.01e-03	8.23e-04	1.03e-06	0.186	5.85e-04	1.13e-03	8.27e-04	1.27e-06	0.181	1.73e-14
	poly3	0.610	0.654	-0.044	0.008 [0.001, 0.016]	yes	1.03e-03	8.42e-04	1.07e-06	0.212	6.11e-04	1.12e-03	8.20e-04	1.26e-06	0.184	8.85e-13
	knn	0.651	0.626	0.025	-0.019 [-0.033, -0.009]	no	1.08e-03	8.85e-04	1.16e-06	0.196	6.51e-04	1.06e-03	7.61e-04	1.13e-06	0.153	-1.32e-05
	svr ★	0.640	0.717	-0.077	0.072 [0.062, 0.082]	yes	9.35e-04	7.51e-04	8.74e-07	0.161	4.63e-04	1.08e-03	7.53e-04	1.16e-06	0.166	-1.47e-04
	kernelridge	0.642	0.660	-0.018	0.015 [-0.000, 0.027]	no	1.02e-03	8.43e-04	1.05e-06	0.231	6.11e-04	1.07e-03	7.70e-04	1.16e-06	0.208	-1.05e-06
	gbm	0.613	0.600	0.013	-0.046 [-0.055, -0.037]	no	1.11e-03	9.23e-04	1.24e-06	0.217	6.91e-04	1.12e-03	8.15e-04	1.25e-06	0.166	-7.09e-08
	rf	0.692	0.581	0.111	-0.064 [-0.088, -0.044]	no	1.14e-03	9.46e-04	1.29e-06	0.218	7.35e-04	9.97e-04	7.14e-04	9.95e-07	0.134	4.54e-07
	extratrees	0.682	0.607	0.075	-0.038 [-0.057, -0.024]	no	1.10e-03	9.17e-04	1.21e-06	0.209	7.00e-04	1.01e-03	7.25e-04	1.03e-06	0.136	6.46e-20
	mlp	0.656	0.644	0.013	-0.002 [-0.020, 0.014]	no	1.05e-03	8.70e-04	1.10e-06	0.191	6.50e-04	1.05e-03	7.55e-04	1.11e-06	0.151	-3.14e-06
	gp	0.631	0.605	0.026	-0.040 [-0.057, -0.027]	no	1.10e-03	9.09e-04	1.22e-06	0.204	6.87e-04	1.09e-03	7.85e-04	1.19e-06	0.158	5.93e-06

best-by-test-R2 = **svr** (0.717) · linear 0.681 · persistence 0.646 · H=1800s · 1787.8s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI Sigma=60s | horizon=30m | metric=AI

	learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
	linear	0.578	0.675	-0.098	0.030 [0.022, 0.036]	yes	1.00e-03	8.06e-04	1.00e-06	0.183	5.24e-04	1.17e-03	8.66e-04	1.36e-06	0.199	7.54e-16
	ridge	0.578	0.675	-0.098	0.030 [0.022, 0.036]	yes	1.00e-03	8.06e-04	1.00e-06	0.183	5.24e-04	1.17e-03	8.66e-04	1.36e-06	0.199	-7.04e-17
	lasso	0.577	0.673	-0.095	0.027 [0.020, 0.034]	yes	1.01e-03	8.10e-04	1.01e-06	0.184	5.29e-04	1.17e-03	8.66e-04	1.36e-06	0.198	4.96e-19
	elasticnet	0.577	0.670	-0.093	0.024 [0.017, 0.031]	yes	1.01e-03	8.14e-04	1.02e-06	0.185	5.34e-04	1.17e-03	8.66e-04	1.37e-06	0.198	4.54e-19
	huber ★	0.571	0.710	-0.139	0.065 [0.061, 0.069]	yes	9.46e-04	7.46e-04	8.96e-07	0.170	4.07e-04	1.18e-03	8.54e-04	1.38e-06	0.208	-1.31e-04
	poly2	0.600	0.654	-0.054	0.008 [-0.003, 0.018]	no	1.04e-03	8.40e-04	1.07e-06	0.191	6.01e-04	1.14e-03	8.34e-04	1.29e-06	0.182	1.94e-14
	poly3	0.611	0.645	-0.033	-0.001 [-0.014, 0.011]	no	1.05e-03	8.61e-04	1.10e-06	0.217	6.35e-04	1.12e-03	8.16e-04	1.25e-06	0.197	1.30e-12
	knn	0.648	0.614	0.034	-0.031 [-0.047, -0.018]	no	1.09e-03	9.02e-04	1.19e-06	0.201	6.75e-04	1.07e-03	7.65e-04	1.14e-06	0.154	-1.53e-05
	svr	0.637	0.708	-0.071	0.063 [0.052, 0.072]	yes	9.50e-04	7.65e-04	9.03e-07	0.165	4.79e-04	1.08e-03	7.57e-04	1.17e-06	0.166	-1.48e-04
	kernelridge	0.641	0.645	-0.004	-0.001 [-0.018, 0.012]	no	1.05e-03	8.64e-04	1.10e-06	0.284	6.33e-04	1.08e-03	7.72e-04	1.16e-06	0.235	5.83e-09
	gbm	0.610	0.597	0.013	-0.048 [-0.059, -0.039]	no	1.12e-03	9.28e-04	1.24e-06	0.218	6.98e-04	1.12e-03	8.19e-04	1.26e-06	0.167	-4.28e-08
	rf	0.689	0.559	0.130	-0.087 [-0.112, -0.063]	no	1.17e-03	9.70e-04	1.36e-06	0.225	7.62e-04	1.00e-03	7.18e-04	1.00e-06	0.134	5.36e-07
	extratrees	0.680	0.587	0.093	-0.058 [-0.079, -0.040]	no	1.13e-03	9.43e-04	1.28e-06	0.217	7.34e-04	1.02e-03	7.28e-04	1.03e-06	0.136	1.54e-20
	mlp	0.654	0.631	0.023	-0.015 [-0.036, 0.004]	no	1.07e-03	8.87e-04	1.14e-06	0.196	6.71e-04	1.06e-03	7.59e-04	1.12e-06	0.151	3.88e-06
	gp	0.633	0.606	0.027	-0.039 [-0.055, -0.026]	no	1.10e-03	9.13e-04	1.22e-06	0.205	6.96e-04	1.09e-03	7.84e-04	1.19e-06	0.158	7.66e-06

best-by-test-R2 = **huber** (0.710) · linear 0.675 · persistence 0.646 · H=1800s · 1904.1s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

horizon=1h Sigma=60s · horizon=1h (H=3600s) — metrics: LE + AI

LE Sigma=60s | horizon=1h | metric=LE

	learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
	linear	0.510	0.627	-0.118	0.032 [0.025, 0.038]	yes	1.48e-03	1.20e-03	2.19e-06	0.192	8.21e-04	1.71e-03	1.25e-03	2.91e-06	0.205	8.50e-16
	ridge	0.510	0.627	-0.118	0.032 [0.025, 0.038]	yes	1.48e-03	1.20e-03	2.19e-06	0.192	8.21e-04	1.71e-03	1.25e-03	2.91e-06	0.205	5.49e-17
	lasso	0.510	0.626	-0.116	0.030 [0.023, 0.036]	yes	1.49e-03	1.21e-03	2.21e-06	0.192	8.23e-04	1.71e-03	1.25e-03	2.91e-06	0.205	-2.94e-19
	elasticnet	0.510	0.624	-0.114	0.028 [0.021, 0.034]	yes	1.49e-03	1.21e-03	2.22e-06	0.193	8.25e-04	1.71e-03	1.25e-03	2.92e-06	0.205	-4.88e-19
	huber ★	0.500	0.683	-0.183	0.087 [0.083, 0.092]	yes	1.37e-03	1.08e-03	1.87e-06	0.171	6.16e-04	1.72e-03	1.23e-03	2.97e-06	0.217	-2.24e-04
	poly2	0.534	0.598	-0.064	0.003 [-0.011, 0.015]	no	1.54e-03	1.25e-03	2.37e-06	0.203	9.19e-04	1.66e-03	1.21e-03	2.77e-06	0.195	3.02e-14
	poly3	0.541	0.577	-0.037	-0.018 [-0.027, -0.009]	no	1.58e-03	1.29e-03	2.49e-06	0.225	9.67e-04	1.65e-03	1.20e-03	2.73e-06	0.204	3.20e-13
	knn	0.587	0.550	0.037	-0.045 [-0.061, -0.033]	no	1.63e-03	1.34e-03	2.65e-06	0.213	1.01e-03	1.57e-03	1.12e-03	2.45e-06	0.165	-1.73e-05
	svr	0.569	0.682	-0.113	0.087 [0.074, 0.098]	yes	1.37e-03	1.10e-03	1.87e-06	0.167	6.69e-04	1.60e-03	1.10e-03	2.56e-06	0.184	-2.74e-04
	kernelridge	0.577	0.593	-0.016	-0.003 [-0.018, 0.010]	no	1.55e-03	1.28e-03	2.40e-06	0.290	9.36e-04	1.59e-03	1.13e-03	2.51e-06	0.266	-6.47e-06
	gbm	0.559	0.541	0.018	-0.054 [-0.066, -0.043]	no	1.64e-03	1.37e-03	2.70e-06	0.225	1.05e-03	1.62e-03	1.17e-03	2.62e-06	0.172	-1.84e-07
	rf	0.635	0.488	0.147	-0.107 [-0.137, -0.082]	no	1.74e-03	1.45e-03	3.01e-06	0.238	1.15e-03	1.47e-03	1.04e-03	2.17e-06	0.142	6.20e-08
	extratrees	0.625	0.527	0.098	-0.068 [-0.089, -0.052]	no	1.67e-03	1.40e-03	2.78e-06	0.225	1.08e-03	1.49e-03	1.06e-03	2.23e-06	0.144	-9.69e-21
	mlp	0.592	0.540	0.052	-0.055 [-0.080, -0.034]	no	1.65e-03	1.38e-03	2.71e-06	0.219	1.08e-03	1.56e-03	1.12e-03	2.43e-06	0.161	2.77e-05
	gp	0.566	0.587	-0.022	-0.008 [-0.023, 0.004]	no	1.56e-03	1.28e-03	2.43e-06	0.201	9.50e-04	1.61e-03	1.14e-03	2.58e-06	0.172	-3.12e-05

best-by-test-R2 = **huber** (0.683) · linear 0.627 · persistence 0.596 · H=3600s · 1740.9s. te.* = OOS test losses; tr.* = in-sample train losses (RMSE/MAE/bias in vol units, MSE=vol^2, QLIKE on variance).

AI Sigma=60s | horizon=1h | metric=AI


learner	train R2	test R2	gap	dR2 [CI95]	sig	te.RMSE	te.MAE	te.MSE	te.QLIKE	te.bias	tr.RMSE	tr.MAE	tr.MSE	tr.QLIKE	tr.bias
linear	0.506	0.620	-0.113	0.024 [0.015, 0.033]	yes	1.50e-03	1.22e-03	2.24e-06	0.195	8.35e-04	1.71e-03	1.26e-03	2.93e-06	0.207	-1.83e-15
ridge	0.506	0.620	-0.113	0.024 [0.015, 0.033]	yes	1.50e-03	1.22e-03	2.24e-06	0.195	8.35e-04	1.71e-03	1.26e-03	2.93e-06	0.207	5.59e-17
lasso	0.506	0.618	-0.111	0.022 [0.013, 0.031]	yes	1.50e-03	1.22e-03	2.25e-06	0.196	8.39e-04	1.71e-03	1.26e-03	2.94e-06	0.207	-1.97e-19
elasticnet	0.506	0.615	-0.109	0.020 [0.011, 0.028]	yes	1.51e-03	1.22e-03	2.27e-06	0.196	8.44e-04	1.71e-03	1.26e-03	2.94e-06	0.207	-4.59e-19
huber ★	0.497	0.675	-0.179	0.079 [0.075, 0.085]	yes	1.38e-03	1.09e-03	1.91e-06	0.175	6.25e-04	1.73e-03	1.24e-03	2.99e-06	0.220	-2.24e-04
poly2	0.528	0.582	-0.054	-0.014 [-0.030, -0.000]	no	1.57e-03	1.28e-03	2.46e-06	0.207	9.44e-04	1.67e-03	1.22e-03	2.80e-06	0.195	-2.96e-14
poly3	0.544	0.569	-0.025	-0.027 [-0.044, -0.011]	no	1.59e-03	1.32e-03	2.54e-06	0.230	1.01e-03	1.65e-03	1.19e-03	2.71e-06	0.208	-3.04e-13
knn	0.586	0.533	0.053	-0.063 [-0.082, -0.046]	no	1.66e-03	1.38e-03	2.75e-06	0.219	1.05e-03	1.57e-03	1.12e-03	2.46e-06	0.165	-2.08e-05
svr	0.568	0.659	-0.091	0.063 [0.050, 0.075]	yes	1.42e-03	1.14e-03	2.01e-06	0.176	7.24e-04	1.60e-03	1.10e-03	2.57e-06	0.183	-2.71e-04
kernelridge	0.577	0.578	-0.001	-0.018 [-0.036, -0.003]	no	1.58e-03	1.31e-03	2.49e-06	0.296	9.76e-04	1.59e-03	1.13e-03	2.52e-06	0.270	-5.61e-06
gbm	0.558	0.540	0.018	-0.056 [-0.070, -0.045]	no	1.65e-03	1.38e-03	2.71e-06	0.226	1.06e-03	1.62e-03	1.17e-03	2.63e-06	0.172	-2.47e-07
rf	0.634	0.452	0.182	-0.143 [-0.178, -0.113]	no	1.80e-03	1.50e-03	3.23e-06	0.249	1.21e-03	1.47e-03	1.05e-03	2.17e-06	0.142	9.98e-08
extratrees	0.624	0.501	0.123	-0.095 [-0.123, -0.071]	no	1.71e-03	1.44e-03	2.94e-06	0.235	1.14e-03	1.50e-03	1.06e-03	2.24e-06	0.144	3.23e-21
mlp	0.591	0.517	0.073	-0.078 [-0.110, -0.052]	no	1.69e-03	1.40e-03	2.84e-06	0.225	1.10e-03	1.56e-03	1.11e-03	2.43e-06	0.163	-1.47e-05
gp	0.566	0.579	-0.013	-0.017 [-0.034, -0.003]	no	1.58e-03	1.30e-03	2.48e-06	0.205	9.77e-04	1.61e-03	1.14e-03	2.58e-06	0.171	-3.04e-05

Capacity - does a flexible learner beat linear? Per cell we compare the BEST learner's test R2 to the LINEAR learner's test R2. The heatmap shows the capacity gain (best - linear test R2) over (horizon x cov-duration), one panel per metric; the table lists the best learner per cell. A near-0 gain means the predictive map is essentially linear at that horizon; a positive gain means capacity helps.

Sigma=30s

Capacity gain (best - linear test R2) [metric=LE]

LE



In plain terms

How much a flexible learner beats plain linear at each cov-dur x horizon.

Math / formula

gain = R2_OOS(best-by-test) - R2_OOS(linear); per (Sigma, H).

How to read it

cell text = best learner + gain; bright/high = capacity helps, ~0 = linear is enough.


Insight

metric LE; empty cells = not finished yet.

Sigma=30s

Capacity gain (best - linear test R2) [metric=AI]

AI



In plain terms

How much a flexible learner beats plain linear at each cov-dur x horizon.

Math / formula

gain = R2_OOS(best-by-test) - R2_OOS(linear); per (Sigma, H).

How to read it

cell text = best learner + gain; bright/high = capacity helps, ~0 = linear is enough.

Insight

metric AI; empty cells = not finished yet.


Best learner + capacity per cell (every finished cov-dur x horizon x metric).

Sigma	horizon	metric	best learner	best test R2	linear test R2	capacity gain	persistence R2	best overfit gap
30s	nextwin	LE	svr	0.502	0.471	0.031	0.427	-0.050
30s	nextwin	AI	svr	0.502	0.468	0.034	0.427	-0.052
30s	30s	LE	svr	0.505	0.468	0.037	0.423	-0.055
30s	30s	AI	svr	0.504	0.465	0.039	0.423	-0.056
30s	45s	LE	svr	0.495	0.468	0.027	0.422	-0.045
30s	45s	AI	svr	0.497	0.465	0.032	0.422	-0.049
30s	60s	LE	svr	0.558	0.533	0.025	0.485	-0.044
30s	60s	AI	svr	0.560	0.529	0.031	0.485	-0.048
30s	2m	LE	svr	0.629	0.587	0.042	0.536	-0.055
30s	2m	AI	svr	0.624	0.583	0.042	0.536	-0.053
30s	5m	LE	svr	0.682	0.627	0.055	0.572	-0.056
30s	5m	AI	svr	0.671	0.622	0.050	0.572	-0.048
30s	10m	LE	svr	0.682	0.630	0.052	0.573	-0.053
30s	10m	AI	svr	0.674	0.625	0.048	0.573	-0.046
30s	15m	LE	svr	0.678	0.625	0.053	0.565	-0.059
30s	15m	AI	svr	0.668	0.620	0.048	0.565	-0.051
30s	30m	LE	huber	0.652	0.601	0.051	0.534	-0.166
30s	30m	AI	huber	0.645	0.595	0.050	0.534	-0.164
30s	1h	LE	huber	0.621	0.549	0.072	0.485	-0.204
30s	1h	AI	huber	0.614	0.543	0.071	0.485	-0.202
45s	nextwin	LE	svr	0.576	0.554	0.022	0.523	-0.048
45s	nextwin	AI	svr	0.573	0.550	0.023	0.523	-0.048
45s	45s	LE	svr	0.578	0.552	0.026	0.519	-0.048
45s	45s	AI	svr	0.574	0.549	0.026	0.519	-0.047

Sigma	horizon	metric	best learner	best test R2	linear test R2	capacity gain	persistence R2	best overfit gap
45s	60s	LE	svr	0.579	0.552	0.026	0.519	-0.048
45s	60s	AI	svr	0.575	0.548	0.027	0.519	-0.047
45s	2m	LE	svr	0.635	0.605	0.030	0.571	-0.058
45s	2m	AI	svr	0.631	0.601	0.030	0.571	-0.055
45s	5m	LE	svr	0.702	0.667	0.034	0.630	-0.048
45s	5m	AI	svr	0.699	0.663	0.036	0.630	-0.048
45s	10m	LE	svr	0.727	0.677	0.050	0.637	-0.065
45s	10m	AI	svr	0.717	0.672	0.045	0.637	-0.057
45s	15m	LE	svr	0.722	0.672	0.050	0.630	-0.070
45s	15m	AI	svr	0.708	0.667	0.041	0.630	-0.058
45s	30m	LE	svr	0.691	0.647	0.045	0.599	-0.077
45s	30m	AI	huber	0.682	0.641	0.041	0.599	-0.150
45s	1h	LE	huber	0.656	0.594	0.062	0.549	-0.192
45s	1h	AI	huber	0.648	0.586	0.062	0.549	-0.188
60s	nextwin	LE	svr	0.626	0.610	0.015	0.587	-0.046
60s	nextwin	AI	svr	0.624	0.607	0.017	0.587	-0.047
60s	60s	LE	svr	0.628	0.612	0.016	0.588	-0.045
60s	60s	AI	svr	0.625	0.608	0.017	0.588	-0.045
60s	2m	LE	svr	0.676	0.656	0.020	0.631	-0.054
60s	2m	AI	svr	0.671	0.652	0.018	0.631	-0.052
60s	5m	LE	svr	0.733	0.703	0.031	0.676	-0.056
60s	5m	AI	svr	0.725	0.698	0.027	0.676	-0.051
60s	10m	LE	svr	0.748	0.711	0.037	0.682	-0.061
60s	10m	AI	svr	0.736	0.706	0.029	0.682	-0.051
60s	15m	LE	svr	0.737	0.707	0.031	0.676	-0.058
60s	15m	AI	svr	0.732	0.702	0.031	0.676	-0.055
60s	30m	LE	svr	0.717	0.681	0.036	0.646	-0.077
60s	30m	AI	huber	0.710	0.675	0.035	0.646	-0.139
60s	1h	LE	huber	0.683	0.627	0.055	0.596	-0.183
60s	1h	AI	huber	0.675	0.620	0.055	0.596	-0.179

LE vs AI - is the answer metric-robust? **AI** **LE** For each (cov-dur x horizon) with BOTH metrics finished, do LE and AI agree on (a) the best learner, and (b) the SIGN of the dial's dR2 over persistence? AI is the headline / LE the diagnostic (Step-6 convention). **metric_robust** = same dR2 sign AND both significant-or-both-not.

Sigma=30s LE vs AI dial edge (best-learner dR2)



In plain terms Do the two SPD rulers (Log-Euclidean vs affine-invariant) agree on the dial's OOS edge?

Math / formula point = one (Sigma, H) cell; x = LE dR2, y = AI dR2; the diagonal is perfect agreement.

How to read it points on/near the diagonal and in the same quadrant = metric-robust; off-diagonal = fragile.

Insight clustering along y=x and in the ++ quadrant would confirm the AI dial reproduces the LE verdict.

Sigma=30s LE vs AI test-R2 per (cov-dur x horizon x learner) **AI** **LE**



In plain terms Every finished (cov-dur x horizon x learner) cell plotted as LE vs AI test R2.

Math / formula point = one (Sigma, H, model); x = LE test R2, y = AI test R2; dashed = AI=LE.

How to read it tight clustering on the diagonal = the SPD metric barely changes accuracy (metric-robust).

Insight color = learner; each spans all cov-dur x horizon cells where both metrics finished.

LE-vs-AI agreement per (cov-dur x horizon). same best learner? same dR2 sign? metric-robust?

Sigma	horizon	best LE	best AI	same best	dR2 LE	dR2 AI	same sign	metric-robust
30s	nextwin	svr	svr	True	0.075	0.075	True	True
30s	30s	svr	svr	True	0.083	0.081	True	True
30s	45s	svr	svr	True	0.073	0.075	True	True
30s	60s	svr	svr	True	0.073	0.075	True	True
30s	2m	svr	svr	True	0.094	0.089	True	True
30s	5m	svr	svr	True	0.110	0.100	True	True
30s	10m	svr	svr	True	0.110	0.101	True	True
30s	15m	svr	svr	True	0.113	0.103	True	True
30s	30m	huber	huber	True	0.118	0.111	True	True
30s	1h	huber	huber	True	0.136	0.129	True	True
45s	nextwin	svr	svr	True	0.053	0.050	True	True
45s	45s	svr	svr	True	0.059	0.055	True	True
45s	60s	svr	svr	True	0.059	0.056	True	True
45s	2m	svr	svr	True	0.064	0.060	True	True
45s	5m	svr	svr	True	0.071	0.068	True	True
45s	10m	svr	svr	True	0.090	0.080	True	True

Sigma	horizon	best LE	best AI	same best	dR2 LE	dR2 AI	same sign	metric-robust
45s	15m	svr	svr	True	0.092	0.078	True	True
45s	30m	svr	huber	False	0.093	0.083	True	True
45s	1h	huber	huber	True	0.107	0.099	True	True
60s	nextwin	svr	svr	True	0.038	0.037	True	True
60s	60s	svr	svr	True	0.040	0.038	True	True
60s	2m	svr	svr	True	0.045	0.040	True	True
60s	5m	svr	svr	True	0.058	0.049	True	True
60s	10m	svr	svr	True	0.066	0.054	True	True
60s	15m	svr	svr	True	0.061	0.056	True	True
60s	30m	svr	huber	False	0.072	0.065	True	True
60s	1h	huber	huber	True	0.087	0.079	True	True

Complete Step-7 performance picture. The full-matrix tab renders every learner; here we render the per-CELL summary + the per-YEAR breakdown of the best learner (train/test R2, dR2 over persistence, QLIKE, RMSE) for every finished (cov-dur x horizon x metric). Every row is labelled with cov-duration + horizon + metric. QLIKE is on variance; RMSE/MAE/bias in vol units.

Best-learner loss battery per cell (train R2 / test R2 / gap / dR2 [CI] / sig / RMSE / MAE / QLIKE / bias / n_test).

Sigma	horizon	metric	best learner	train R2	test R2	gap	dR2 [CI95]	sig	RMSE	MAE	QLIKE	bias	n_test
30s	nextwin	LE	svr	0.453	0.502	-0.050	0.075 [0.070,0.080]	yes	2.24e-04	1.79e-04	0.777	5.10e-05	624446
30s	nextwin	AI	svr	0.450	0.502	-0.052	0.075 [0.070,0.079]	yes	2.24e-04	1.79e-04	0.747	5.04e-05	624446
30s	30s	LE	svr	0.451	0.505	-0.055	0.083 [0.078,0.087]	yes	2.18e-04	1.75e-04	0.968	4.53e-05	597064
30s	30s	AI	svr	0.448	0.504	-0.056	0.081 [0.077,0.086]	yes	2.18e-04	1.76e-04	0.921	4.49e-05	597064
30s	45s	LE	svr	0.450	0.495	-0.045	0.073 [0.066,0.079]	yes	2.20e-04	1.77e-04	0.739	5.01e-05	596378
30s	45s	AI	svr	0.448	0.497	-0.049	0.075 [0.070,0.079]	yes	2.19e-04	1.77e-04	0.706	4.93e-05	596378
30s	60s	LE	svr	0.514	0.558	-0.044	0.073 [0.067,0.080]	yes	2.58e-04	2.04e-04	0.365	8.90e-05	617379
30s	60s	AI	svr	0.512	0.560	-0.048	0.075 [0.070,0.081]	yes	2.58e-04	2.04e-04	0.360	8.92e-05	617379
30s	2m	LE	svr	0.575	0.629	-0.055	0.094 [0.087,0.101]	yes	3.08e-04	2.43e-04	0.309	1.20e-04	621119
30s	2m	AI	svr	0.571	0.624	-0.053	0.089 [0.082,0.096]	yes	3.10e-04	2.44e-04	0.302	1.24e-04	621119
30s	5m	LE	svr	0.626	0.682	-0.056	0.110 [0.102,0.118]	yes	4.26e-04	3.38e-04	0.226	1.98e-04	617436
30s	5m	AI	svr	0.624	0.671	-0.048	0.100 [0.091,0.108]	yes	4.33e-04	3.44e-04	0.230	2.06e-04	617436
30s	10m	LE	svr	0.629	0.682	-0.053	0.110 [0.100,0.119]	yes	5.88e-04	4.68e-04	0.200	2.88e-04	610631
30s	10m	AI	svr	0.627	0.674	-0.046	0.101 [0.090,0.111]	yes	5.96e-04	4.76e-04	0.204	3.02e-04	610631
30s	15m	LE	svr	0.619	0.678	-0.059	0.113 [0.103,0.123]	yes	7.18e-04	5.76e-04	0.194	3.67e-04	603869
30s	15m	AI	svr	0.617	0.668	-0.051	0.103 [0.091,0.114]	yes	7.30e-04	5.86e-04	0.198	3.83e-04	603869
30s	30m	LE	huber	0.486	0.652	-0.166	0.118 [0.113,0.123]	yes	1.04e-03	8.28e-04	0.197	4.84e-04	583830
30s	30m	AI	huber	0.481	0.645	-0.164	0.111 [0.106,0.116]	yes	1.05e-03	8.36e-04	0.201	4.88e-04	583830
30s	1h	LE	huber	0.417	0.621	-0.204	0.136 [0.129,0.143]	yes	1.50e-03	1.19e-03	0.200	7.18e-04	542965
30s	1h	AI	huber	0.412	0.614	-0.202	0.129 [0.123,0.136]	yes	1.51e-03	1.20e-03	0.204	7.25e-04	542965
45s	nextwin	LE	svr	0.528	0.576	-0.048	0.053 [0.046,0.060]	yes	2.38e-04	1.87e-04	0.415	6.15e-05	416834
45s	nextwin	AI	svr	0.525	0.573	-0.048	0.050 [0.043,0.056]	yes	2.39e-04	1.88e-04	0.412	6.15e-05	416834
45s	45s	LE	svr	0.530	0.578	-0.048	0.059 [0.055,0.063]	yes	2.31e-04	1.84e-04	0.410	6.18e-05	398151
45s	45s	AI	svr	0.527	0.574	-0.047	0.055 [0.050,0.060]	yes	2.32e-04	1.84e-04	0.411	6.11e-05	398151
45s	60s	LE	svr	0.530	0.579	-0.048	0.059 [0.055,0.064]	yes	2.31e-04	1.84e-04	0.402	6.26e-05	397447
45s	60s	AI	svr	0.528	0.575	-0.047	0.056 [0.051,0.060]	yes	2.32e-04	1.84e-04	0.400	6.35e-05	397447
45s	2m	LE	svr	0.578	0.635	-0.058	0.064 [0.059,0.070]	yes	2.74e-04	2.15e-04	0.316	9.38e-05	411541
45s	2m	AI	svr	0.575	0.631	-0.055	0.060 [0.054,0.065]	yes	2.76e-04	2.17e-04	0.316	9.41e-05	411541
45s	5m	LE	svr	0.653	0.702	-0.048	0.071 [0.064,0.079]	yes	3.93e-04	3.10e-04	0.219	1.70e-04	411857
45s	5m	AI	svr	0.651	0.699	-0.048	0.068 [0.060,0.075]	yes	3.95e-04	3.13e-04	0.221	1.75e-04	411857
45s	10m	LE	svr	0.662	0.727	-0.065	0.090 [0.083,0.098]	yes	5.38e-04	4.27e-04	0.179	2.47e-04	407145
45s	10m	AI	svr	0.660	0.717	-0.057	0.080 [0.071,0.088]	yes	5.48e-04	4.37e-04	0.184	2.63e-04	407145
45s	15m	LE	svr	0.652	0.722	-0.070	0.092 [0.084,0.102]	yes	6.67e-04	5.31e-04	0.171	3.21e-04	403083
45s	15m	AI	svr	0.650	0.708	-0.058	0.078 [0.069,0.087]	yes	6.84e-04	5.46e-04	0.178	3.43e-04	403083
45s	30m	LE	svr	0.614	0.691	-0.077	0.093 [0.081,0.104]	yes	9.78e-04	7.85e-04	0.174	4.96e-04	389737
45s	30m	AI	huber	0.532	0.682	-0.150	0.083 [0.079,0.088]	yes	9.93e-04	7.87e-04	0.183	4.46e-04	389737
45s	1h	LE	huber	0.464	0.656	-0.192	0.107 [0.102,0.113]	yes	1.43e-03	1.13e-03	0.184	6.64e-04	362482
45s	1h	AI	huber	0.460	0.648	-0.188	0.099 [0.094,0.105]	yes	1.44e-03	1.14e-03	0.188	6.72e-04	362482
60s	nextwin	LE	svr	0.579	0.626	-0.046	0.038 [0.031,0.045]	yes	2.49e-04	1.94e-04	0.375	6.55e-05	312652
60s	nextwin	AI	svr	0.577	0.624	-0.047	0.037 [0.032,0.042]	yes	2.49e-04	1.94e-04	0.371	6.48e-05	312652
60s	60s	LE	svr	0.583	0.628	-0.045	0.040 [0.033,0.047]	yes	2.41e-04	1.89e-04	0.356	6.62e-05	298209
60s	60s	AI	svr	0.580	0.625	-0.045	0.038 [0.031,0.044]	yes	2.41e-04	1.90e-04	0.356	6.62e-05	298209
60s	2m	LE	svr	0.622	0.676	-0.054	0.045 [0.040,0.051]	yes	2.90e-04	2.28e-04	0.283	1.03e-04	309044
60s	2m	AI	svr	0.619	0.671	-0.052	0.040 [0.034,0.045]	yes	2.93e-04	2.30e-04	0.285	1.03e-04	309044
60s	5m	LE	svr	0.677	0.733	-0.056	0.058 [0.052,0.065]	yes	3.90e-04	3.07e-04	0.200	1.57e-04	309043
60s	5m	AI	svr	0.674	0.725	-0.051	0.049 [0.042,0.057]	yes	3.96e-04	3.12e-04	0.203	1.67e-04	309043
60s	10m	LE	svr	0.687	0.748	-0.061	0.066 [0.059,0.074]	yes	5.23e-04	4.13e-04	0.167	2.28e-04	305725
60s	10m	AI	svr	0.685	0.736	-0.051	0.054 [0.045,0.061]	yes	5.36e-04	4.26e-04	0.173	2.49e-04	305725
60s	15m	LE	svr	0.680	0.737	-0.058	0.061 [0.052,0.070]	yes	6.48e-04	5.16e-04	0.162	3.07e-04	302315
60s	15m	AI	svr	0.677	0.732	-0.055	0.056 [0.046,0.065]	yes	6.54e-04	5.22e-04	0.165	3.16e-04	302315
60s	30m	LE	svr	0.640	0.717	-0.077	0.072 [0.062,0.082]	yes	9.35e-04	7.51e-04	0.161	4.63e-04	292276
60s	30m	AI	huber	0.571	0.710	-0.139	0.065 [0.061,0.069]	yes	9.46e-04	7.46e-04	0.170	4.07e-04	292276
60s	1h	LE	huber	0.500	0.683	-0.183	0.087 [0.083,0.092]	yes	1.37e-03	1.08e-03	0.171	6.16e-04	271823
60s	1h	AI	huber	0.497	0.675	-0.179	0.079 [0.075,0.085]	yes	1.38e-03	1.09e-03	0.175	6.25e-04	271823

Per-year OOS breakdown of the best learner (purged walk-forward test years). train R2 / test R2 / persistence R2 / dR2 / QLIKE / RMSE.

	Sigma	horizon	metric	learner	test year	train R2	test R2	persistence R2	dR2	QLIKE	RMSE	
	30s	nextwin	LE	svr	2023	0.449	0.590	0.479	0.110	0.858	2.37e-04	
	30s	nextwin	LE	svr	2024	0.464	0.418	0.347	0.071	0.954	2.11e-04	
	30s	nextwin	LE	svr	2025	0.446	0.440	0.415	0.024	0.543	2.23e-04	
	30s	nextwin	AI	svr	2023	0.451	0.586	0.479	0.106	0.882	2.38e-04	
	30s	nextwin	AI	svr	2024	0.460	0.414	0.347	0.067	0.877	2.12e-04	
	30s	nextwin	AI	svr	2025	0.443	0.447	0.415	0.032	0.512	2.22e-04	
	30s	30s	LE	svr	2023	0.448	0.591	0.473	0.118	1.840	2.31e-04	
	30s	30s	LE	svr	2024	0.463	0.416	0.341	0.075	0.714	2.06e-04	
	30s	30s	LE	svr	2025	0.442	0.448	0.414	0.034	0.450	2.16e-04	
	30s	30s	AI	svr	2023	0.450	0.589	0.473	0.115	1.659	2.31e-04	
	30s	30s	AI	svr	2024	0.460	0.409	0.341	0.068	0.723	2.07e-04	
	30s	30s	AI	svr	2025	0.438	0.453	0.414	0.040	0.466	2.14e-04	
	30s	45s	LE	svr	2023	0.449	0.583	0.473	0.110	0.805	2.33e-04	
	30s	45s	LE	svr	2024	0.461	0.412	0.341	0.071	0.990	2.07e-04	
	30s	45s	LE	svr	2025	0.442	0.429	0.414	0.015	0.438	2.19e-04	
	30s	45s	AI	svr	2023	0.450	0.579	0.473	0.106	0.729	2.34e-04	
	30s	45s	AI	svr	2024	0.458	0.408	0.341	0.067	0.942	2.07e-04	
	30s	45s	AI	svr	2025	0.440	0.445	0.414	0.032	0.456	2.16e-04	
	30s	60s	LE	svr	2023	0.500	0.642	0.532	0.110	0.380	2.74e-04	
	30s	60s	LE	svr	2024	0.528	0.487	0.408	0.079	0.352	2.39e-04	
	30s	60s	LE	svr	2025	0.510	0.483	0.473	0.010	0.363	2.61e-04	
	30s	60s	AI	svr	2023	0.503	0.639	0.532	0.107	0.369	2.75e-04	
	30s	60s	AI	svr	2024	0.525	0.478	0.408	0.070	0.349	2.41e-04	
	30s	60s	AI	svr	2025	0.506	0.503	0.473	0.029	0.362	2.56e-04	
	30s	2m	LE	svr	2023	0.558	0.721	0.580	0.140	0.334	3.18e-04	
	30s	2m	LE	svr	2024	0.591	0.535	0.465	0.070	0.292	2.93e-04	
	30s	2m	LE	svr	2025	0.570	0.559	0.521	0.038	0.302	3.13e-04	
	30s	2m	AI	svr	2023	0.557	0.702	0.580	0.122	0.327	3.28e-04	
	30s	2m	AI	svr	2024	0.586	0.544	0.465	0.079	0.281	2.90e-04	
	30s	2m	AI	svr	2025	0.566	0.563	0.521	0.042	0.300	3.12e-04	
	30s	5m	LE	svr	2023	0.604	0.770	0.612	0.158	0.252	4.32e-04	
	30s	5m	LE	svr	2024	0.642	0.605	0.508	0.096	0.203	4.00e-04	
	30s	5m	LE	svr	2025	0.625	0.599	0.556	0.043	0.224	4.44e-04	
	30s	5m	AI	svr	2023	0.608	0.755	0.612	0.143	0.264	4.46e-04	
	30s	5m	AI	svr	2024	0.638	0.583	0.508	0.075	0.206	4.10e-04	
	30s	5m	AI	svr	2025	0.620	0.604	0.556	0.048	0.221	4.41e-04	
	30s	10m	LE	svr	2023	0.599	0.781	0.609	0.172	0.220	5.84e-04	
	30s	10m	LE	svr	2024	0.649	0.589	0.514	0.075	0.180	5.62e-04	
	30s	10m	LE	svr	2025	0.628	0.591	0.559	0.032	0.200	6.18e-04	
	30s	10m	AI	svr	2023	0.603	0.767	0.609	0.158	0.230	6.03e-04	
	30s	10m	AI	svr	2024	0.647	0.560	0.514	0.046	0.187	5.82e-04	
	30s	10m	AI	svr	2025	0.623	0.610	0.559	0.051	0.197	6.03e-04	
	30s	15m	LE	svr	2023	0.585	0.774	0.599	0.175	0.214	7.21e-04	
	30s	15m	LE	svr	2024	0.639	0.578	0.509	0.069	0.175	6.90e-04	
	30s	15m	LE	svr	2025	0.619	0.599	0.553	0.046	0.191	7.43e-04	
	30s	15m	AI	svr	2023	0.588	0.760	0.599	0.161	0.224	7.42e-04	
	30s	15m	AI	svr	2024	0.636	0.549	0.509	0.040	0.180	7.14e-04	
	30s	15m	AI	svr	2025	0.614	0.610	0.553	0.057	0.189	7.32e-04	
	30s	30m	LE	huber	2023	0.440	0.672	0.563	0.109	0.257	1.21e-03	
	30s	30m	LE	huber	2024	0.497	0.631	0.484	0.147	0.167	8.96e-04	
	30s	30m	LE	huber	2025	0.498	0.634	0.526	0.108	0.167	9.86e-04	
	30s	30m	AI	huber	2023	0.438	0.660	0.563	0.097	0.263	1.23e-03	
	30s	30m	AI	huber	2024	0.491	0.625	0.484	0.141	0.170	9.03e-04	
	30s	30m	AI	huber	2025	0.492	0.636	0.526	0.110	0.169	9.84e-04	
	30s	1h	LE	huber	2023	0.354	0.636	0.507	0.128	0.267	1.76e-03	
	30s	1h	LE	huber	2024	0.429	0.608	0.439	0.169	0.167	1.27e-03	
	30s	1h	LE	huber	2025	0.434	0.605	0.484	0.122	0.166	1.42e-03	
	30s	1h	AI	huber	2023	0.354	0.622	0.507	0.115	0.274	1.80e-03	
	30s	1h	AI	huber	2024	0.423	0.603	0.439	0.163	0.170	1.28e-03	
	30s	1h	AI	huber	2025	0.429	0.610	0.484	0.126	0.167	1.41e-03	
	45s	nextwin	LE	svr	2023	0.523	0.651	0.571	0.080	0.416	2.54e-04	

	Sigma	horizon	metric	learner	test year	train R2	test R2	persistence R2	dR2	QLIKE	RMSE	
	45s	nextwin	LE	svr	2024	0.539	0.506	0.447	0.059	0.447	2.22e-04	
	45s	nextwin	LE	svr	2025	0.521	0.516	0.511	0.006	0.384	2.38e-04	
	45s	nextwin	AI	svr	2023	0.523	0.648	0.571	0.077	0.407	2.55e-04	
	45s	nextwin	AI	svr	2024	0.535	0.498	0.447	0.051	0.441	2.24e-04	
	45s	nextwin	AI	svr	2025	0.517	0.519	0.511	0.008	0.389	2.38e-04	
	45s	45s	LE	svr	2023	0.523	0.649	0.566	0.084	0.416	2.48e-04	
	45s	45s	LE	svr	2024	0.542	0.502	0.440	0.061	0.449	2.16e-04	
	45s	45s	LE	svr	2025	0.524	0.531	0.513	0.018	0.367	2.28e-04	
	45s	45s	AI	svr	2023	0.524	0.647	0.566	0.081	0.415	2.48e-04	
	45s	45s	AI	svr	2024	0.539	0.493	0.440	0.052	0.445	2.18e-04	
	45s	45s	AI	svr	2025	0.519	0.528	0.513	0.015	0.373	2.28e-04	
	45s	60s	LE	svr	2023	0.525	0.654	0.565	0.088	0.407	2.46e-04	
	45s	60s	LE	svr	2024	0.541	0.496	0.440	0.056	0.434	2.18e-04	
	45s	60s	LE	svr	2025	0.524	0.528	0.512	0.016	0.367	2.28e-04	
	45s	60s	AI	svr	2023	0.525	0.650	0.565	0.085	0.399	2.47e-04	
	45s	60s	AI	svr	2024	0.537	0.487	0.440	0.047	0.431	2.19e-04	
	45s	60s	AI	svr	2025	0.521	0.529	0.512	0.017	0.371	2.28e-04	
	45s	2m	LE	svr	2023	0.561	0.714	0.615	0.098	0.323	2.87e-04	
	45s	2m	LE	svr	2024	0.592	0.561	0.500	0.061	0.308	2.56e-04	
	45s	2m	LE	svr	2025	0.574	0.571	0.558	0.013	0.318	2.78e-04	
	45s	2m	AI	svr	2023	0.561	0.708	0.615	0.093	0.325	2.90e-04	
	45s	2m	AI	svr	2024	0.589	0.549	0.500	0.048	0.306	2.60e-04	
	45s	2m	AI	svr	2025	0.571	0.573	0.558	0.015	0.318	2.77e-04	
	45s	5m	LE	svr	2023	0.630	0.792	0.668	0.125	0.234	3.91e-04	
	45s	5m	LE	svr	2024	0.669	0.619	0.572	0.047	0.201	3.75e-04	
	45s	5m	LE	svr	2025	0.652	0.619	0.616	0.003	0.222	4.14e-04	
	45s	5m	AI	svr	2023	0.631	0.784	0.668	0.116	0.241	3.99e-04	
	45s	5m	AI	svr	2024	0.666	0.615	0.572	0.043	0.200	3.77e-04	
	45s	5m	AI	svr	2025	0.648	0.625	0.616	0.010	0.221	4.10e-04	
	45s	10m	LE	svr	2023	0.634	0.812	0.670	0.142	0.193	5.34e-04	
	45s	10m	LE	svr	2024	0.677	0.653	0.585	0.068	0.160	5.10e-04	
	45s	10m	LE	svr	2025	0.663	0.646	0.624	0.022	0.183	5.68e-04	
	45s	10m	AI	svr	2023	0.635	0.801	0.670	0.131	0.202	5.50e-04	
	45s	10m	AI	svr	2024	0.675	0.625	0.585	0.040	0.167	5.31e-04	
	45s	10m	AI	svr	2025	0.659	0.651	0.624	0.027	0.183	5.64e-04	
	45s	15m	LE	svr	2023	0.619	0.810	0.660	0.150	0.184	6.60e-04	
	45s	15m	LE	svr	2024	0.667	0.652	0.582	0.070	0.151	6.27e-04	
	45s	15m	LE	svr	2025	0.655	0.632	0.619	0.014	0.177	7.11e-04	
	45s	15m	AI	svr	2023	0.620	0.796	0.660	0.136	0.196	6.84e-04	
	45s	15m	AI	svr	2024	0.666	0.620	0.582	0.038	0.159	6.55e-04	
	45s	15m	AI	svr	2025	0.650	0.632	0.619	0.013	0.178	7.11e-04	
	45s	30m	LE	svr	2023	0.566	0.787	0.623	0.164	0.189	9.73e-04	
	45s	30m	LE	svr	2024	0.635	0.599	0.557	0.042	0.156	9.33e-04	
	45s	30m	LE	svr	2025	0.620	0.605	0.591	0.014	0.176	1.02e-03	
	45s	30m	AI	huber	2023	0.489	0.694	0.623	0.071	0.239	1.17e-03	
	45s	30m	AI	huber	2024	0.540	0.666	0.557	0.109	0.154	8.52e-04	
	45s	30m	AI	huber	2025	0.544	0.673	0.591	0.082	0.156	9.33e-04	
	45s	1h	LE	huber	2023	0.402	0.671	0.568	0.103	0.243	1.67e-03	
	45s	1h	LE	huber	2024	0.475	0.645	0.511	0.134	0.152	1.21e-03	
	45s	1h	LE	huber	2025	0.483	0.640	0.547	0.093	0.155	1.36e-03	
	45s	1h	AI	huber	2023	0.401	0.654	0.568	0.086	0.251	1.72e-03	
	45s	1h	AI	huber	2024	0.470	0.641	0.511	0.131	0.155	1.21e-03	
	45s	1h	AI	huber	2025	0.478	0.643	0.547	0.096	0.156	1.35e-03	
	60s	nextwin	LE	svr	2023	0.569	0.700	0.632	0.068	0.401	2.62e-04	
	60s	nextwin	LE	svr	2024	0.590	0.565	0.522	0.044	0.378	2.30e-04	
	60s	nextwin	LE	svr	2025	0.576	0.558	0.570	-0.013	0.349	2.53e-04	
	60s	nextwin	AI	svr	2023	0.569	0.699	0.632	0.067	0.391	2.62e-04	
	60s	nextwin	AI	svr	2024	0.587	0.556	0.522	0.034	0.370	2.32e-04	
	60s	nextwin	AI	svr	2025	0.573	0.563	0.570	-0.008	0.352	2.52e-04	
	60s	60s	LE	svr	2023	0.578	0.699	0.631	0.068	0.361	2.56e-04	
	60s	60s	LE	svr	2024	0.595	0.555	0.518	0.037	0.359	2.25e-04	

Sigma	horizon	metric	learner	test year	train R2	test R2	persistence R2	dR2	QLIKE	RMSE
60s	60s	LE	svr	2025	0.577	0.574	0.576	-0.002	0.350	2.40e-04
60s	60s	AI	svr	2023	0.577	0.694	0.631	0.064	0.353	2.58e-04
60s	60s	AI	svr	2024	0.591	0.547	0.518	0.029	0.360	2.27e-04
60s	60s	AI	svr	2025	0.572	0.579	0.576	0.003	0.355	2.38e-04
60s	2m	LE	svr	2023	0.602	0.752	0.672	0.080	0.302	3.01e-04
60s	2m	LE	svr	2024	0.633	0.605	0.570	0.034	0.265	2.72e-04
60s	2m	LE	svr	2025	0.621	0.611	0.614	-0.003	0.281	2.97e-04
60s	2m	AI	svr	2023	0.602	0.748	0.672	0.077	0.303	3.03e-04
60s	2m	AI	svr	2024	0.631	0.588	0.570	0.018	0.269	2.78e-04
60s	2m	AI	svr	2025	0.617	0.612	0.614	-0.002	0.282	2.96e-04
60s	5m	LE	svr	2023	0.653	0.815	0.710	0.105	0.216	3.88e-04
60s	5m	LE	svr	2024	0.691	0.662	0.626	0.036	0.182	3.71e-04
60s	5m	LE	svr	2025	0.678	0.658	0.659	-0.002	0.202	4.12e-04
60s	5m	AI	svr	2023	0.652	0.807	0.710	0.097	0.221	3.95e-04
60s	5m	AI	svr	2024	0.689	0.640	0.626	0.013	0.188	3.83e-04
60s	5m	AI	svr	2025	0.672	0.659	0.659	0.000	0.200	4.11e-04
60s	10m	LE	svr	2023	0.658	0.833	0.713	0.120	0.176	5.09e-04
60s	10m	LE	svr	2024	0.702	0.670	0.639	0.031	0.153	5.04e-04
60s	10m	LE	svr	2025	0.689	0.671	0.666	0.004	0.172	5.55e-04
60s	10m	AI	svr	2023	0.655	0.821	0.713	0.108	0.186	5.27e-04
60s	10m	AI	svr	2024	0.701	0.639	0.639	0.000	0.161	5.28e-04
60s	10m	AI	svr	2025	0.685	0.672	0.666	0.006	0.173	5.54e-04
60s	15m	LE	svr	2023	0.642	0.832	0.704	0.128	0.164	6.19e-04
60s	15m	LE	svr	2024	0.698	0.628	0.636	-0.008	0.155	6.49e-04
60s	15m	LE	svr	2025	0.682	0.669	0.662	0.007	0.166	6.74e-04
60s	15m	AI	svr	2023	0.641	0.826	0.704	0.121	0.169	6.31e-04
60s	15m	AI	svr	2024	0.696	0.618	0.636	-0.018	0.159	6.57e-04
60s	15m	AI	svr	2025	0.678	0.670	0.662	0.008	0.167	6.73e-04
60s	30m	LE	svr	2023	0.593	0.814	0.669	0.146	0.167	9.07e-04
60s	30m	LE	svr	2024	0.661	0.605	0.611	-0.007	0.153	9.27e-04
60s	30m	LE	svr	2025	0.645	0.646	0.635	0.011	0.163	9.70e-04
60s	30m	AI	huber	2023	0.528	0.723	0.669	0.055	0.220	1.11e-03
60s	30m	AI	huber	2024	0.578	0.699	0.611	0.087	0.142	8.09e-04
60s	30m	AI	huber	2025	0.585	0.698	0.635	0.063	0.147	8.96e-04
60s	1h	LE	huber	2023	0.436	0.699	0.614	0.085	0.225	1.60e-03
60s	1h	LE	huber	2024	0.510	0.674	0.564	0.109	0.141	1.16e-03
60s	1h	LE	huber	2025	0.519	0.663	0.590	0.073	0.146	1.31e-03
60s	1h	AI	huber	2023	0.436	0.682	0.614	0.068	0.234	1.64e-03
60s	1h	AI	huber	2024	0.505	0.671	0.564	0.107	0.143	1.16e-03
60s	1h	AI	huber	2025	0.516	0.667	0.590	0.076	0.147	1.30e-03

GridSearch best-params + wall-time per cell.

Sigma	horizon	metric	best learner	best_params	elapsed (s)
30s	nextwin	LE	svr	{"C": 3.0, "gamma": 0.1}	2774.4
30s	nextwin	AI	svr	{"C": 3.0, "gamma": 0.1}	2728.3
30s	30s	LE	svr	{"C": 3.0, "gamma": 0.1}	2676.5
30s	30s	AI	svr	{"C": 3.0, "gamma": 0.1}	2661.7
30s	45s	LE	svr	{"C": 3.0, "gamma": 0.1}	2681.8
30s	45s	AI	svr	{"C": 3.0, "gamma": 0.1}	2656.5
30s	60s	LE	svr	{"C": 3.0, "gamma": 0.1}	2688.0
30s	60s	AI	svr	{"C": 3.0, "gamma": 0.1}	2684.7
30s	2m	LE	svr	{"C": 3.0, "gamma": 0.1}	2735.9
30s	2m	AI	svr	{"C": 3.0, "gamma": 0.1}	2732.6
30s	5m	LE	svr	{"C": 3.0, "gamma": 0.1}	2655.5
30s	5m	AI	svr	{"C": 3.0, "gamma": 0.1}	2636.2
30s	10m	LE	svr	{"C": 3.0, "gamma": 0.1}	2654.0
30s	10m	AI	svr	{"C": 3.0, "gamma": 0.1}	2588.6
30s	15m	LE	svr	{"C": 3.0, "gamma": 0.1}	2605.4
30s	15m	AI	svr	{"C": 3.0, "gamma": 0.1}	2657.2
30s	30m	LE	huber	(no grid / defaults)	2605.3
30s	30m	AI	huber	(no grid / defaults)	2548.2

Sigma	horizon	metric	best learner	best_params	elapsed (s)
30s	1h	LE	huber	(no grid / defaults)	2600.2
30s	1h	AI	huber	(no grid / defaults)	2489.9
45s	nextwin	LE	svr	{"C": 3.0, "gamma": 0.1}	2249.6
45s	nextwin	AI	svr	{"C": 3.0, "gamma": 0.1}	2172.8
45s	45s	LE	svr	{"C": 3.0, "gamma": 0.1}	2060.2
45s	45s	AI	svr	{"C": 3.0, "gamma": 0.1}	2497.0
45s	60s	LE	svr	{"C": 3.0, "gamma": 0.1}	2070.2
45s	60s	AI	svr	{"C": 3.0, "gamma": 0.1}	2088.6
45s	2m	LE	svr	{"C": 3.0, "gamma": 0.1}	2107.2
45s	2m	AI	svr	{"C": 3.0, "gamma": 0.1}	2105.6
45s	5m	LE	svr	{"C": 3.0, "gamma": 0.1}	2117.3
45s	5m	AI	svr	{"C": 3.0, "gamma": 0.1}	2111.2
45s	10m	LE	svr	{"C": 3.0, "gamma": 0.1}	2056.4
45s	10m	AI	svr	{"C": 3.0, "gamma": 0.1}	2079.4
45s	15m	LE	svr	{"C": 3.0, "gamma": 0.1}	2082.4
45s	15m	AI	svr	{"C": 3.0, "gamma": 0.1}	2115.5
45s	30m	LE	svr	{"C": 3.0, "gamma": 0.1}	2044.7
45s	30m	AI	huber	(no grid / defaults)	2049.5
45s	1h	LE	huber	(no grid / defaults)	1998.6
45s	1h	AI	huber	(no grid / defaults)	2026.4
60s	nextwin	LE	svr	{"C": 3.0, "gamma": 0.1}	1801.4
60s	nextwin	AI	svr	{"C": 3.0, "gamma": 0.1}	1840.1
60s	60s	LE	svr	{"C": 3.0, "gamma": 0.1}	1771.6
60s	60s	AI	svr	{"C": 3.0, "gamma": 0.1}	1814.1
60s	2m	LE	svr	{"C": 3.0, "gamma": 0.1}	1780.9
60s	2m	AI	svr	{"C": 3.0, "gamma": 0.1}	1836.1
60s	5m	LE	svr	{"C": 3.0, "gamma": 0.1}	1803.2
60s	5m	AI	svr	{"C": 1.0, "gamma": 0.1}	1850.7
60s	10m	LE	svr	{"C": 3.0, "gamma": 0.1}	1803.3
60s	10m	AI	svr	{"C": 3.0, "gamma": 0.1}	1867.5
60s	15m	LE	svr	{"C": 3.0, "gamma": 0.1}	1777.1
60s	15m	AI	svr	{"C": 3.0, "gamma": 0.1}	1859.8
60s	30m	LE	svr	{"C": 3.0, "gamma": 0.1}	1787.8
60s	30m	AI	huber	(no grid / defaults)	1904.1
60s	1h	LE	huber	(no grid / defaults)	1740.9
60s	1h	AI	huber	(no grid / defaults)	1764.1

Step 7 (Track A) — horizon x cov-duration bake-off (PARTIAL render: 54 cell(s) finished)

Holding the 30/45/60 s masked covariance state fixed, the fixed Step-6 predictor `[rv, s_slow, s_fast, v]` is baked off with 15 learners against a realized-vol target at growing horizons, under LE and AI features. The headline object is **dR2(horizon)** - the dial's OOS edge over persistence - which is small/coincident at next-window and tests whether it grows (leading) with H. Current peaks: 30s/AI: peak dR2 0.129 @ 1h [sig]; 45s/AI: peak dR2 0.099 @ 1h [sig]; 60s/AI: peak dR2 0.079 @ 1h [sig]. **next-window is COINCIDENT; there is no next-day forecast here (deferred Track B).**

Holistic scoreboard — every finished cov-dur x horizon x metric, color-coded

GREEN = dial adds OOS value / robust / capacity helps / significant			AMBER = small / coincident / borderline / not significant		RED = dial hurts / dR2<0 / metric-fragile		
cell (horizon · metric)	best learner	dR2 over persistence [CI95]	sig	best / linear test R2	capacity gain	QLIKE	LE-vs-AI
Sigma = 30s covariance window							
horizon=nextwin · LE	svr	0.075 [0.070,0.080]	sig	0.502 / 0.471	+0.031	0.777	robust
horizon=nextwin · AI	svr	0.075 [0.070,0.079]	sig	0.502 / 0.468	+0.034	0.747	robust
horizon=30s · LE	svr	0.083 [0.078,0.087]	sig	0.505 / 0.468	+0.037	0.968	robust
horizon=30s · AI	svr	0.081 [0.077,0.086]	sig	0.504 / 0.465	+0.039	0.921	robust
horizon=45s · LE	svr	0.073 [0.066,0.079]	sig	0.495 / 0.468	+0.027	0.739	robust
horizon=45s · AI	svr	0.075 [0.070,0.079]	sig	0.497 / 0.465	+0.032	0.706	robust
horizon=60s · LE	svr	0.073 [0.067,0.080]	sig	0.558 / 0.533	+0.025	0.365	robust
horizon=60s · AI	svr	0.075 [0.070,0.081]	sig	0.560 / 0.529	+0.031	0.360	robust
horizon=2m · LE	svr	0.094 [0.087,0.101]	sig	0.629 / 0.587	+0.042	0.309	robust
horizon=2m · AI	svr	0.089 [0.082,0.096]	sig	0.624 / 0.583	+0.042	0.302	robust
horizon=5m · LE	svr	0.110 [0.102,0.118]	sig	0.682 / 0.627	+0.055	0.226	robust
horizon=5m · AI	svr	0.100 [0.091,0.108]	sig	0.671 / 0.622	+0.050	0.230	robust
horizon=10m · LE	svr	0.110 [0.100,0.119]	sig	0.682 / 0.630	+0.052	0.200	robust
horizon=10m · AI	svr	0.101 [0.090,0.111]	sig	0.674 / 0.625	+0.048	0.204	robust
horizon=15m · LE	svr	0.113 [0.103,0.123]	sig	0.678 / 0.625	+0.053	0.194	robust
horizon=15m · AI	svr	0.103 [0.091,0.114]	sig	0.668 / 0.620	+0.048	0.198	robust
horizon=30m · LE	huber	0.118 [0.113,0.123]	sig	0.652 / 0.601	+0.051	0.197	robust
horizon=30m · AI	huber	0.111 [0.106,0.116]	sig	0.645 / 0.595	+0.050	0.201	robust
horizon=1h · LE	huber	0.136 [0.129,0.143]	sig	0.621 / 0.549	+0.072	0.200	robust
horizon=1h · AI	huber	0.129 [0.123,0.136]	sig	0.614 / 0.543	+0.071	0.204	robust
Sigma = 45s covariance window							
horizon=nextwin · LE	svr	0.053 [0.046,0.060]	sig	0.576 / 0.554	+0.022	0.415	robust
horizon=nextwin · AI	svr	0.050 [0.043,0.056]	sig	0.573 / 0.550	+0.023	0.412	robust
horizon=30s · LE	N/A — horizon < cov window (H=30s ≤ Sigma=45s); not measurable						skip
horizon=30s · AI	N/A — horizon < cov window (H=30s ≤ Sigma=45s); not measurable						skip
horizon=45s · LE	svr	0.059 [0.055,0.063]	sig	0.578 / 0.552	+0.026	0.410	robust
horizon=45s · AI	svr	0.055 [0.050,0.060]	sig	0.574 / 0.549	+0.026	0.411	robust
horizon=60s · LE	svr	0.059 [0.055,0.064]	sig	0.579 / 0.552	+0.026	0.402	robust
horizon=60s · AI	svr	0.056 [0.051,0.060]	sig	0.575 / 0.548	+0.027	0.400	robust
horizon=2m · LE	svr	0.064 [0.059,0.070]	sig	0.635 / 0.605	+0.030	0.316	robust
horizon=2m · AI	svr	0.060 [0.054,0.065]	sig	0.631 / 0.601	+0.030	0.316	robust
horizon=5m · LE	svr	0.071 [0.064,0.079]	sig	0.702 / 0.667	+0.034	0.219	robust
horizon=5m · AI	svr	0.068 [0.060,0.075]	sig	0.699 / 0.663	+0.036	0.221	robust
horizon=10m · LE	svr	0.090 [0.083,0.098]	sig	0.727 / 0.677	+0.050	0.179	robust
horizon=10m · AI	svr	0.080 [0.071,0.088]	sig	0.717 / 0.672	+0.045	0.184	robust
horizon=15m · LE	svr	0.092 [0.084,0.102]	sig	0.722 / 0.672	+0.050	0.171	robust
horizon=15m · AI	svr	0.078 [0.069,0.087]	sig	0.708 / 0.667	+0.041	0.178	robust
horizon=30m · LE	svr	0.093 [0.081,0.104]	sig	0.691 / 0.647	+0.045	0.174	robust
horizon=30m · AI	huber	0.083 [0.079,0.088]	sig	0.682 / 0.641	+0.041	0.183	robust
horizon=1h · LE	huber	0.107 [0.102,0.113]	sig	0.656 / 0.594	+0.062	0.184	robust
horizon=1h · AI	huber	0.099 [0.094,0.105]	sig	0.648 / 0.586	+0.062	0.188	robust
Sigma = 60s covariance window							
horizon=nextwin · LE	svr	0.038 [0.031,0.045]	sig	0.626 / 0.610	+0.015	0.375	robust
horizon=nextwin · AI	svr	0.037 [0.032,0.042]	sig	0.624 / 0.607	+0.017	0.371	robust
horizon=30s · LE	N/A — horizon < cov window (H=30s ≤ Sigma=60s); not measurable						skip
horizon=30s · AI	N/A — horizon < cov window (H=30s ≤ Sigma=60s); not measurable						skip
horizon=45s · LE	N/A — horizon < cov window (H=45s ≤ Sigma=60s); not measurable						skip
horizon=45s · AI	N/A — horizon < cov window (H=45s ≤ Sigma=60s); not measurable						skip
horizon=60s · LE	svr	0.040 [0.033,0.047]	sig	0.628 / 0.612	+0.016	0.356	robust
horizon=60s · AI	svr	0.038 [0.031,0.044]	sig	0.625 / 0.608	+0.017	0.356	robust
horizon=2m · LE	svr	0.045 [0.040,0.051]	sig	0.676 / 0.656	+0.020	0.283	robust
horizon=2m · AI	svr	0.040 [0.034,0.045]	sig	0.671 / 0.652	+0.018	0.285	robust
horizon=5m · LE	svr	0.058 [0.052,0.065]	sig	0.733 / 0.703	+0.031	0.200	robust
horizon=5m · AI	svr	0.049 [0.042,0.057]	sig	0.725 / 0.698	+0.027	0.203	robust
horizon=10m · LE	svr	0.066 [0.059,0.074]	sig	0.748 / 0.711	+0.037	0.167	robust
horizon=10m · AI	svr	0.054 [0.045,0.061]	sig	0.736 / 0.706	+0.029	0.173	robust
horizon=15m · LE	svr	0.061 [0.052,0.070]	sig	0.737 / 0.707	+0.031	0.162	robust

cell (horizon · metric)	best learner	dR2 over persistence [CI95]	sig	best / linear test R2	capacity gain	QLIKE	LE-vs-AI
horizon=15m · AI	svr	0.056 [0.046,0.065]	sig	0.732 / 0.702	+0.031	0.165	robust
horizon=30m · LE	svr	0.072 [0.062,0.082]	sig	0.717 / 0.681	+0.036	0.161	robust
horizon=30m · AI	huber	0.065 [0.061,0.069]	sig	0.710 / 0.675	+0.035	0.170	robust
horizon=1h · LE	huber	0.087 [0.083,0.092]	sig	0.683 / 0.627	+0.055	0.171	robust
horizon=1h · AI	huber	0.079 [0.075,0.085]	sig	0.675 / 0.620	+0.055	0.175	robust

dR2 GREEN = positive AND 95% CI clears 0 (dial leads persistence here); AMBER = positive but not significant (coincident); RED = negative. Capacity gain = best-learner minus linear test R2. LE-vs-AI shows "pending" until the matching AI cell finishes (AI cells run after all LE cells).

What this establishes. Step 7 (Track A) traces the **horizon-of-predictability** curve for the covariance stress dial while holding the covariance state fixed. It converts Step-6's single next-window coincidence check into a full dR2(H) sweep and, for the first time, builds the **AI (affine-invariant) dial** that Steps 5/6 deferred - so the LE-vs-AI robustness claim is tested on a real predictor, not just a detector. The 15-learner bake-off separates "does the dial carry signal" from "does the answer need capacity" at each horizon.

Honest caveats / negatives.

- **next-window is COINCIDENT.** At H=nextwin the dial co-moves with vol and persistence dominates; a small dR2 there is expected, not a failure.
- **Longer horizons test LEADING, but overlapping targets inflate apparent skill** - the block-bootstrap CI over days (not the point estimate) is the honest significance test.
- **No next-day forecast here.** Track B (daily HAR / QLIKE / Diebold-Mariano / MCS) is a separate deferred run; Track A stays intraday (nextwin..1h).
- **The covariance state is unchanged** - no new Sigma window is built; this is a horizon sweep on the existing 30/45/60 s masked cloud.
- **PARTIAL.** This render reflects only the cells finished so far (54); re-run the builder when the detached grid completes for the full scoreboard.